



# Introduction to Neural Networks





# What is Deep Learning?

## ARTIFICIAL INTELLIGENCE

Any technique that enables computers to mimic human behavior



## MACHINE LEARNING

Ability to learn without explicitly being programmed



## DEEP LEARNING

Extract patterns from data using neural networks

3 1 3 4 7 2  
1 7 4 2 3 5



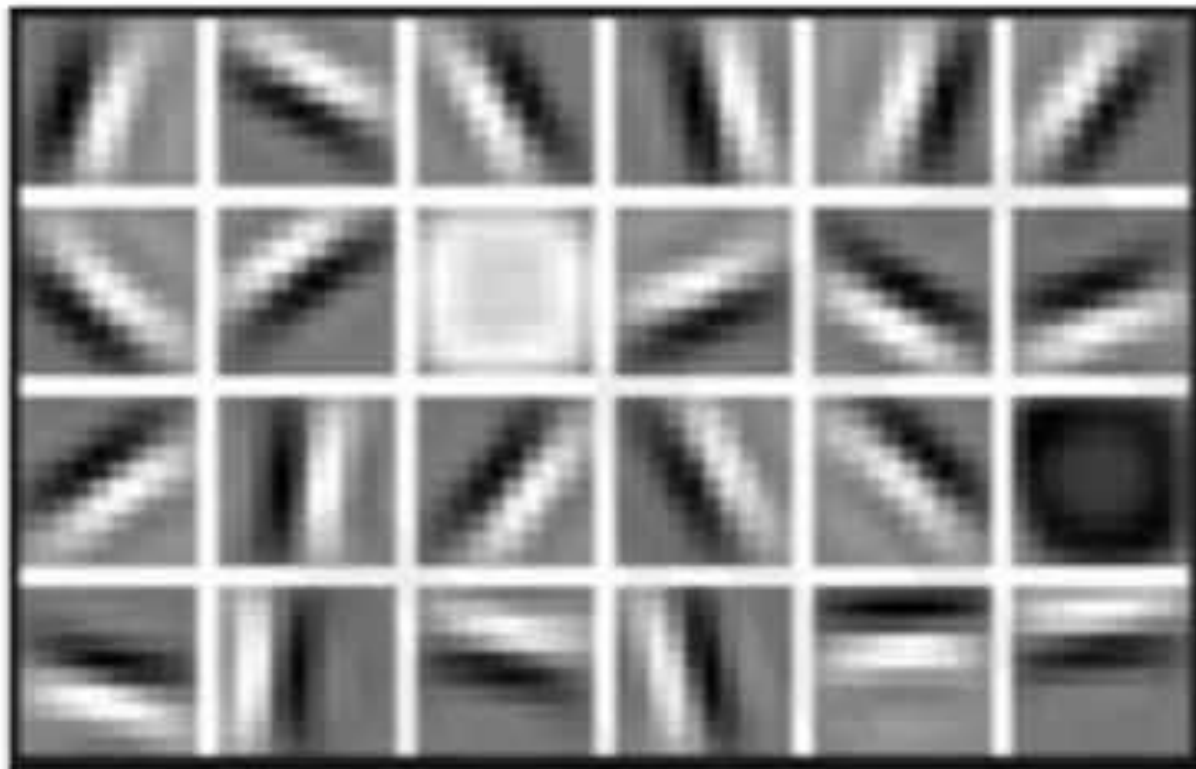
# Why Deep Learning and Why Now?

# Why Deep Learning?

Hand engineered features are time consuming, brittle, and not scalable in practice

Can we learn the **underlying features** directly from data?

Low Level Features



Lines & Edges

Mid Level Features



Eyes & Nose & Ears

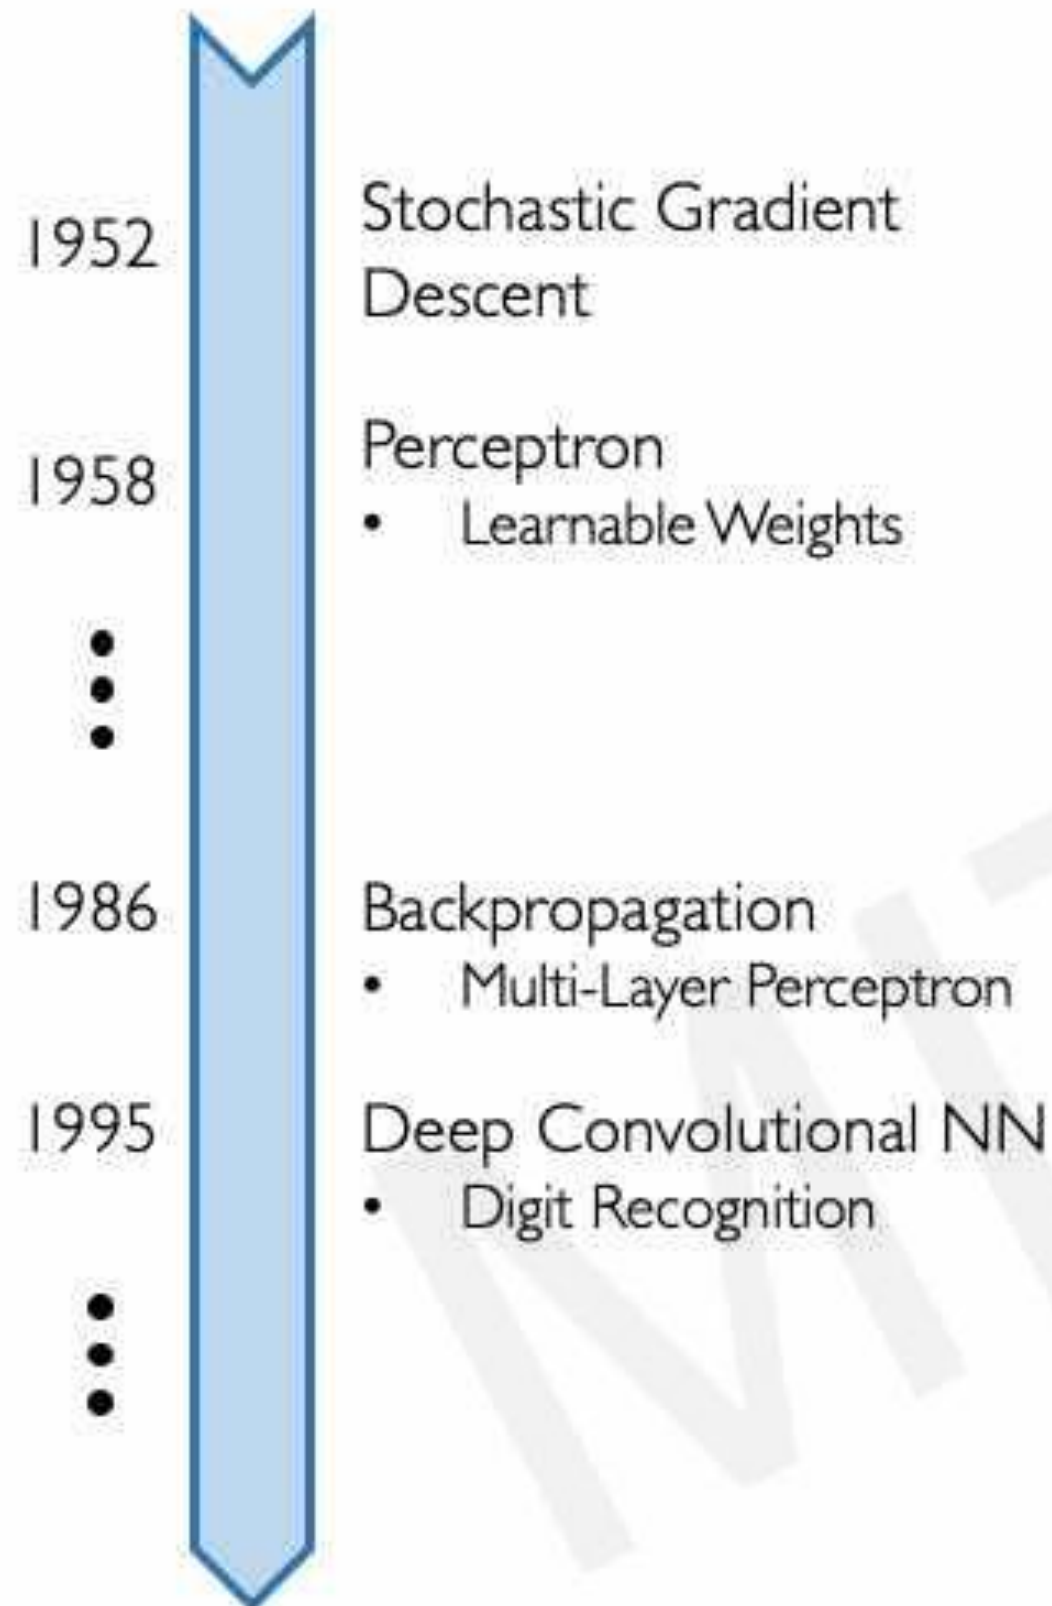
High Level Features



Facial Structure

# Why Now?

Neural Networks date back decades, so why the resurgence?



## 1. Big Data

- Larger Datasets
- Easier Collection & Storage

IMAGENET



WIKIPEDIA  
The Free Encyclopedia



## 2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable

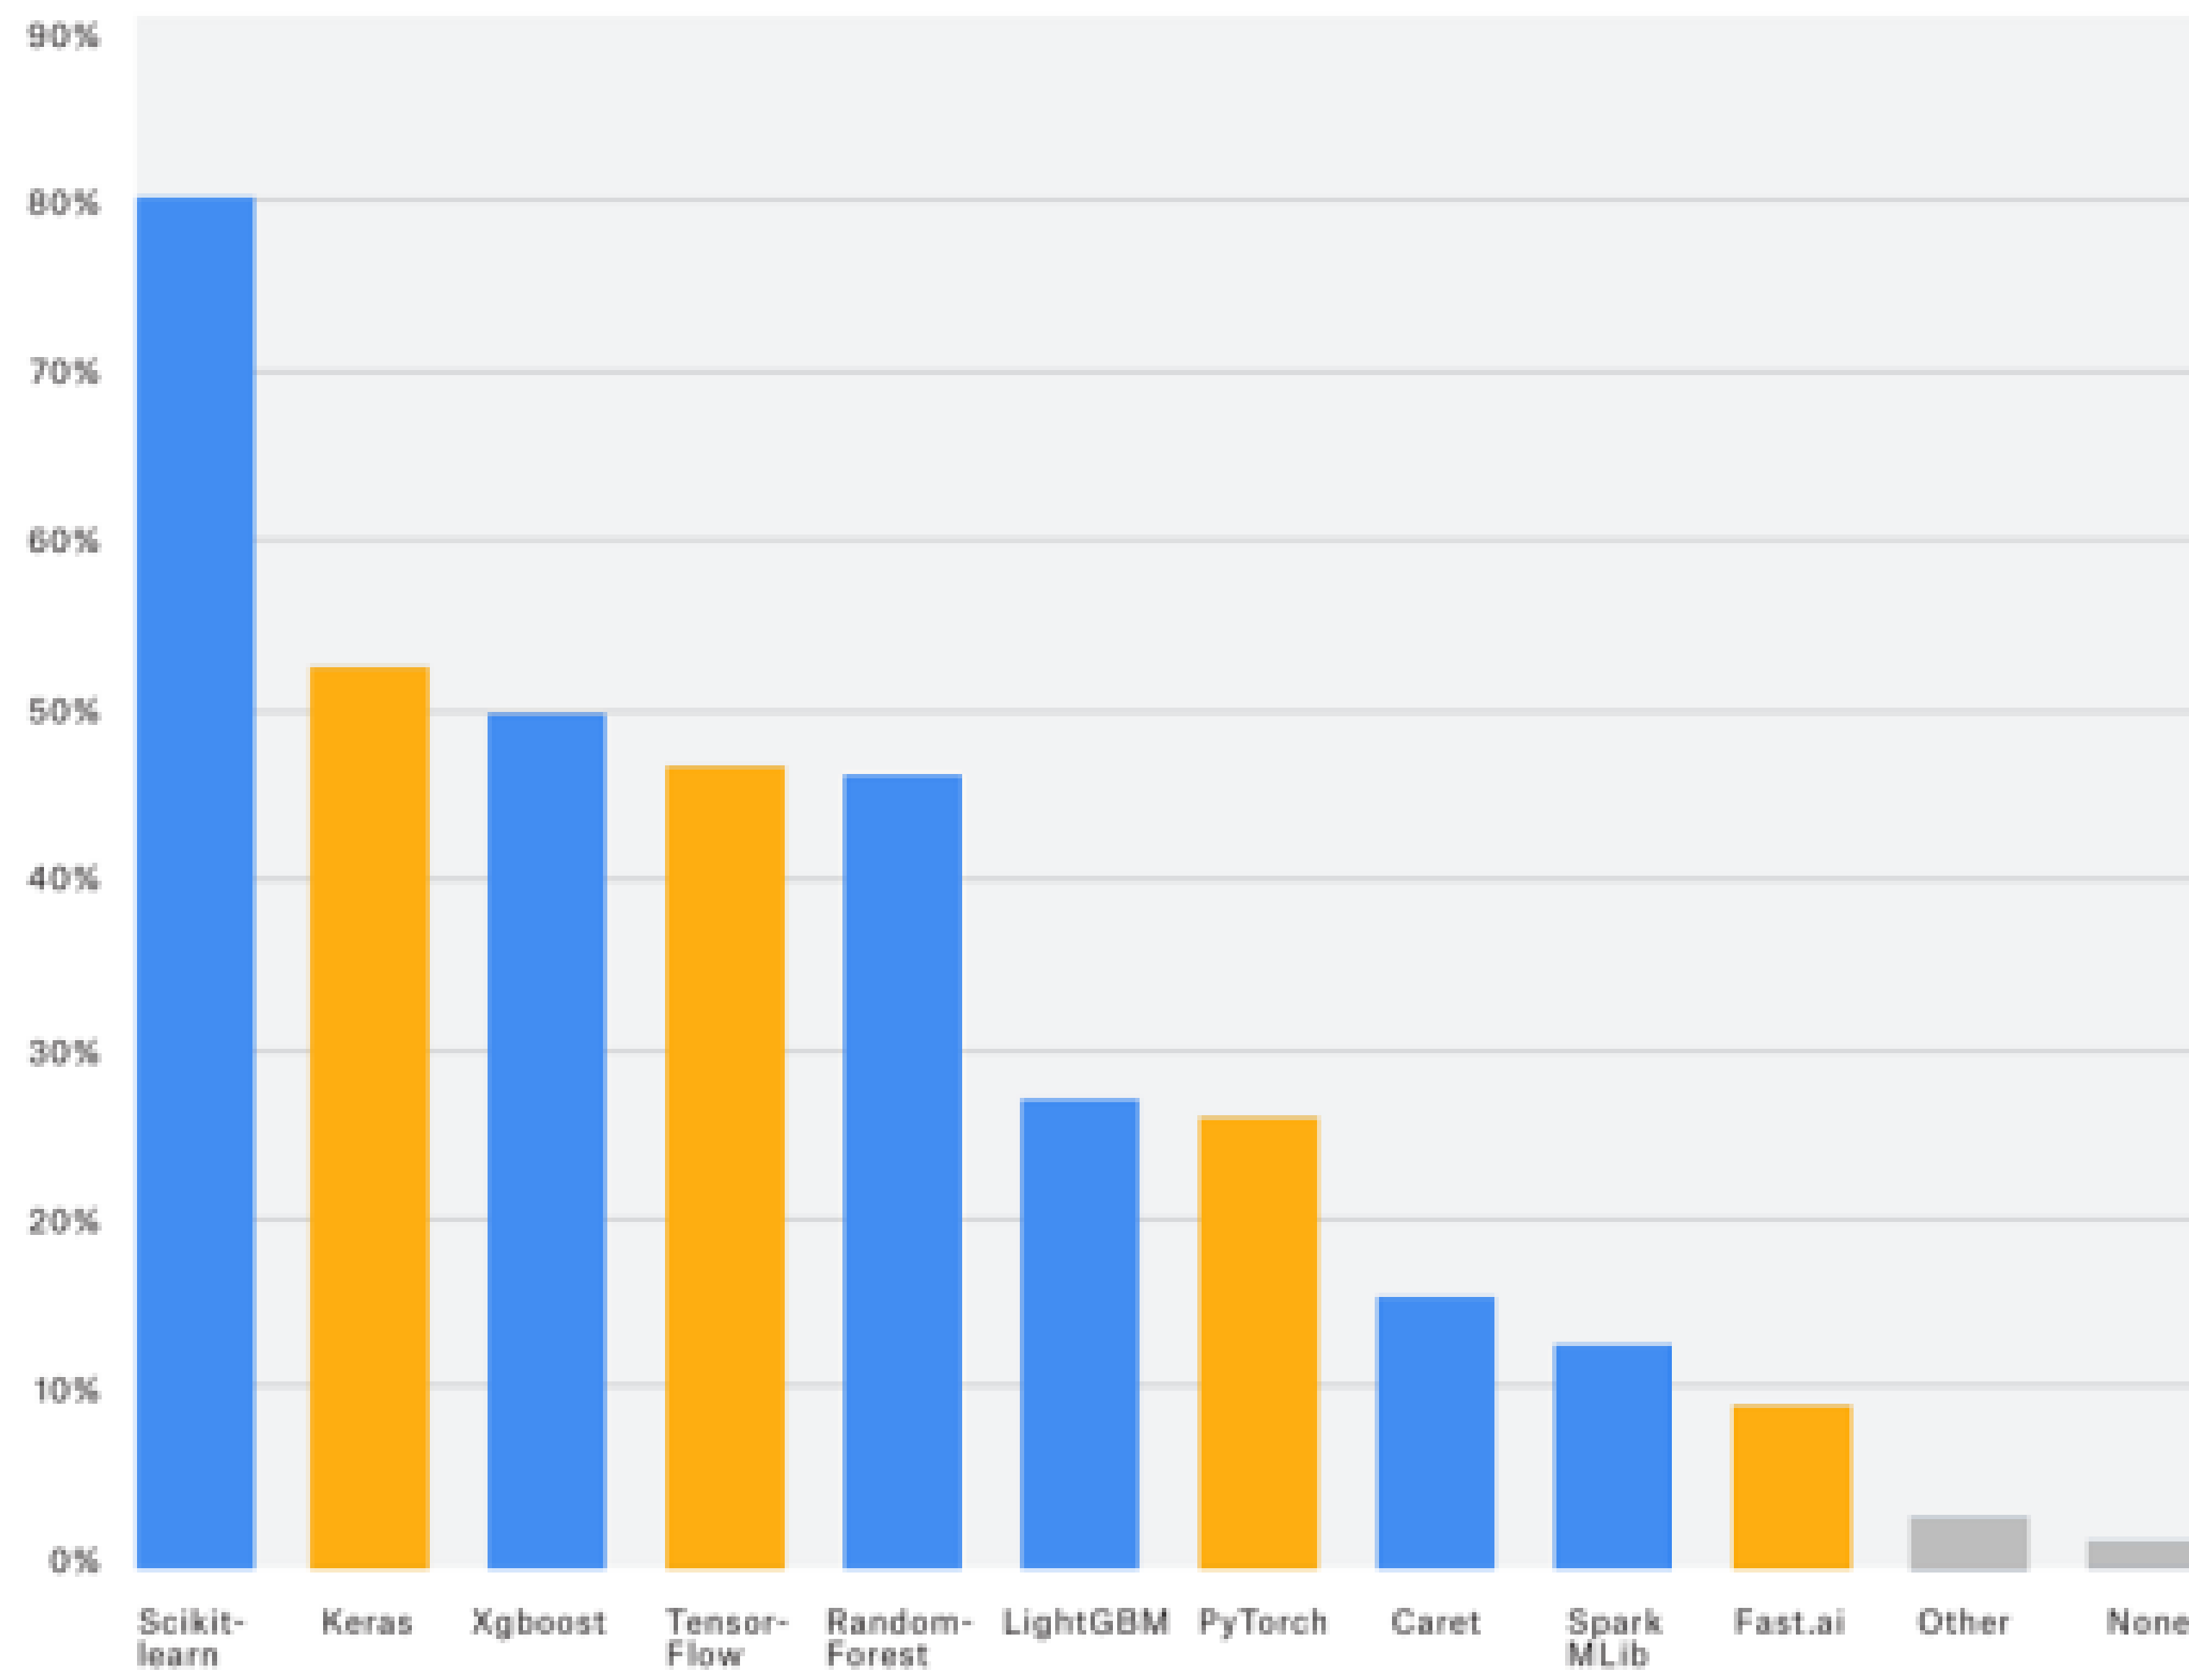


## 3. Software

- Improved Techniques
- New Models
- Toolboxes



# Trends in Deep Learning



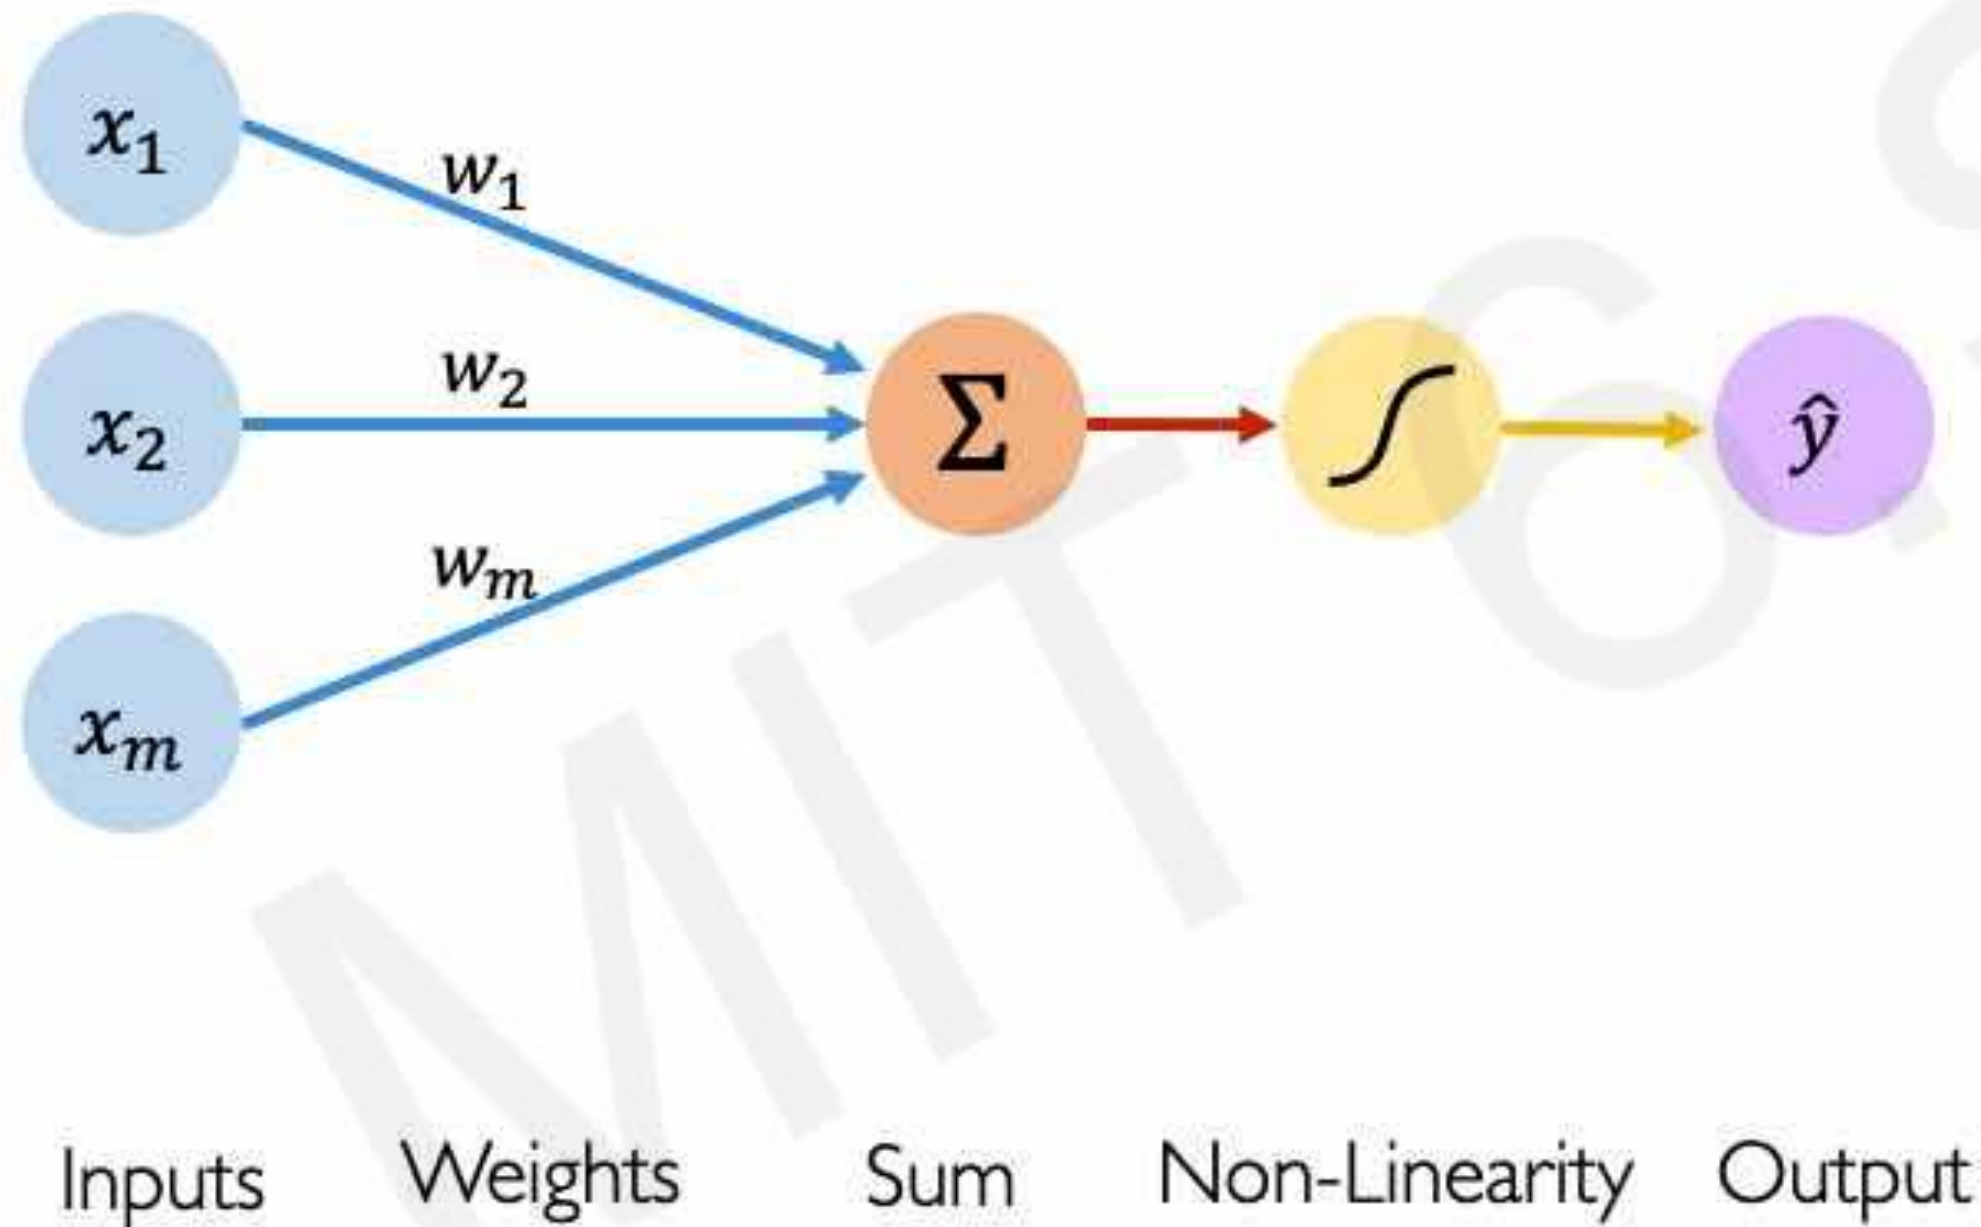


# The Perceptron

The structural building block of deep learning



# The Perceptron: Forward Propagation



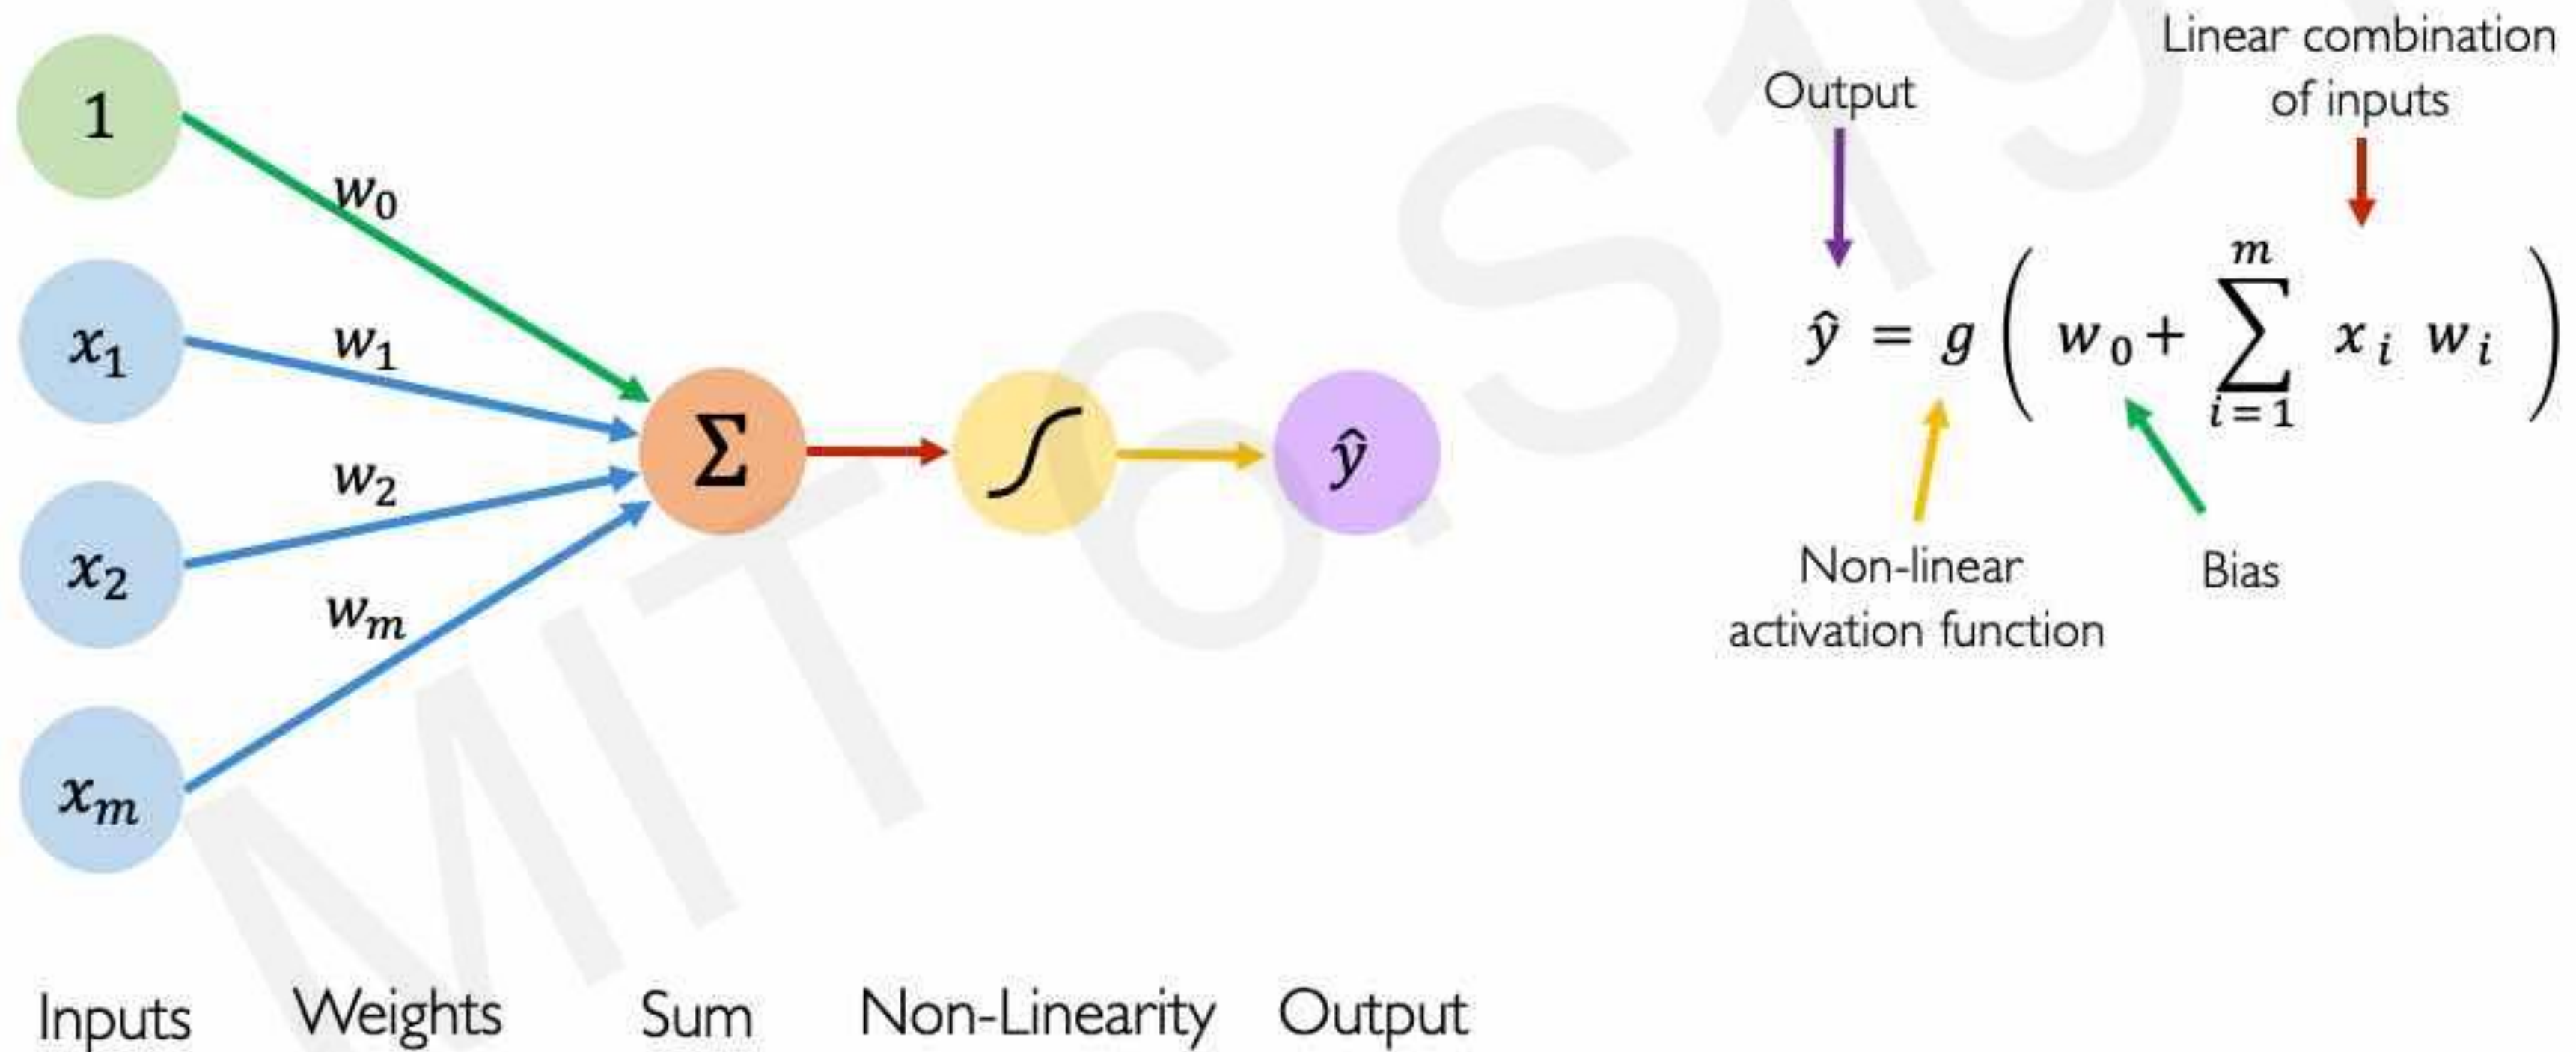
Linear combination of inputs

Output

$$\hat{y} = g \left( \sum_{i=1}^m x_i w_i \right)$$

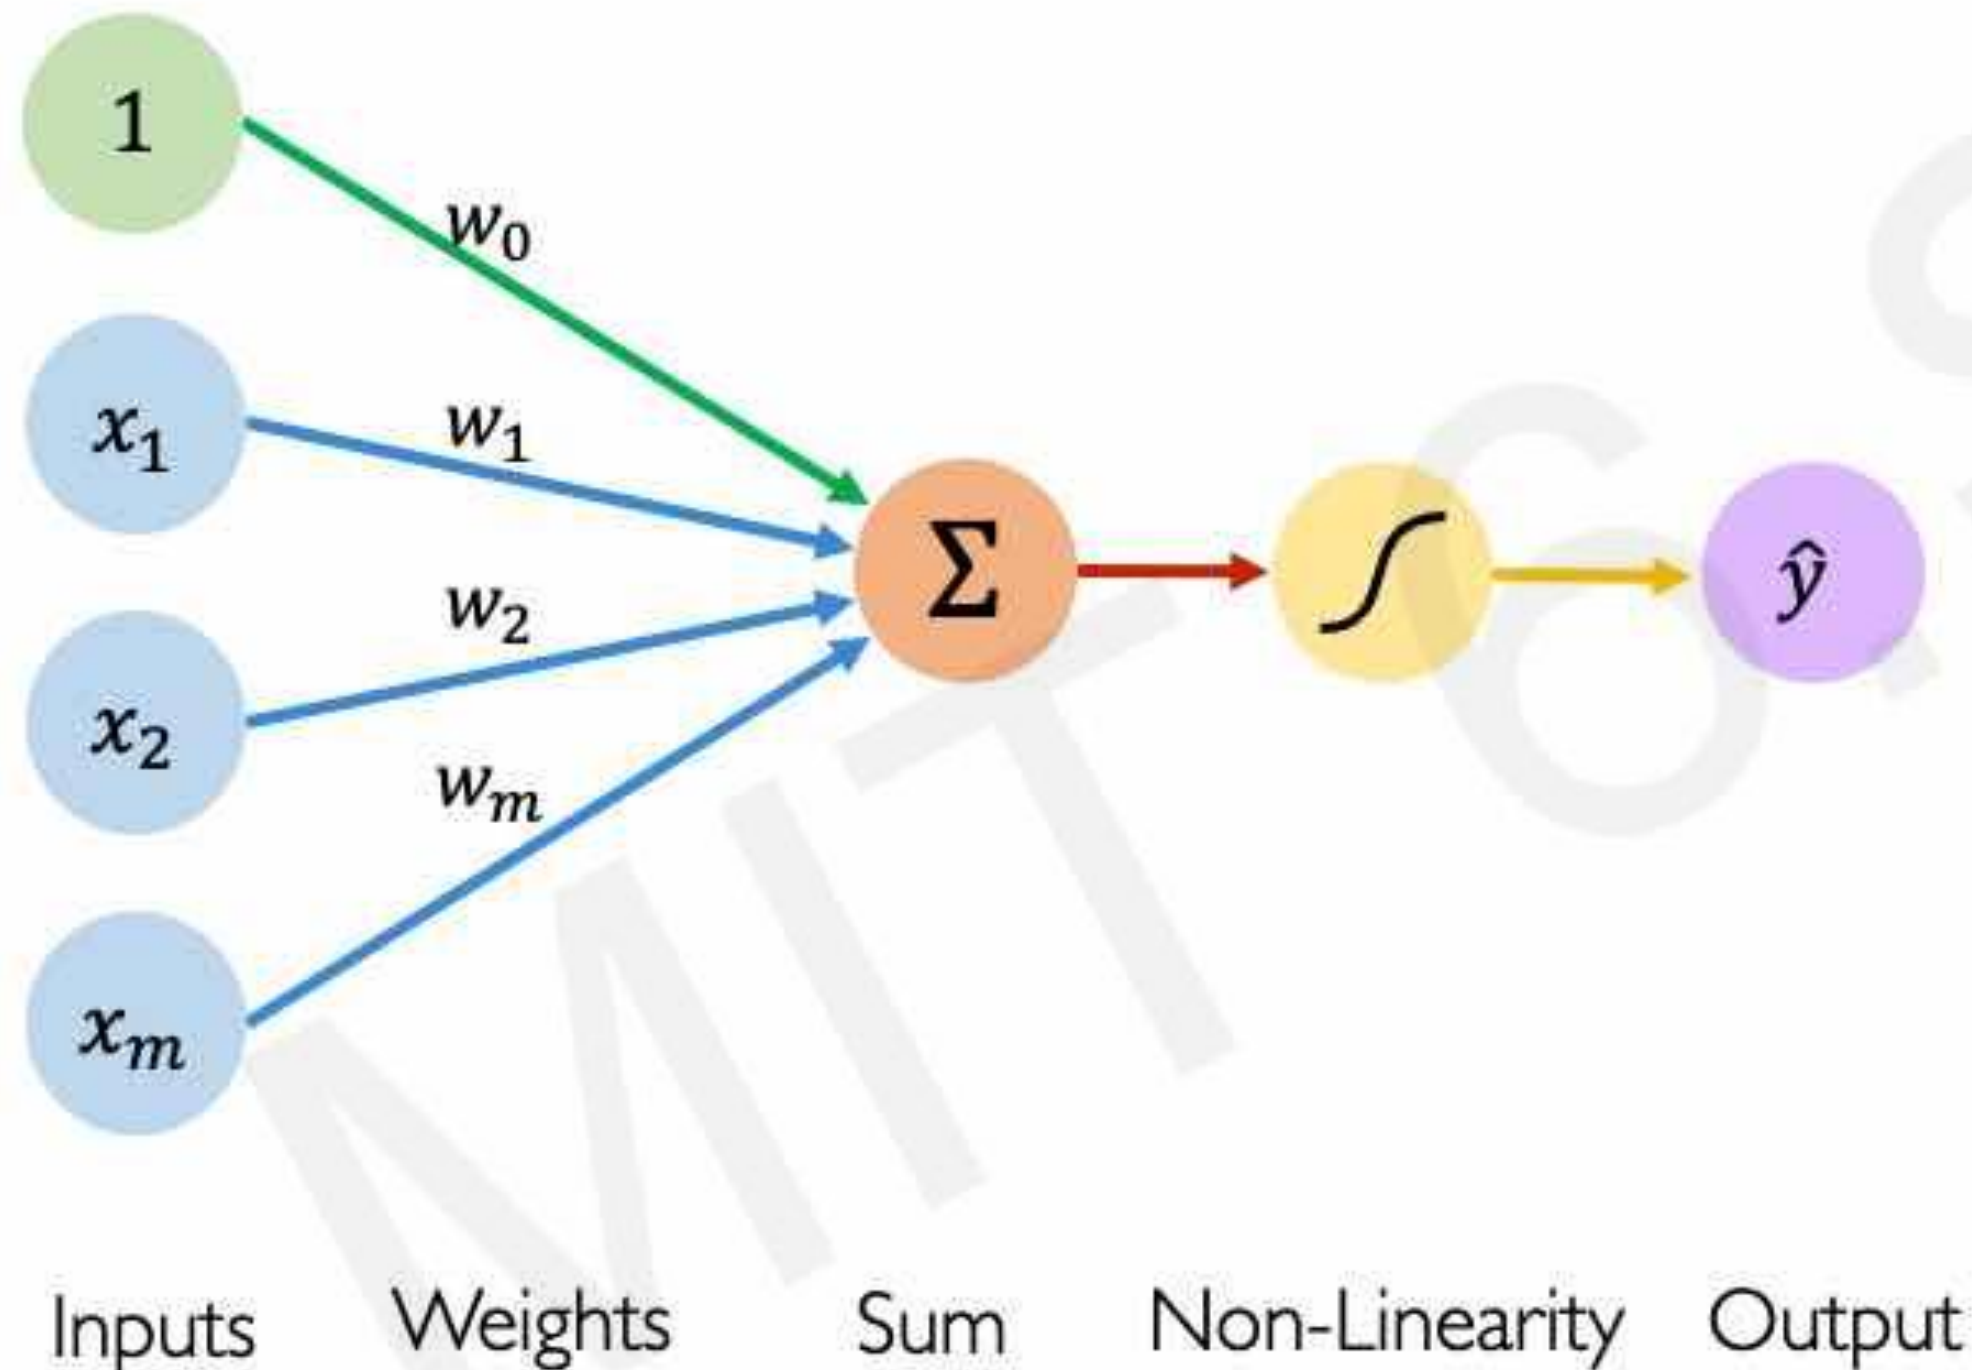
Non-linear activation function

# The Perceptron: Forward Propagation





# The Perceptron: Forward Propagation

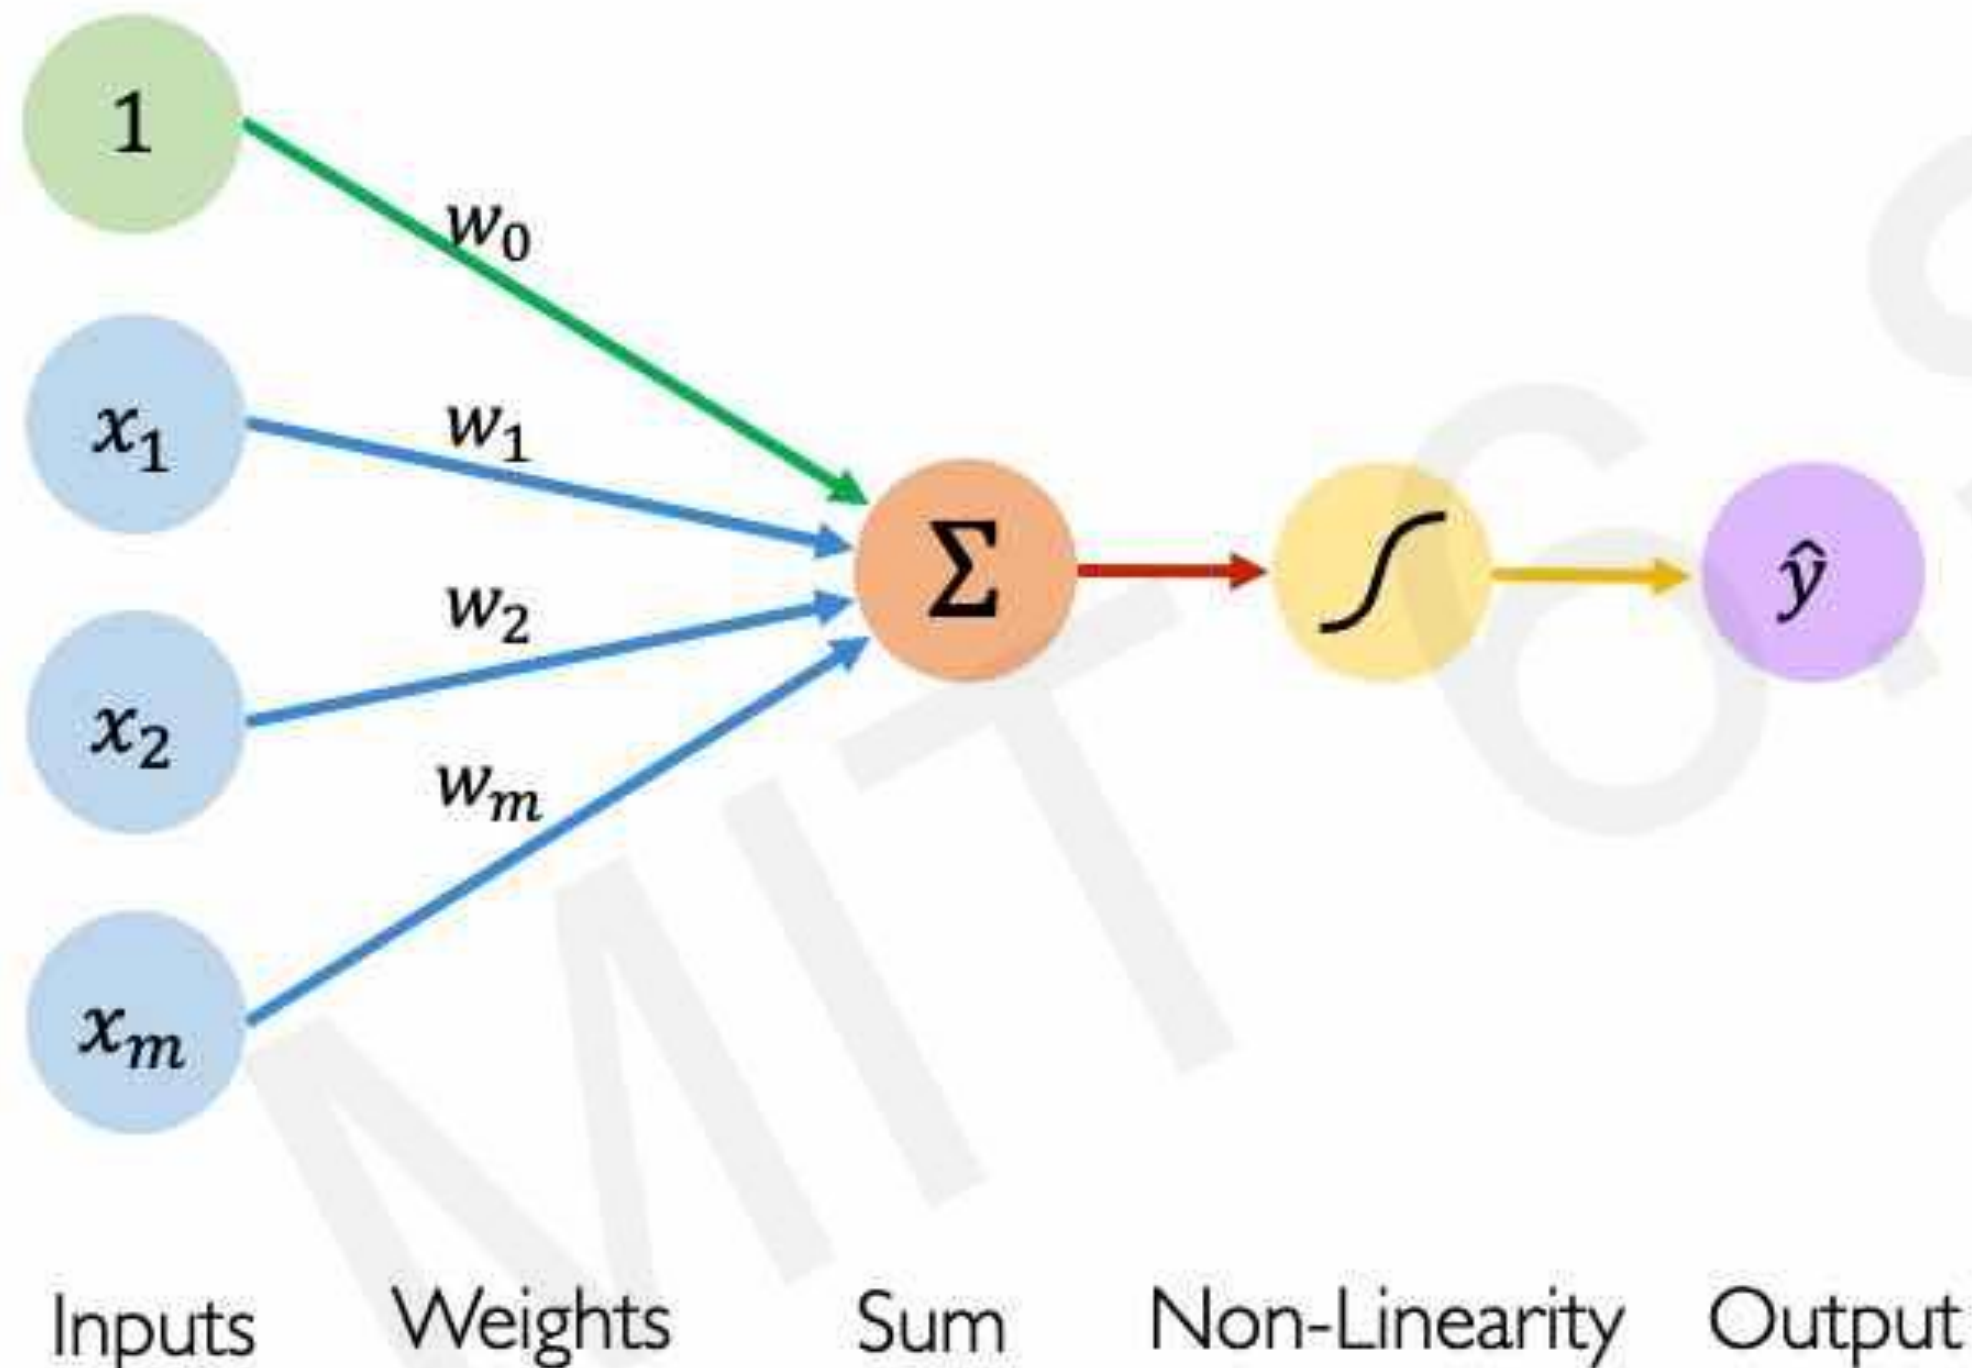


$$\hat{y} = g \left( w_0 + \sum_{i=1}^m x_i w_i \right)$$

$$\hat{y} = g ( w_0 + \mathbf{X}^T \mathbf{W} )$$

where:  $\mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$  and  $\mathbf{W} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}$

# The Perceptron: Forward Propagation

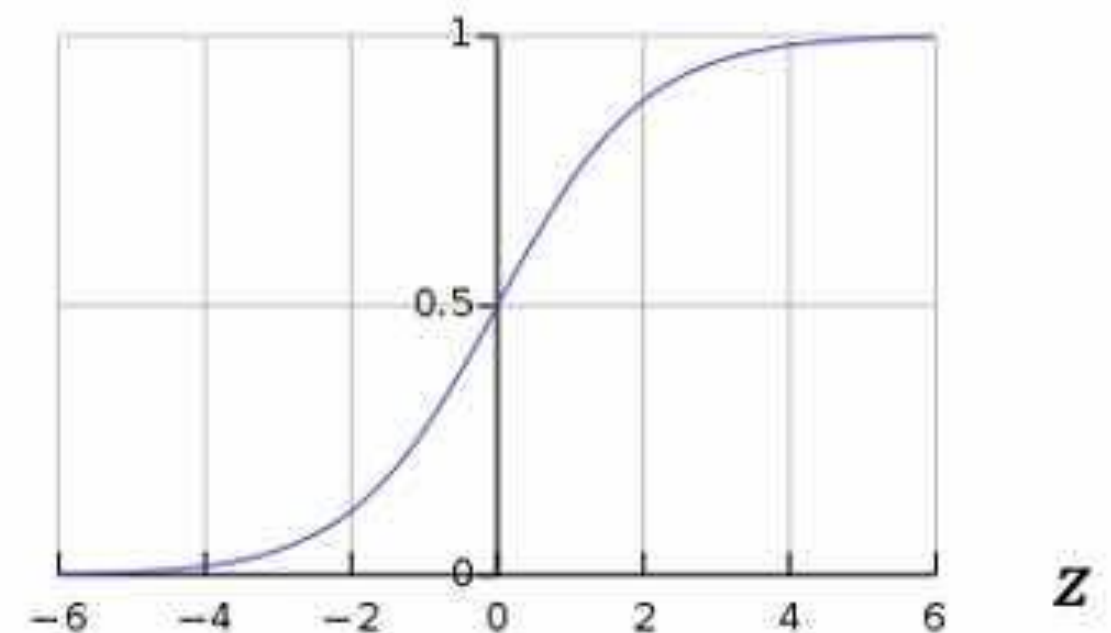


## Activation Functions

$$\hat{y} = g(w_0 + \mathbf{X}^T \mathbf{W})$$

- Example: sigmoid function

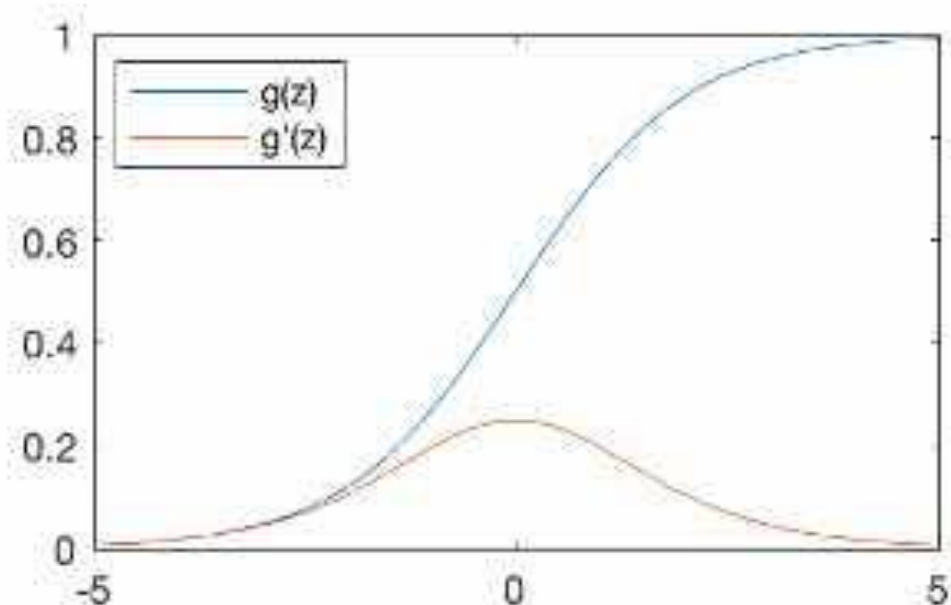
$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$





# Common Activation Functions

Sigmoid Function



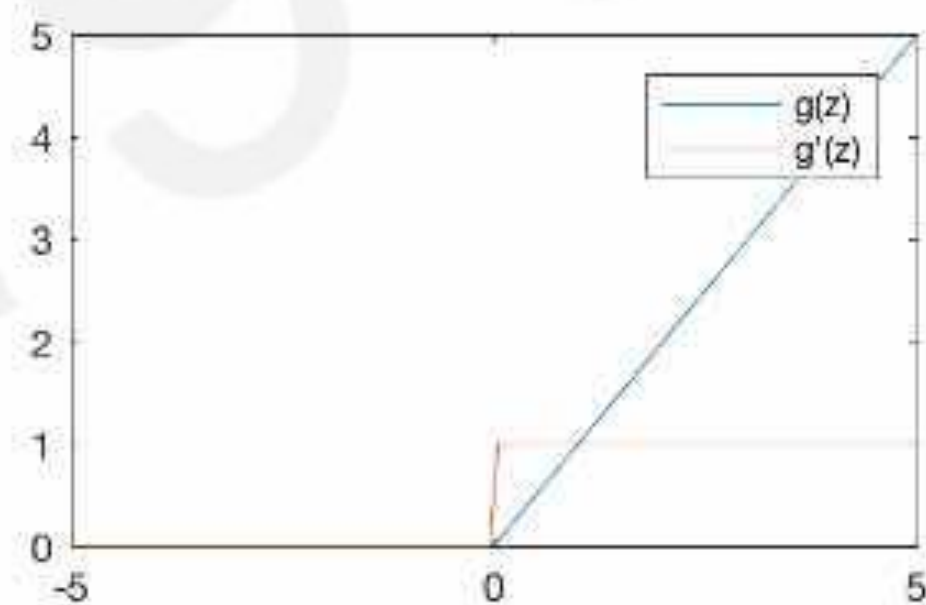
$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$



```
tf.math.sigmoid(z)
```

Rectified Linear Unit (ReLU)



$$g(z) = \max(0, z)$$

$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$

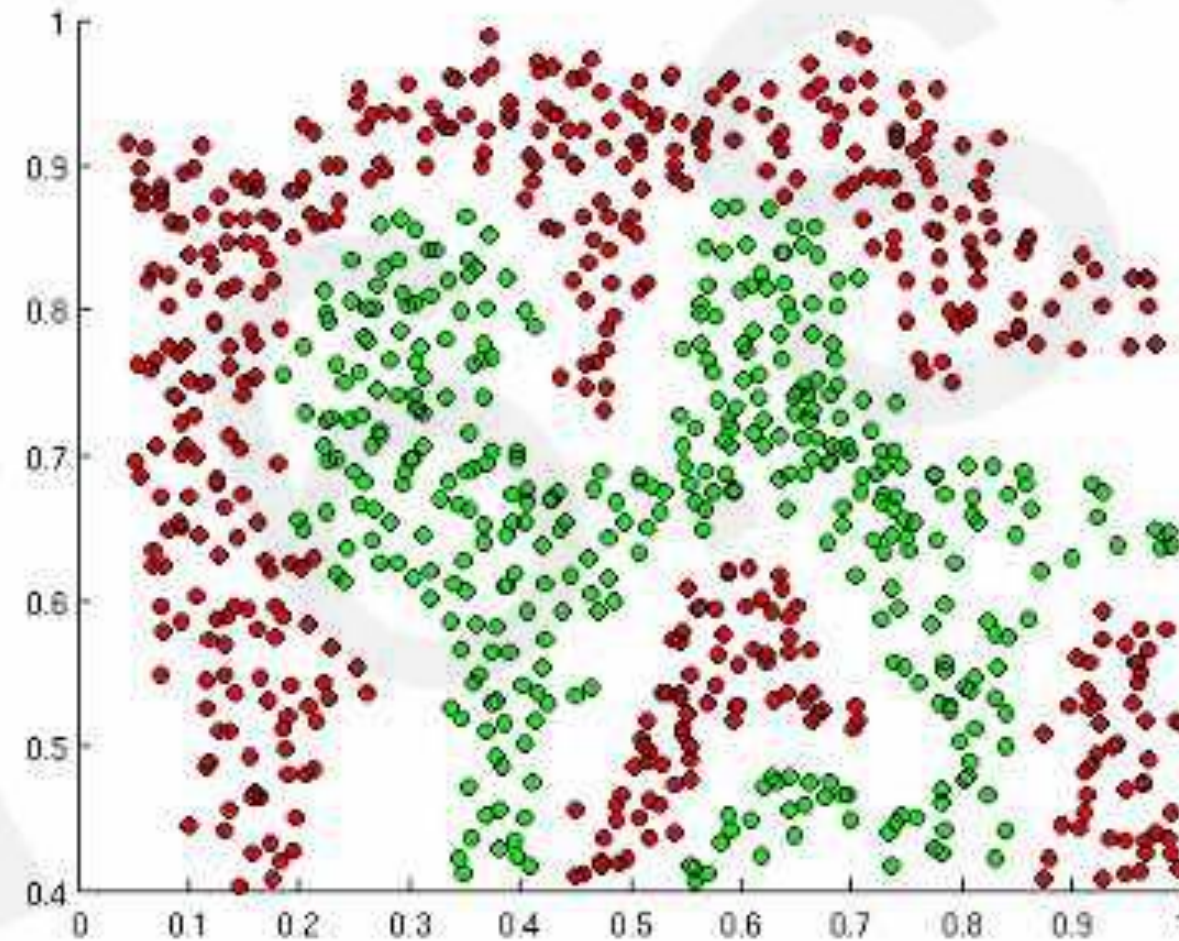


```
tf.nn.relu(z)
```



# Importance of Activation Functions

The purpose of activation functions is to **introduce non-linearities** into the network

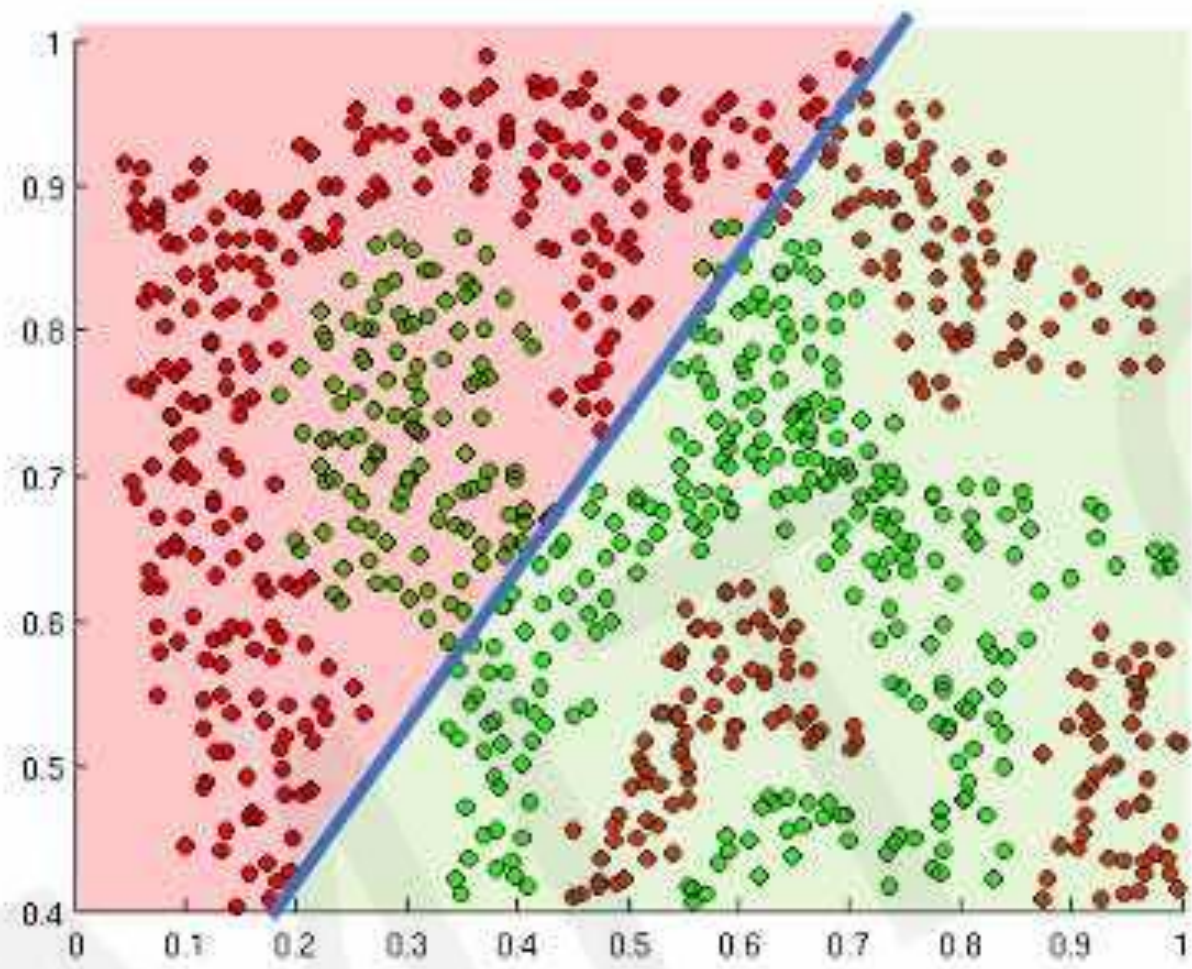


What if we wanted to build a neural network to distinguish green vs red points?



# Importance of Activation Functions

The purpose of activation functions is to **introduce non-linearities** into the network

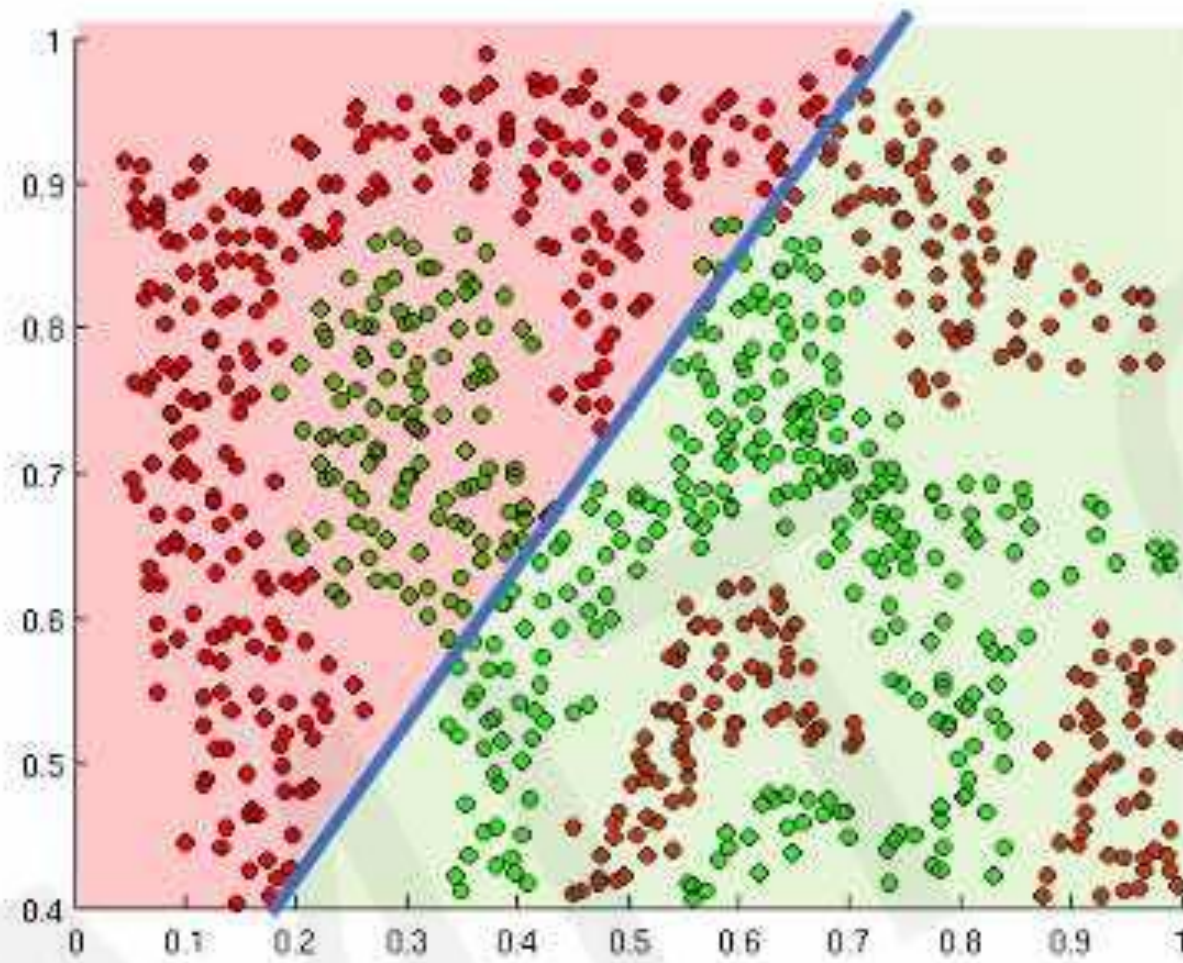


Linear activation functions produce linear decisions no matter the network size

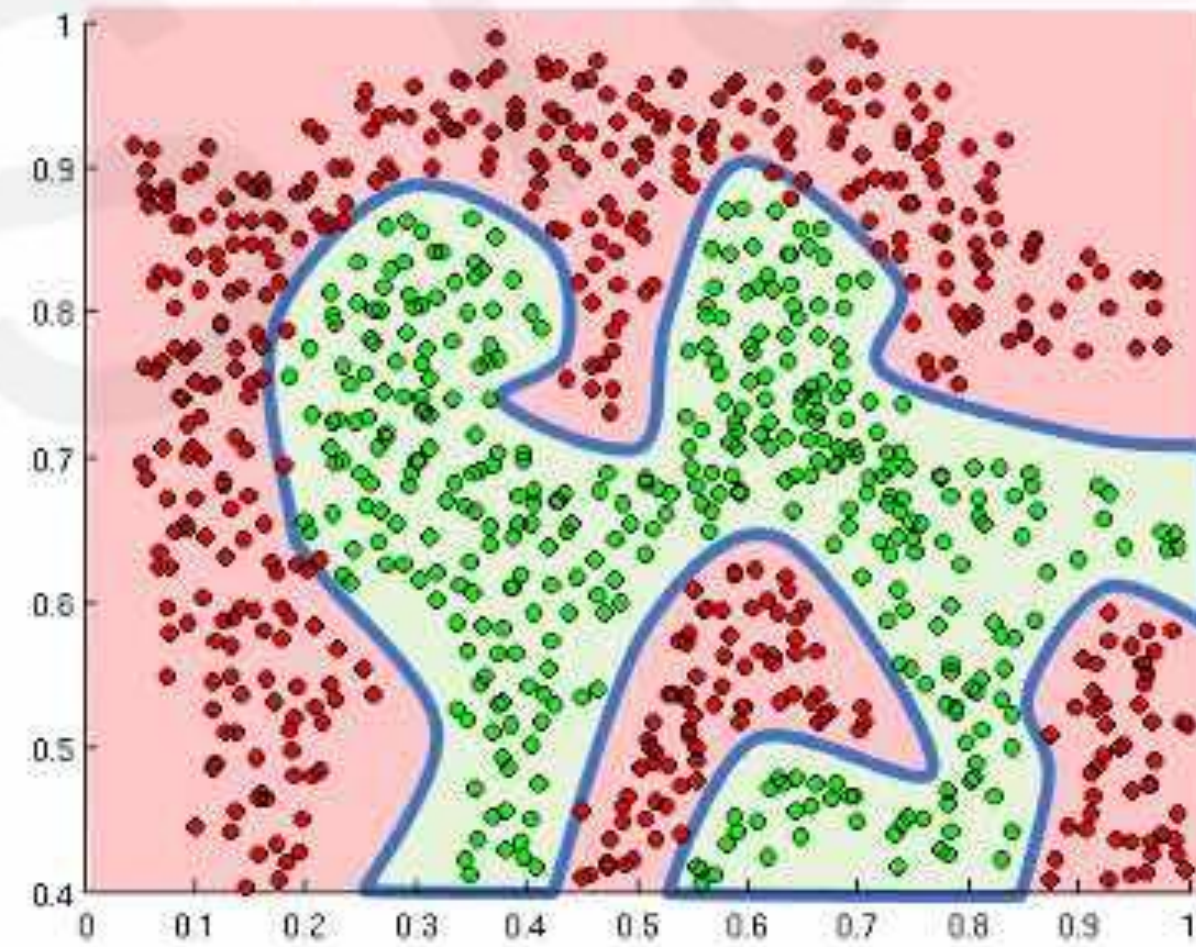


# Importance of Activation Functions

*The purpose of activation functions is to **introduce non-linearities** into the network*



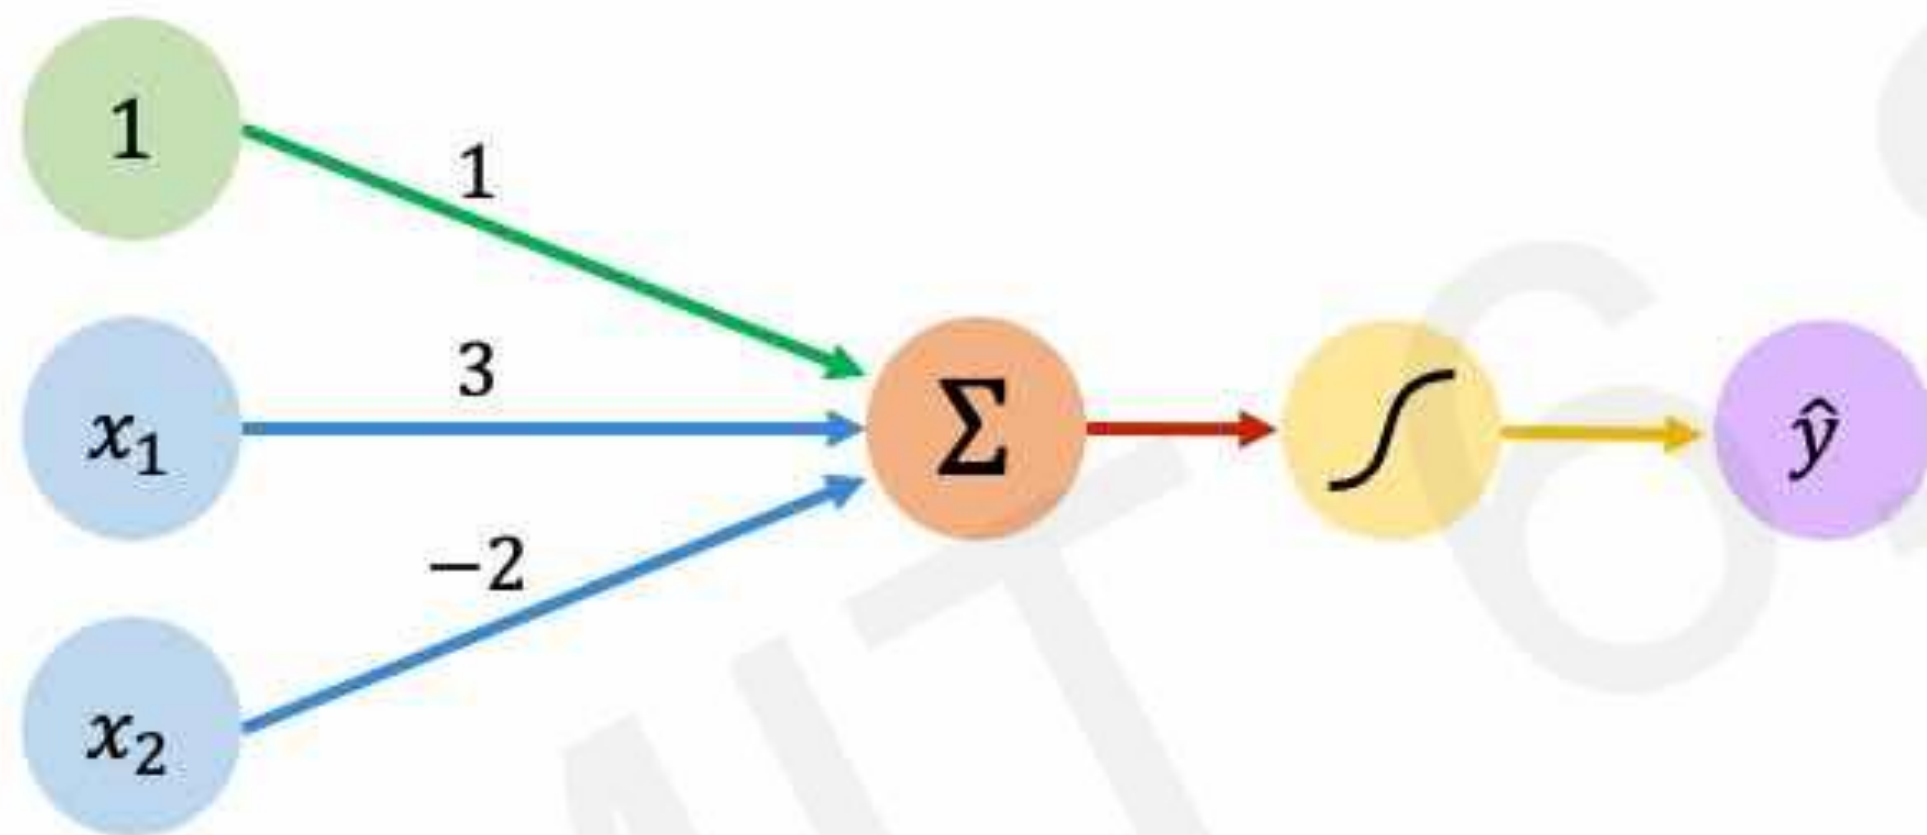
Linear activation functions produce linear decisions no matter the network size



Non-linearities allow us to approximate arbitrarily complex functions



# The Perceptron: Example

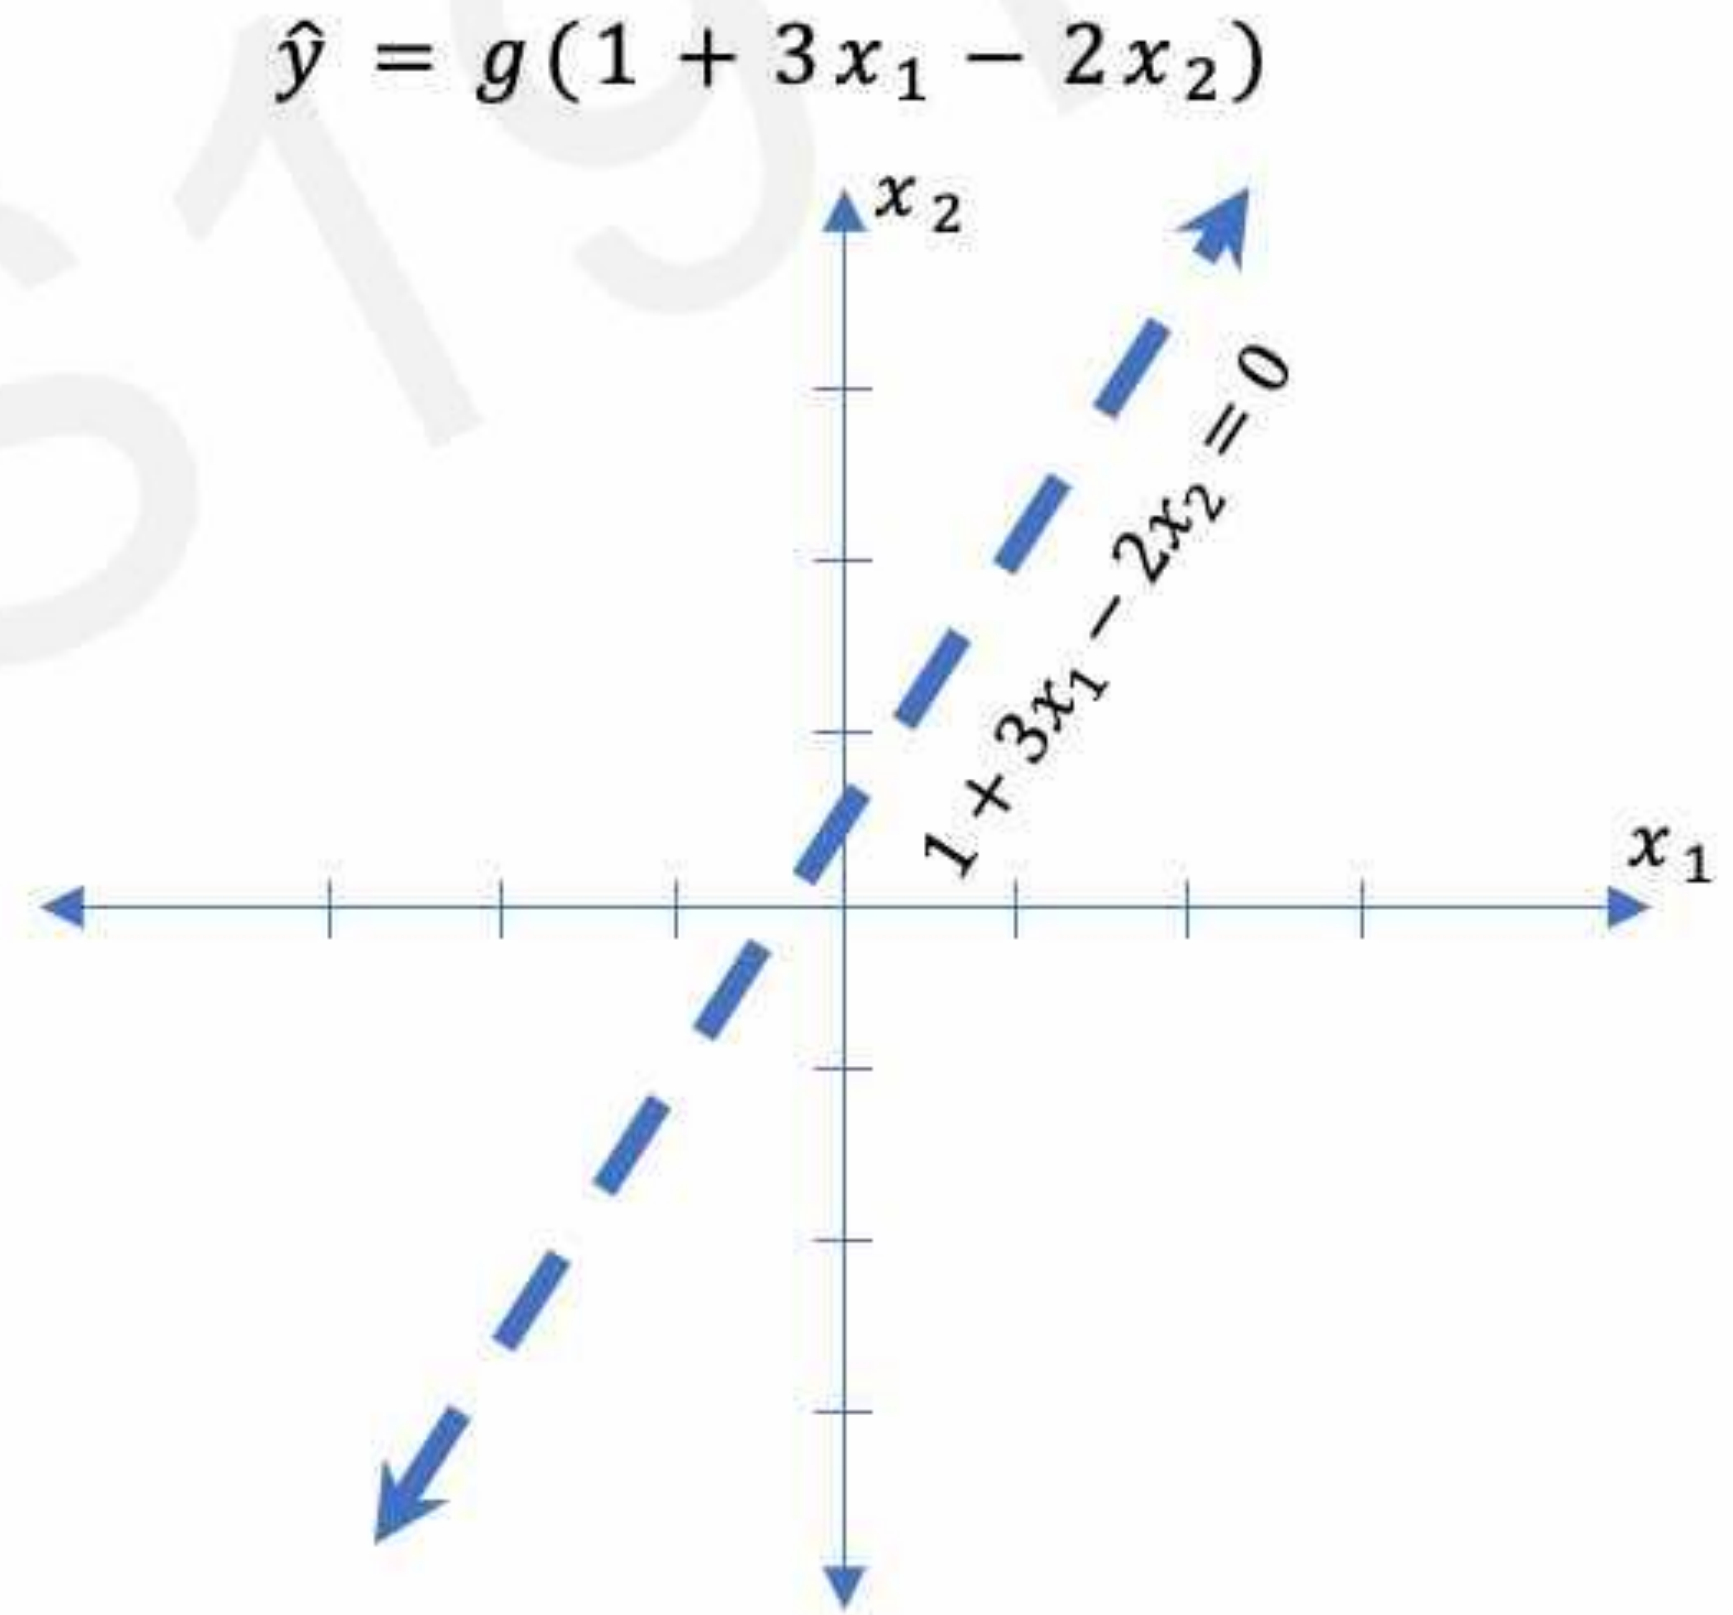
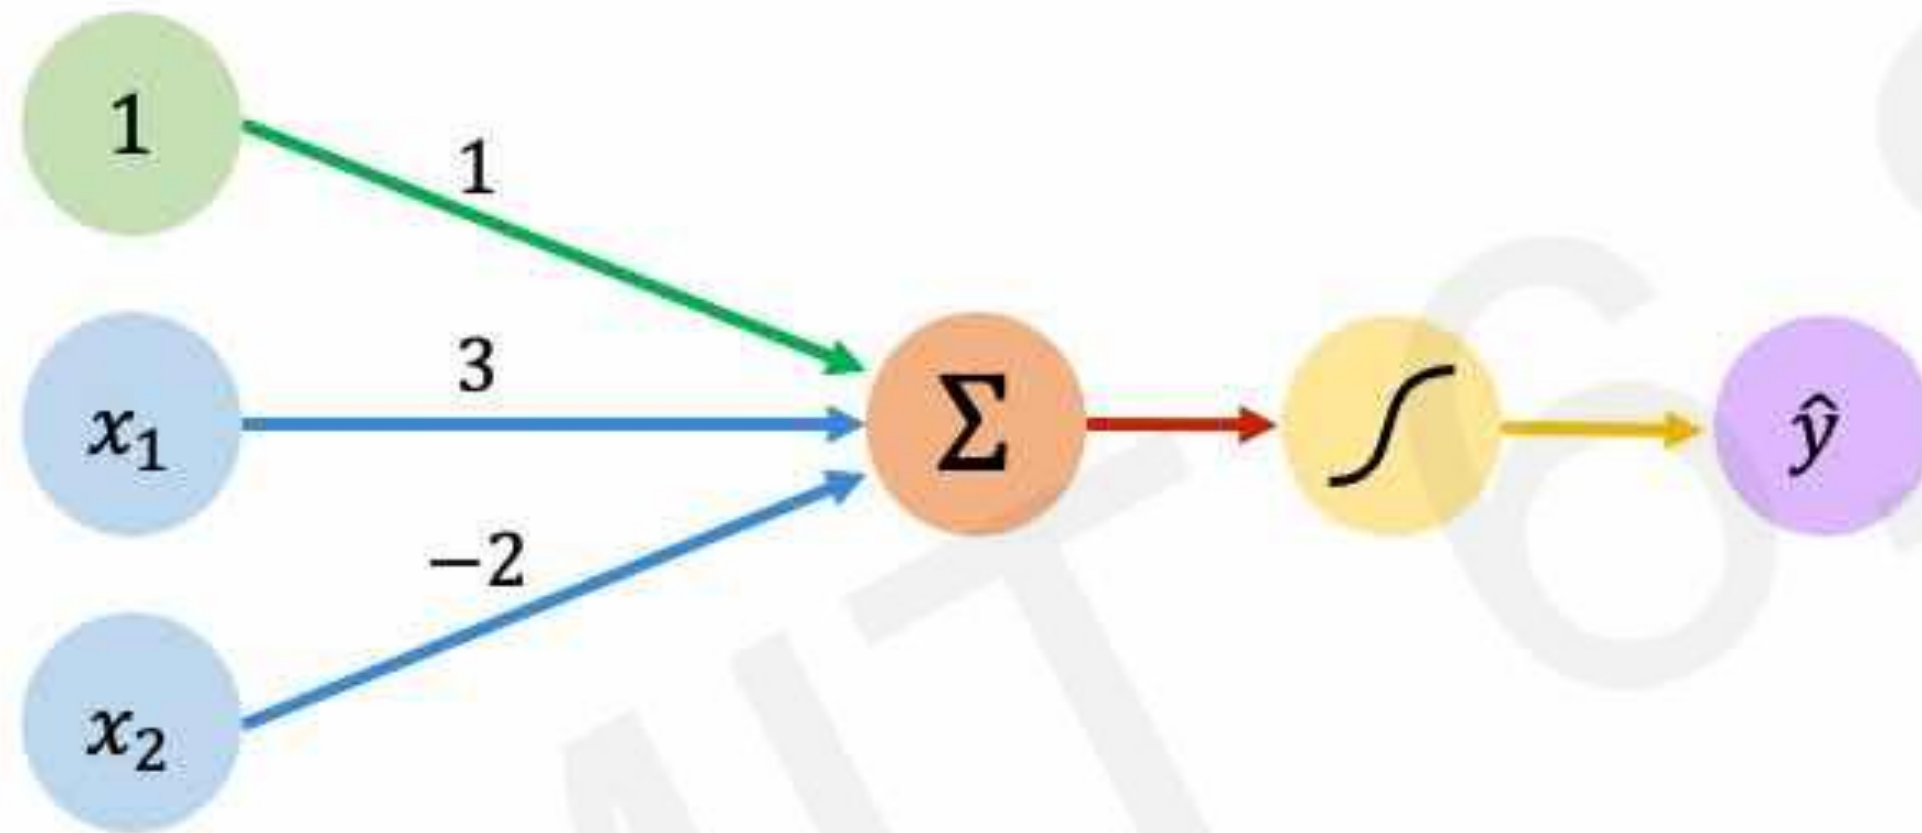


We have:  $w_0 = 1$  and  $\mathbf{W} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$

$$\begin{aligned}\hat{y} &= g(w_0 + \mathbf{X}^T \mathbf{W}) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ \hat{y} &= g(1 + 3x_1 - 2x_2)\end{aligned}$$

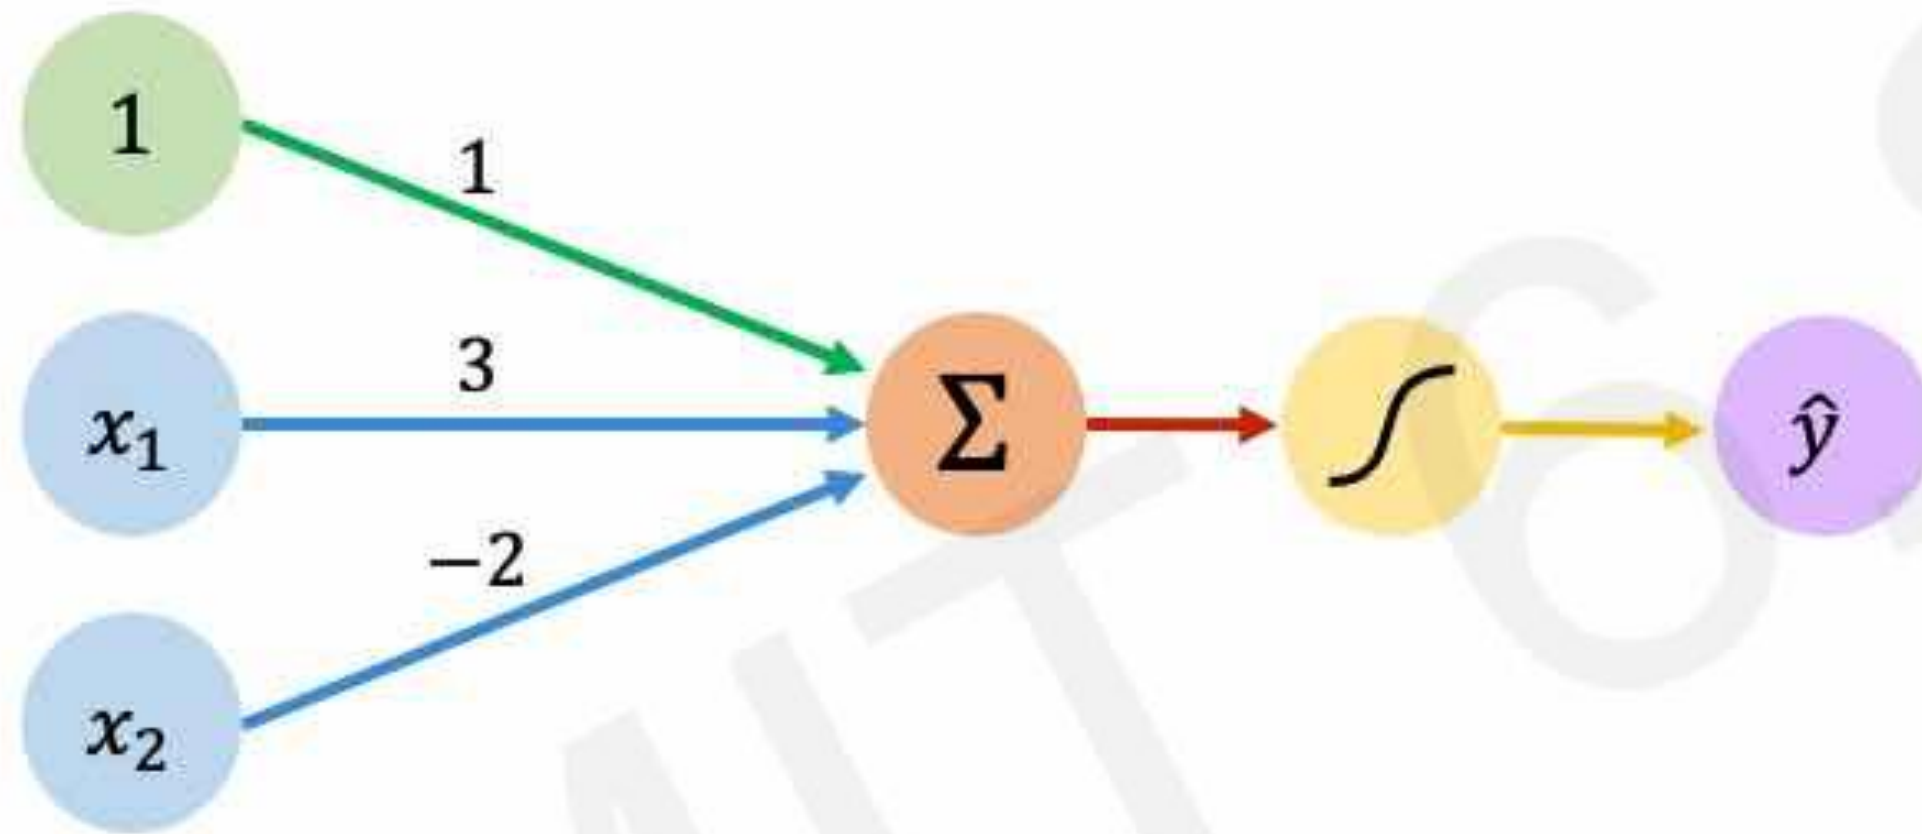
This is just a line in 2D!

# The Perceptron: Example



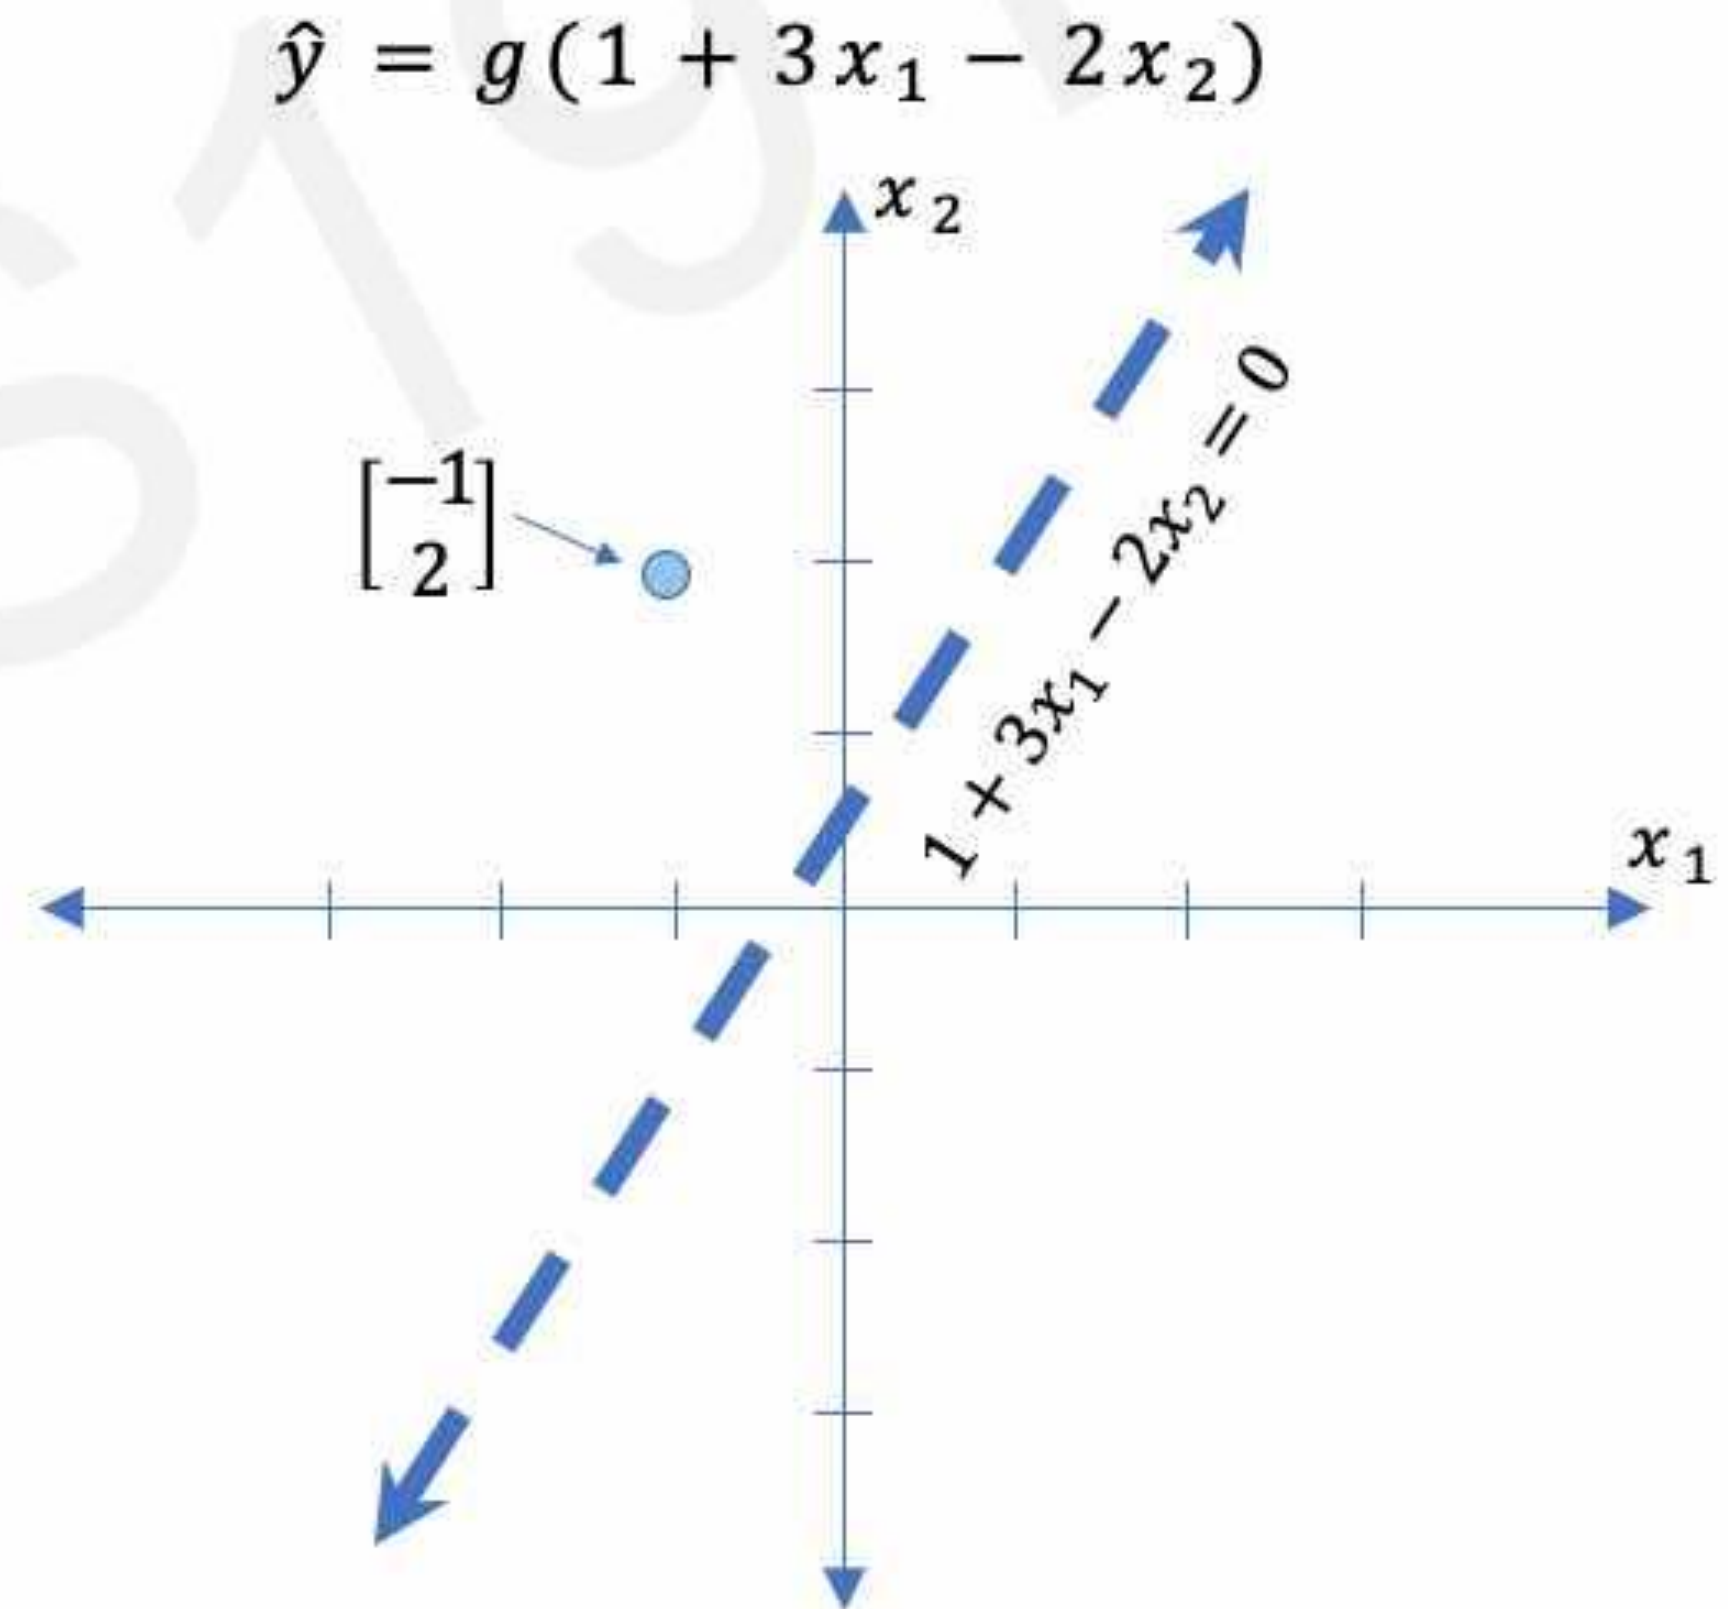


# The Perceptron: Example

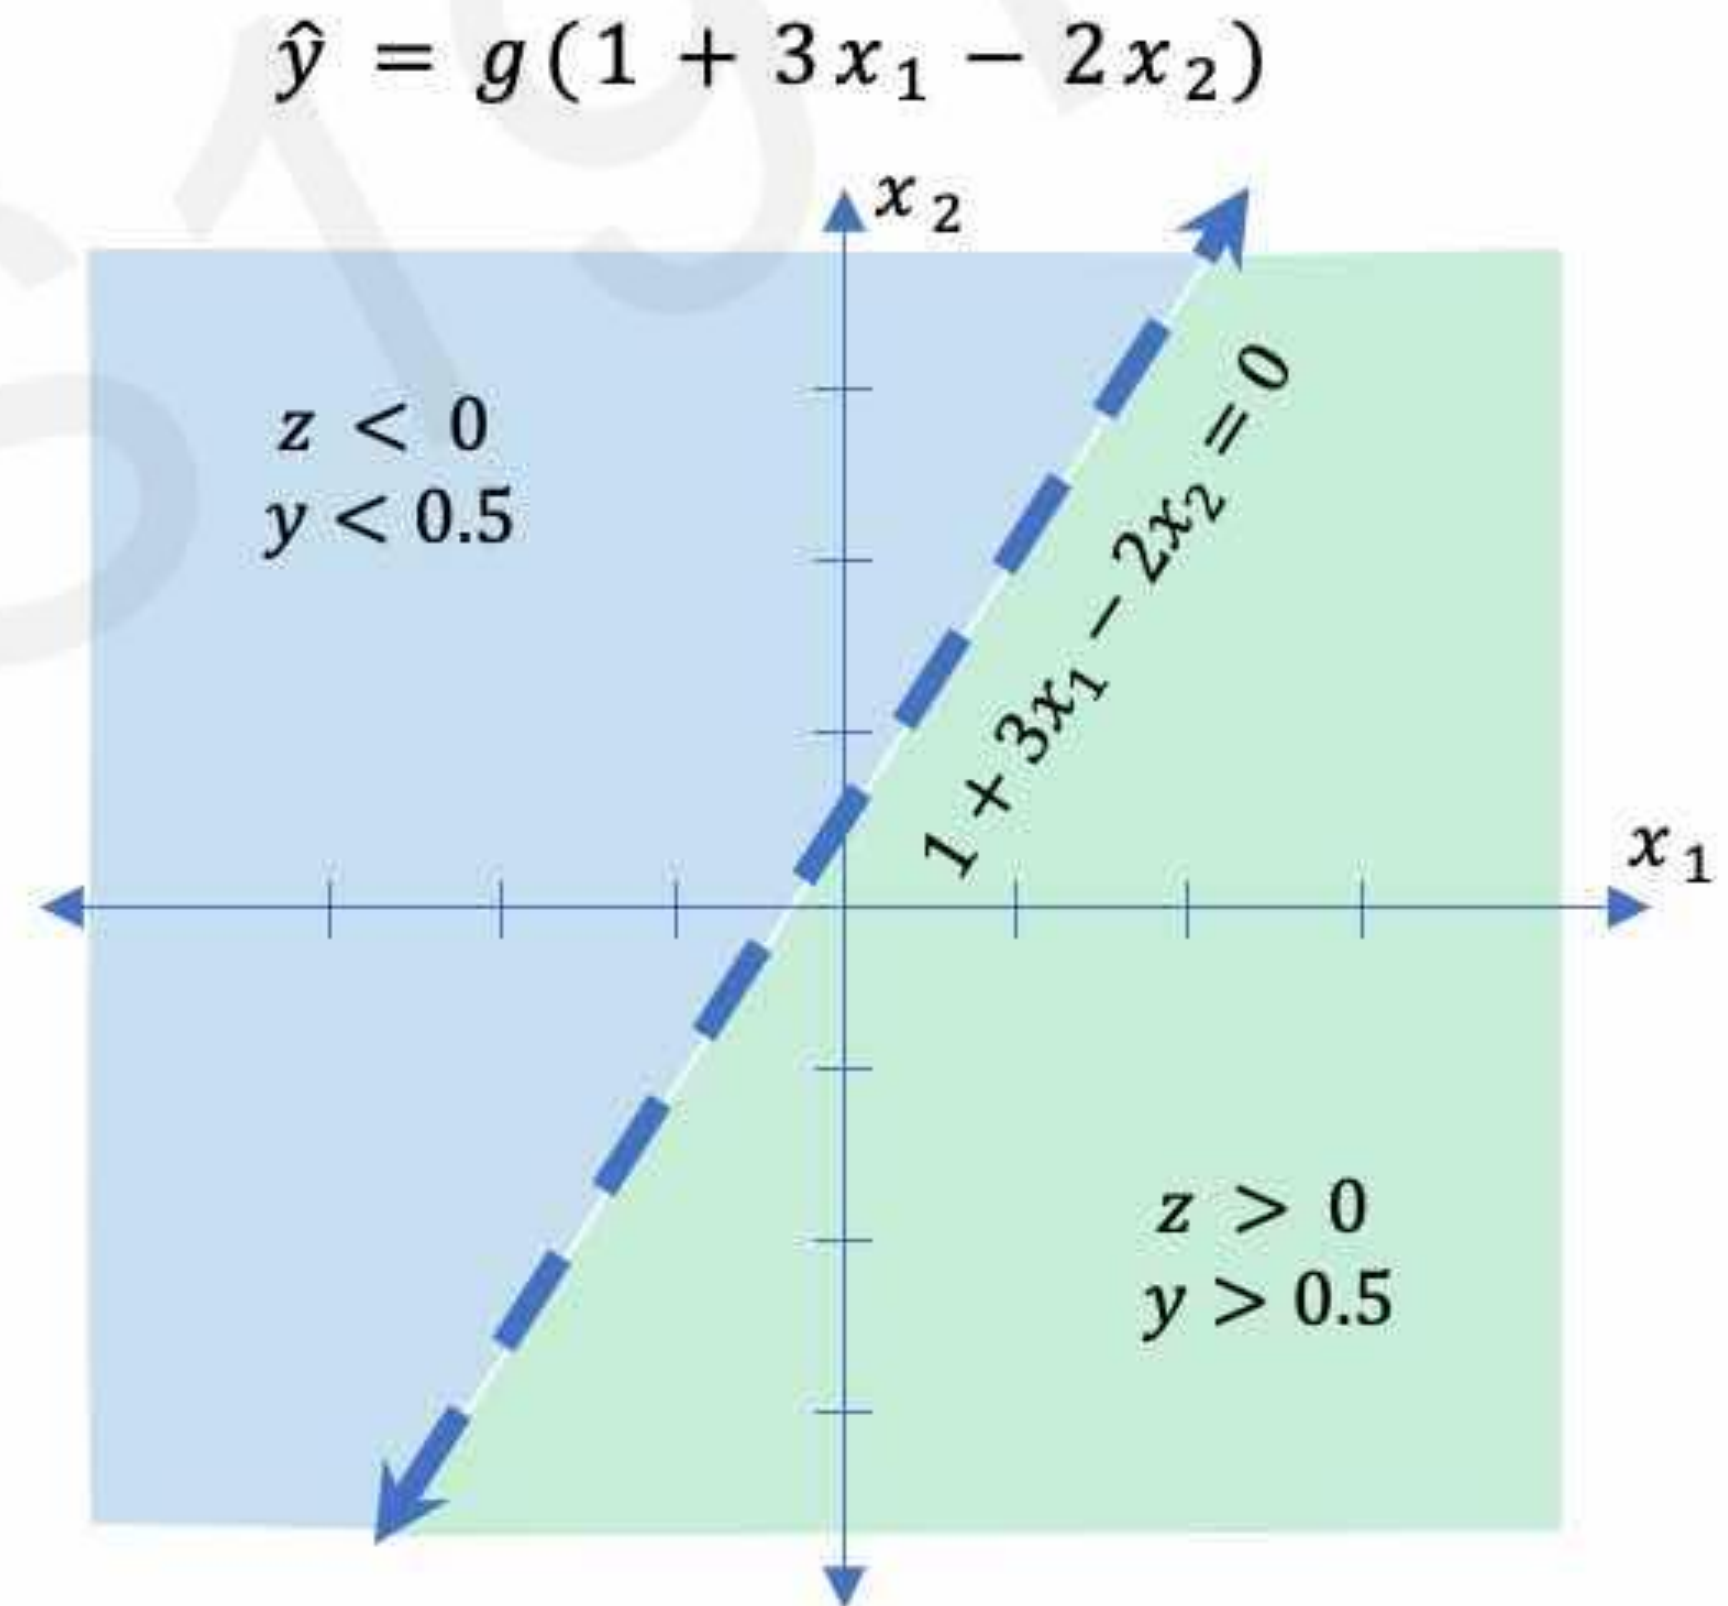
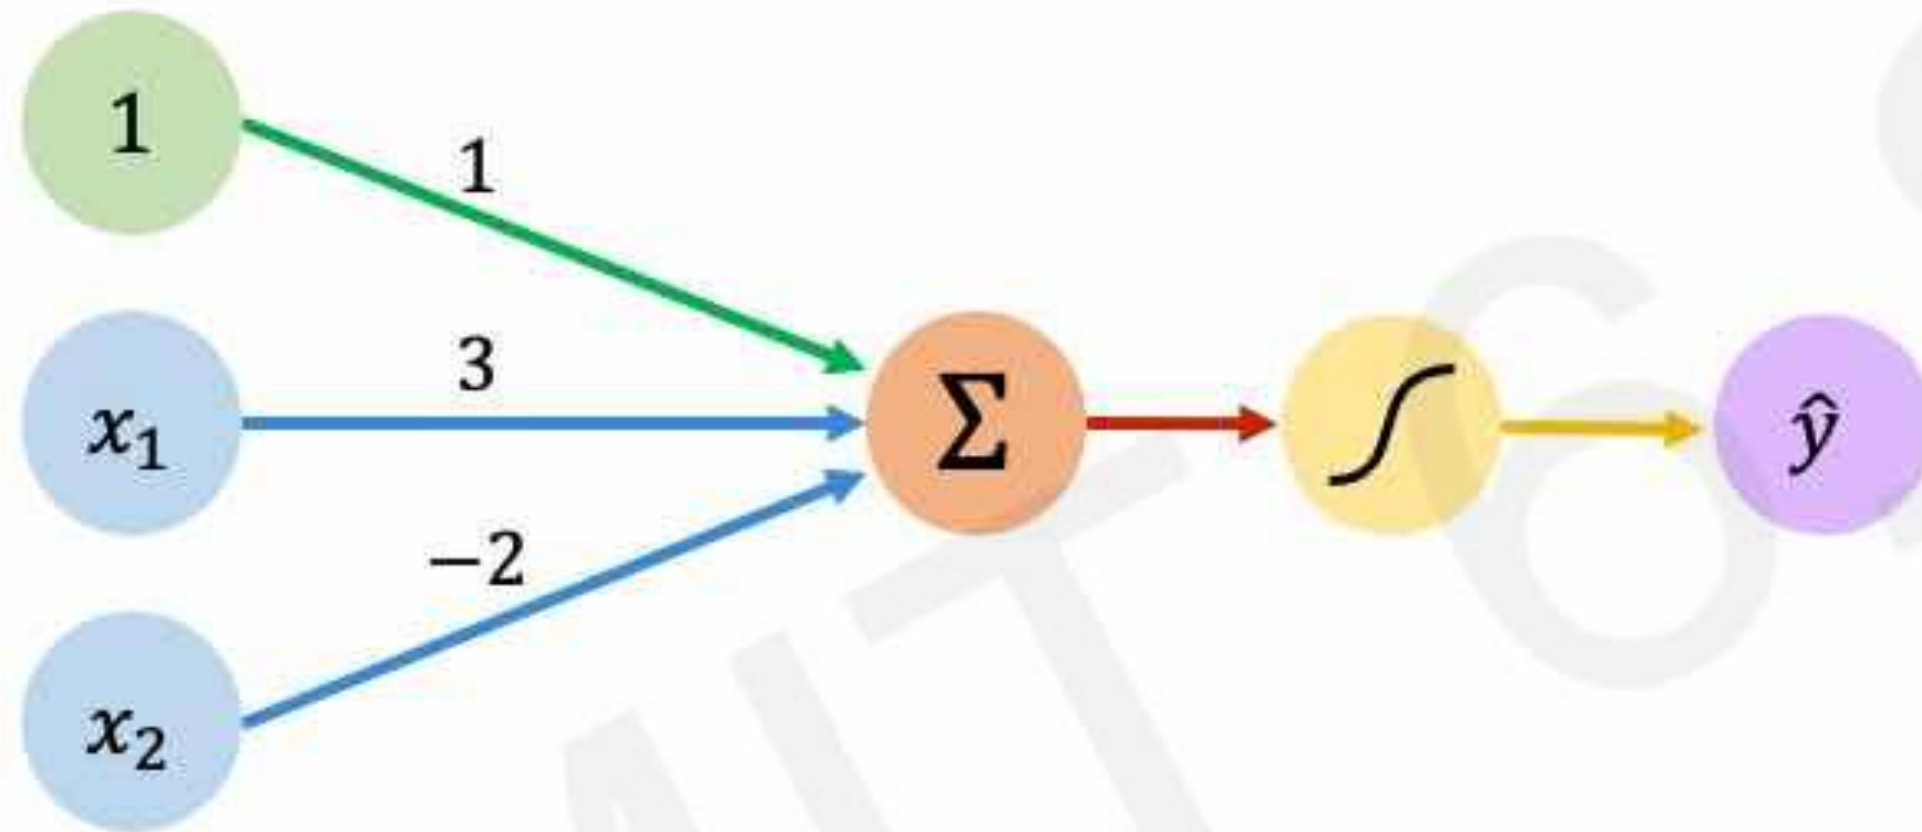


Assume we have input:  $\mathbf{X} = \begin{bmatrix} -1 \\ 2 \end{bmatrix}$

$$\begin{aligned}\hat{y} &= g(1 + (3 * -1) - (2 * 2)) \\ &= g(-6) \approx 0.002\end{aligned}$$

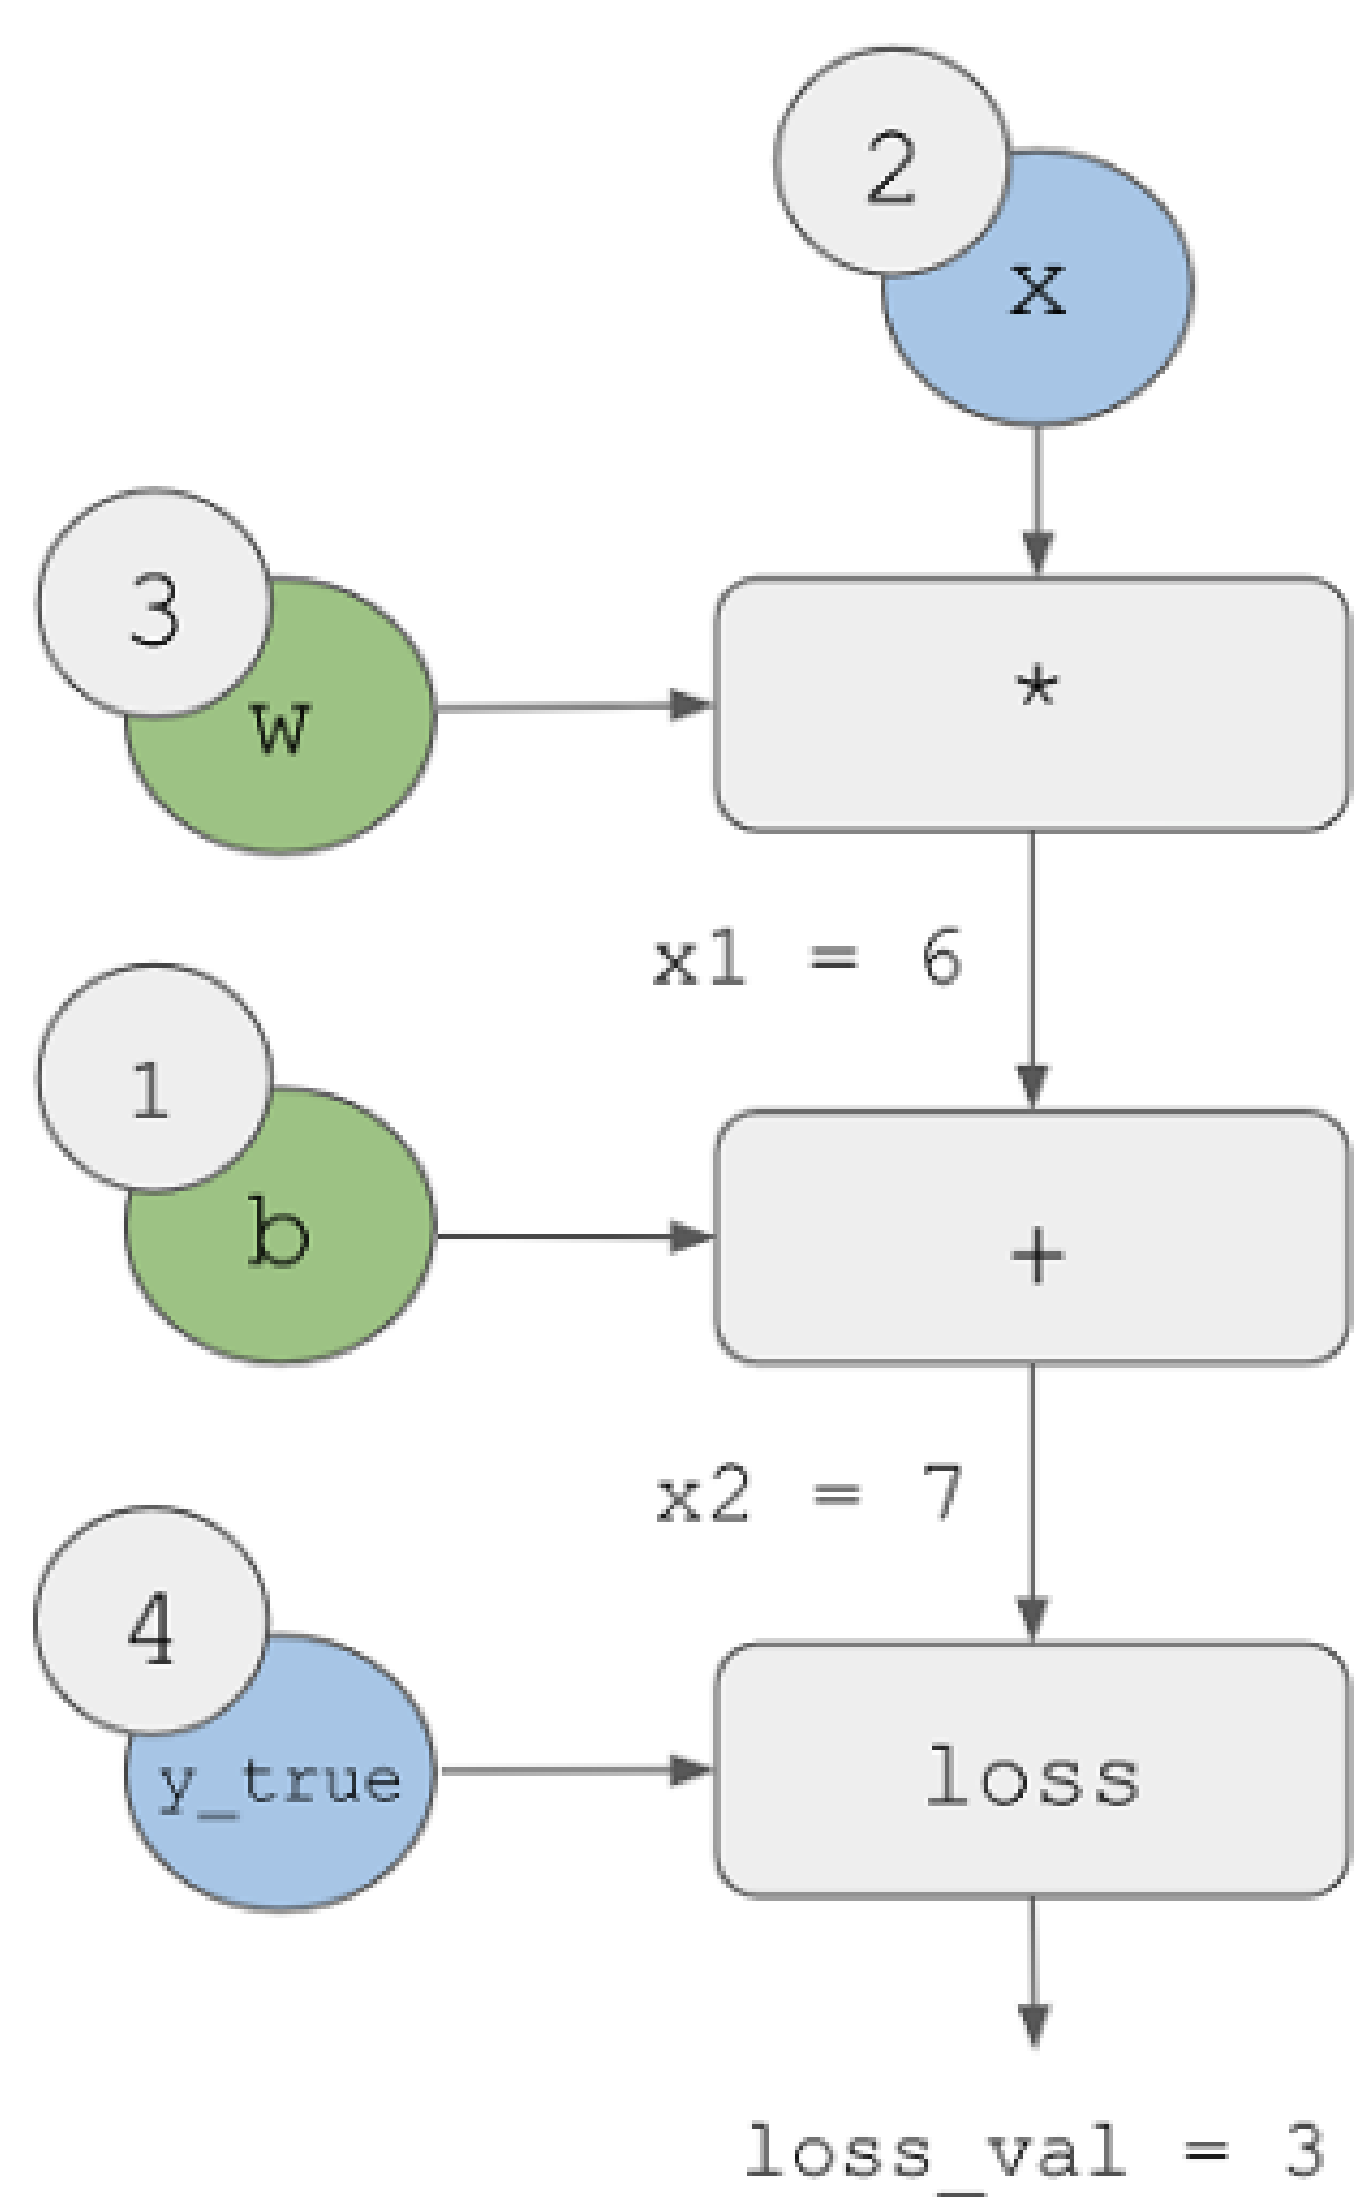
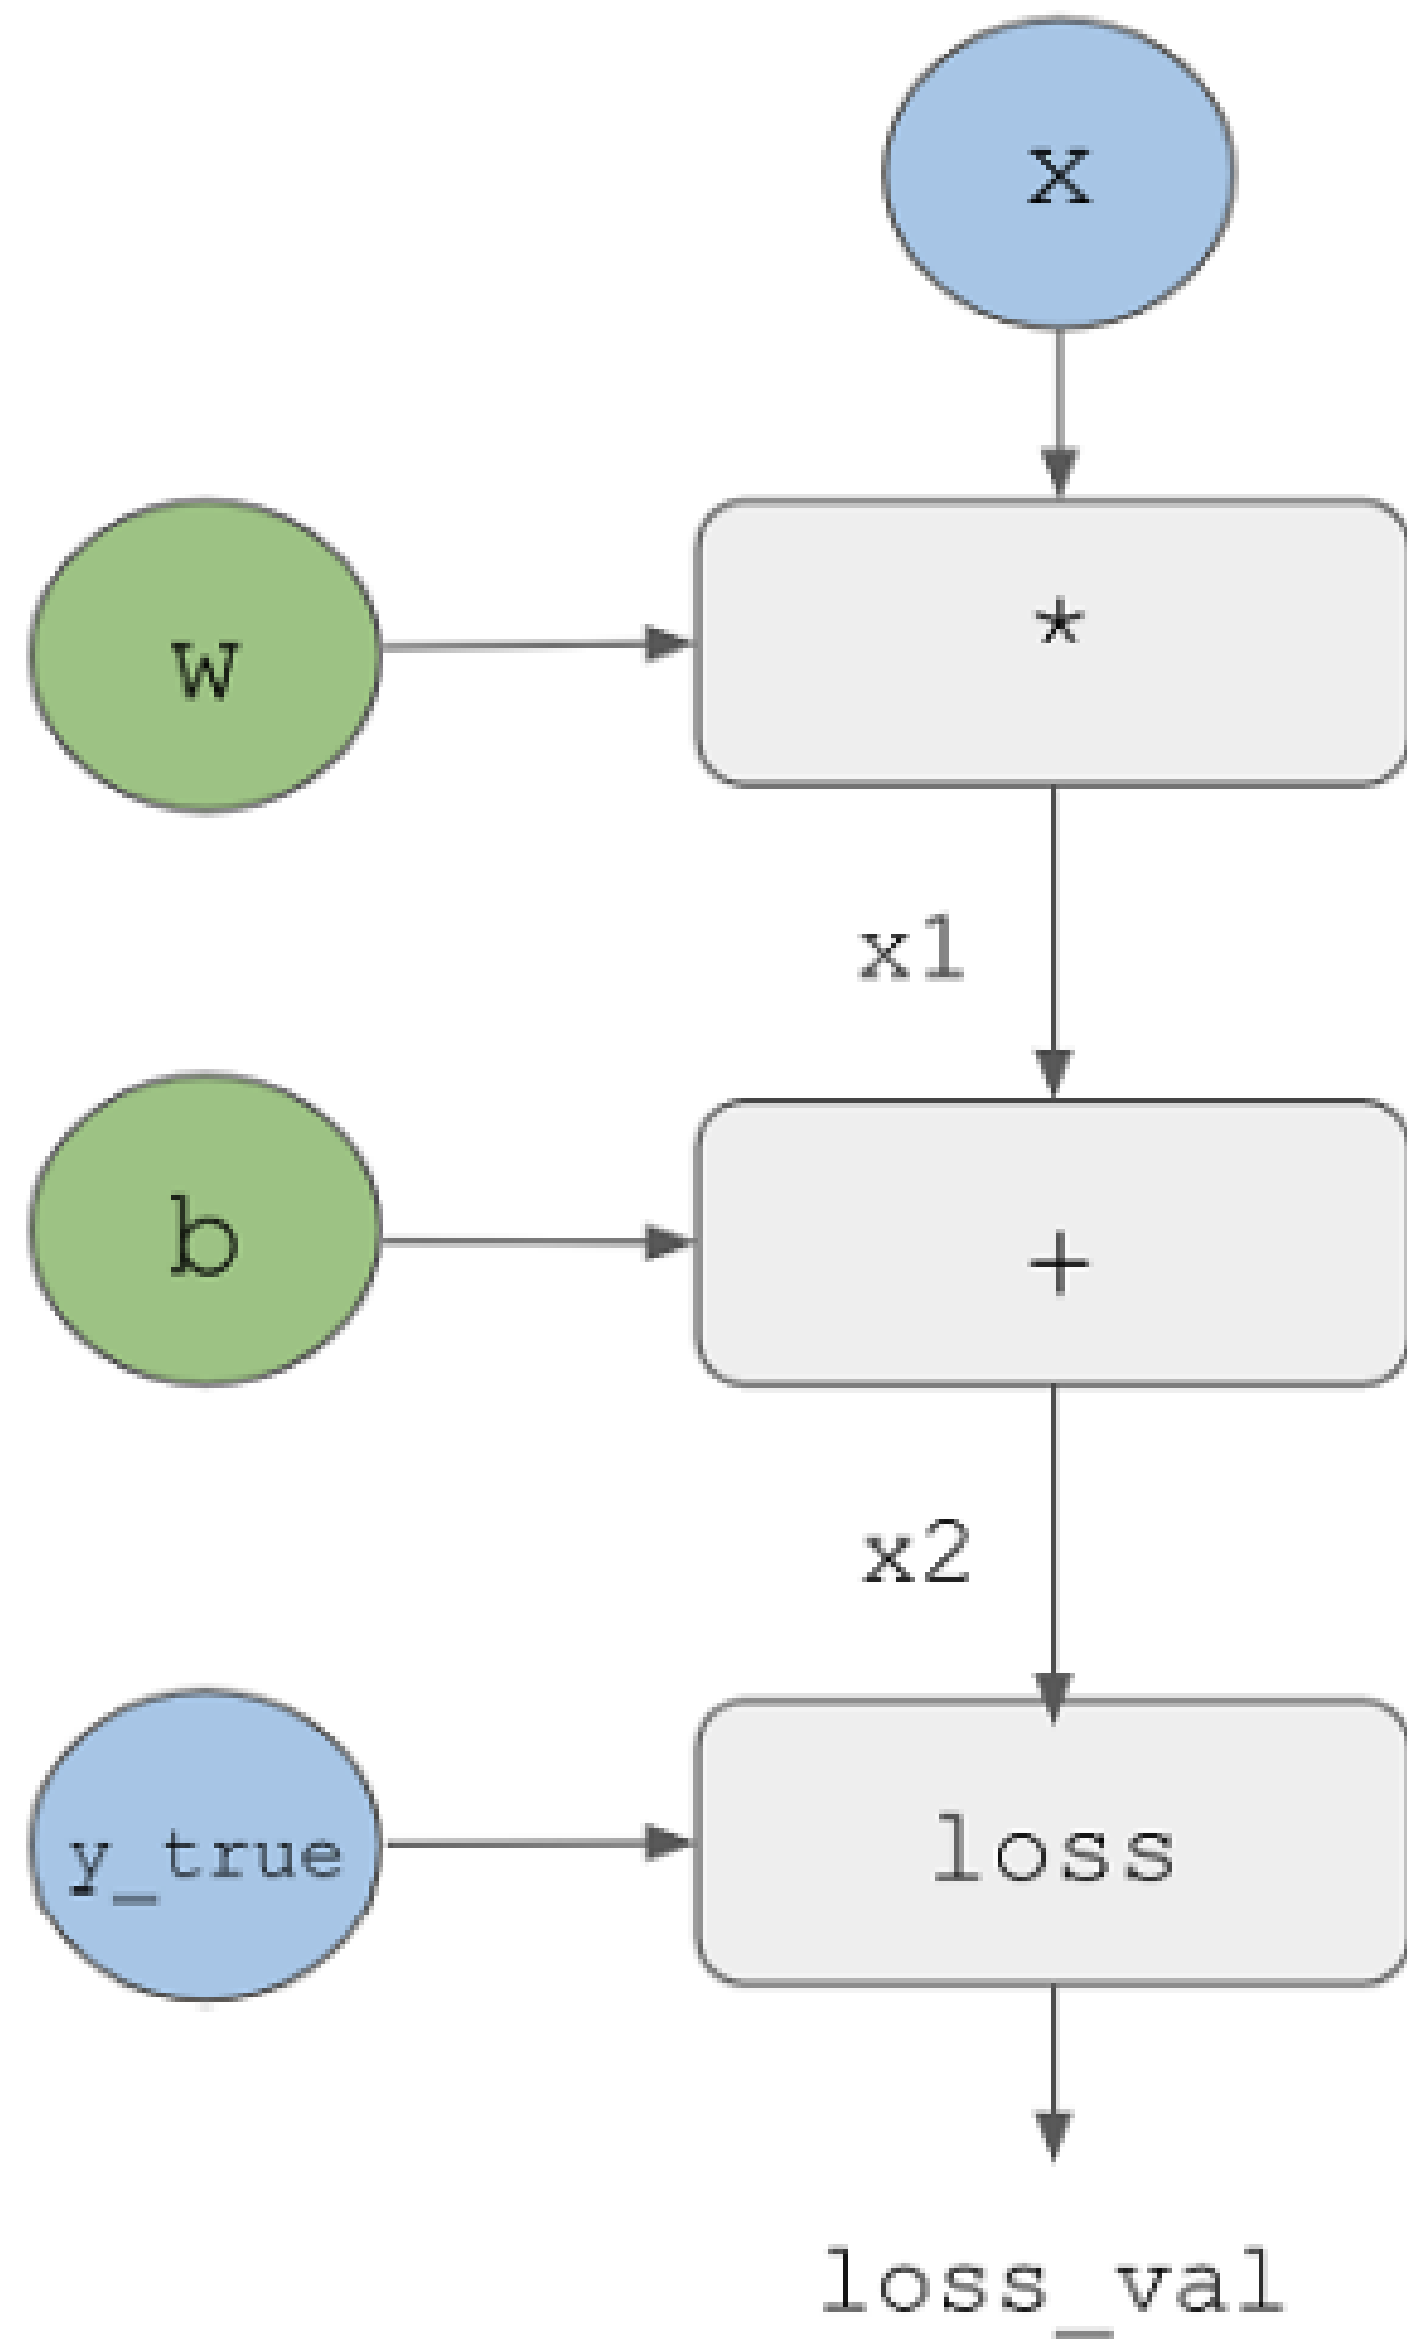


# The Perceptron: Example



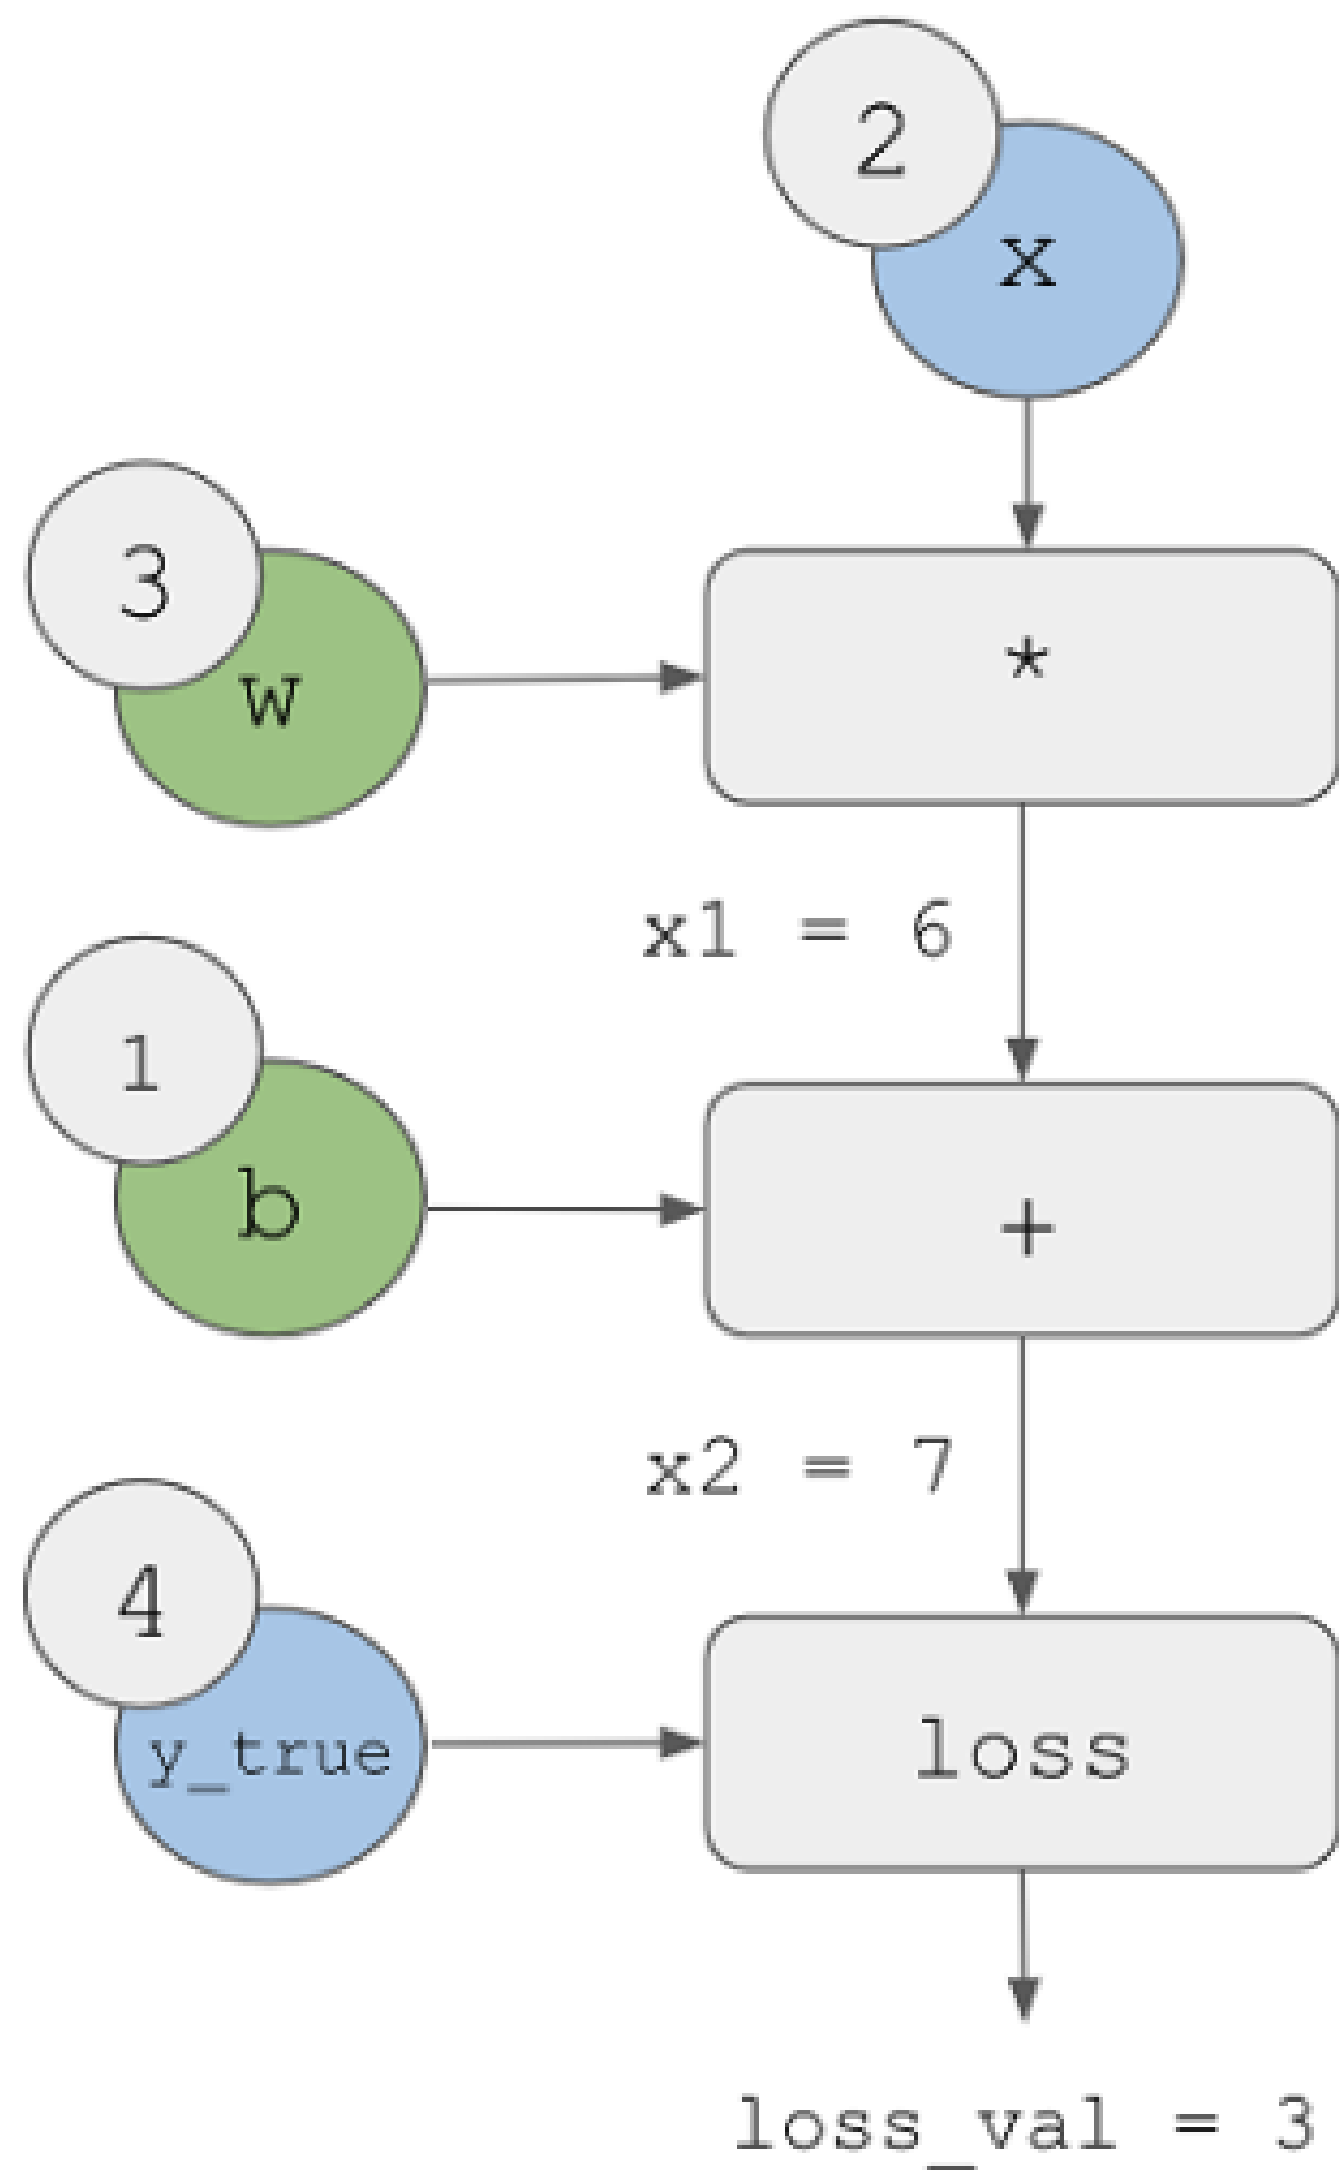


# Building Neural Networks with Perceptrons

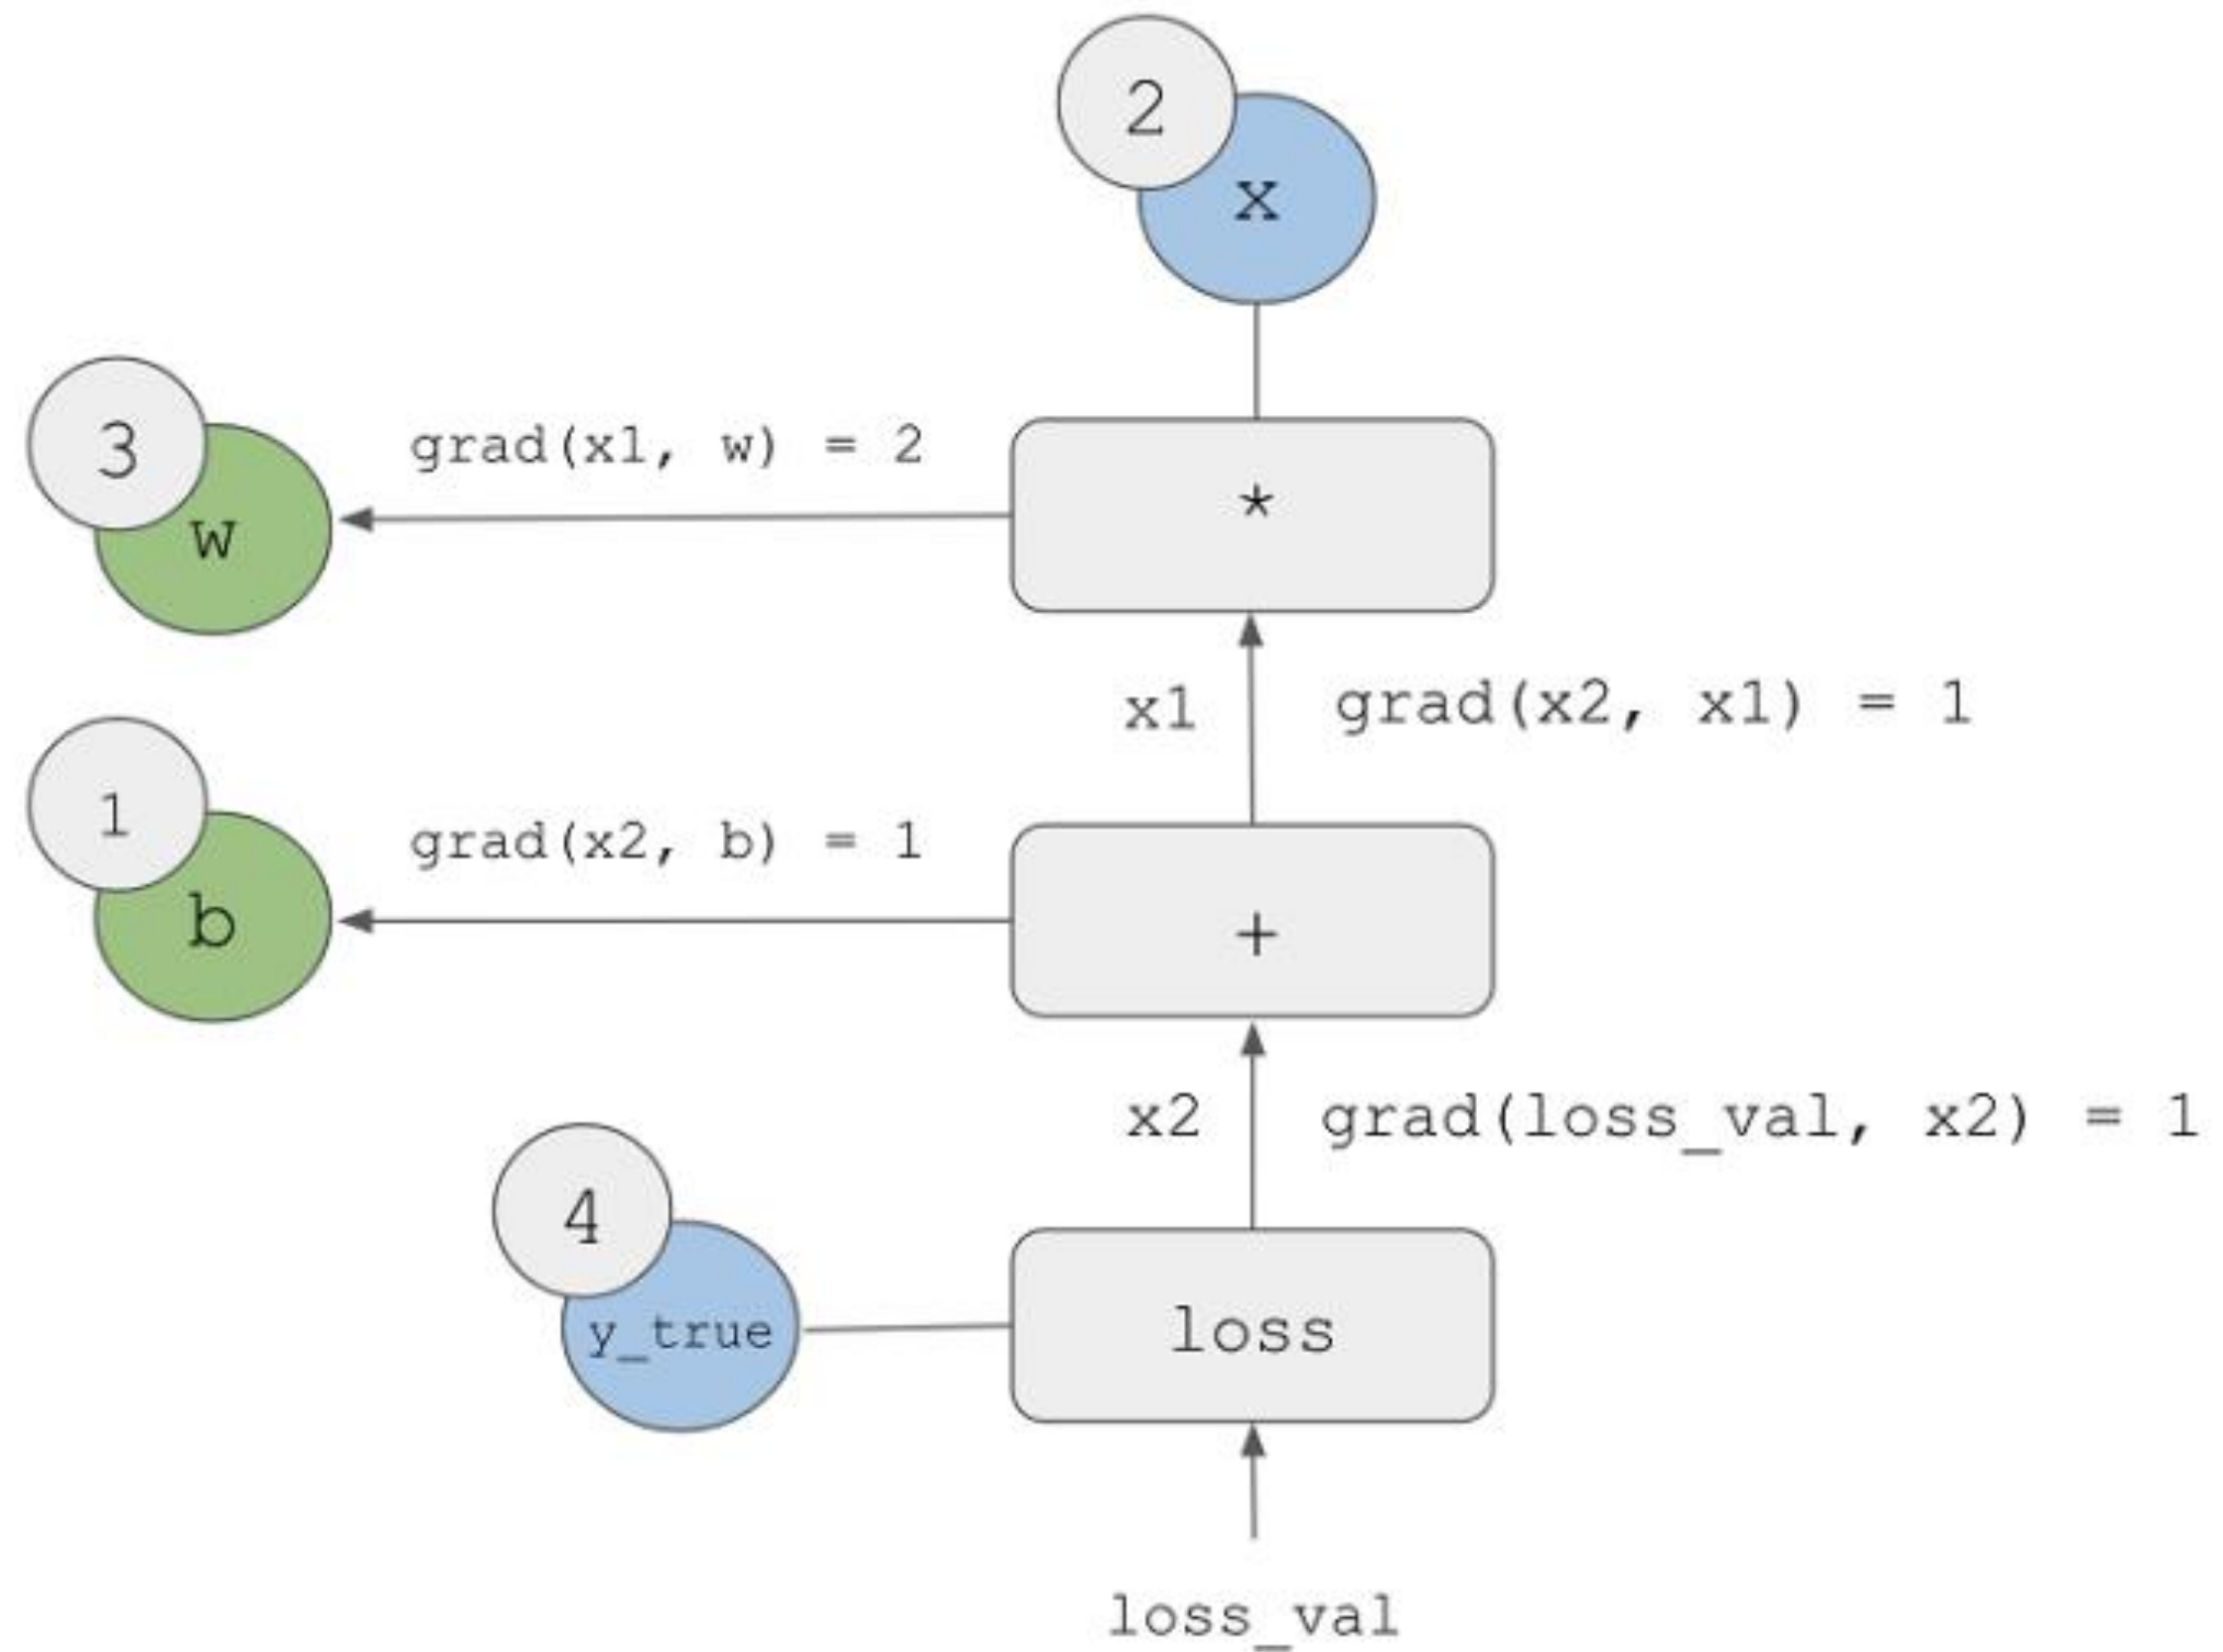


**Figure 2.21 Running a forward pass**





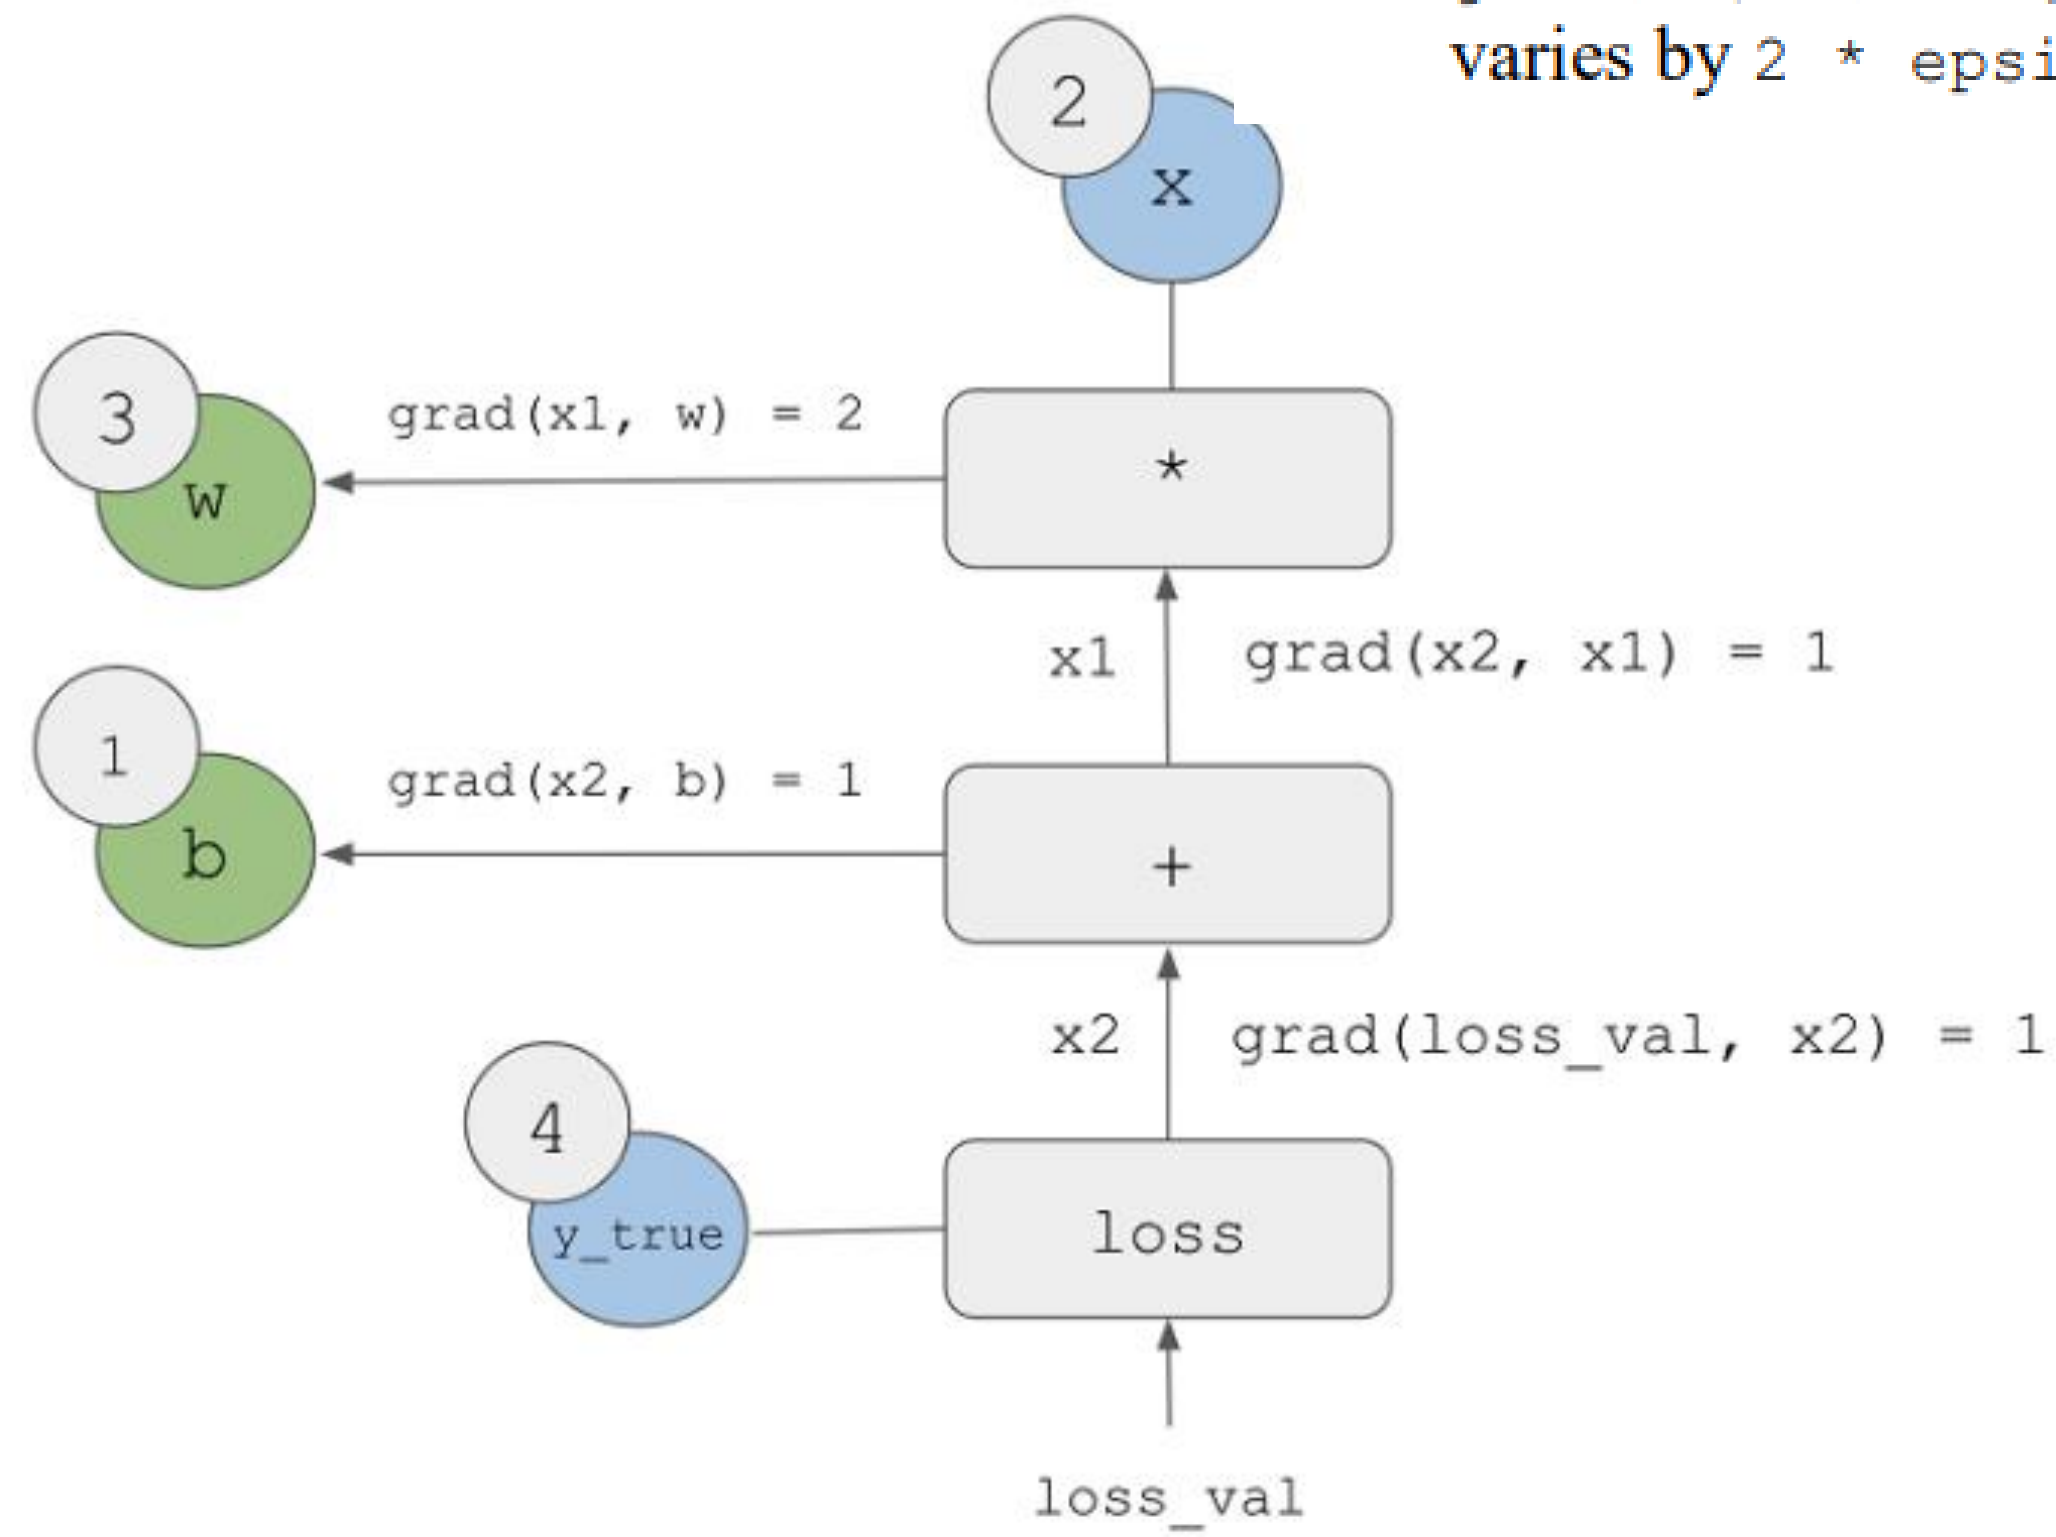
**Figure 2.21 Running a forward pass**



**Figure 2.22 Running a backward pass**

We have:

- $\text{grad}(\text{loss\_val}, x_2) = 1$ , because as  $x_2$  varies by an amount  $\epsilon$ ,  $\text{loss\_val} = \text{abs}(4 - x_2)$  varies by the same amount.
- $\text{grad}(x_2, x_1) = 1$ , because as  $x_1$  varies by an amount  $\epsilon$ ,  $x_2 = x_1 + b = x_1 + 1$  varies by the same amount.
- $\text{grad}(x_2, b) = 1$ , because as  $b$  varies by an amount  $\epsilon$ ,  $x_2 = x_1 + b = 6 + b$  varies by the same amount.
- $\text{grad}(x_1, w) = 2$ , because as  $w$  varies by an amount  $\epsilon$ ,  $x_1 = x * w = 2 * w$  varies by  $2 * \epsilon$ .

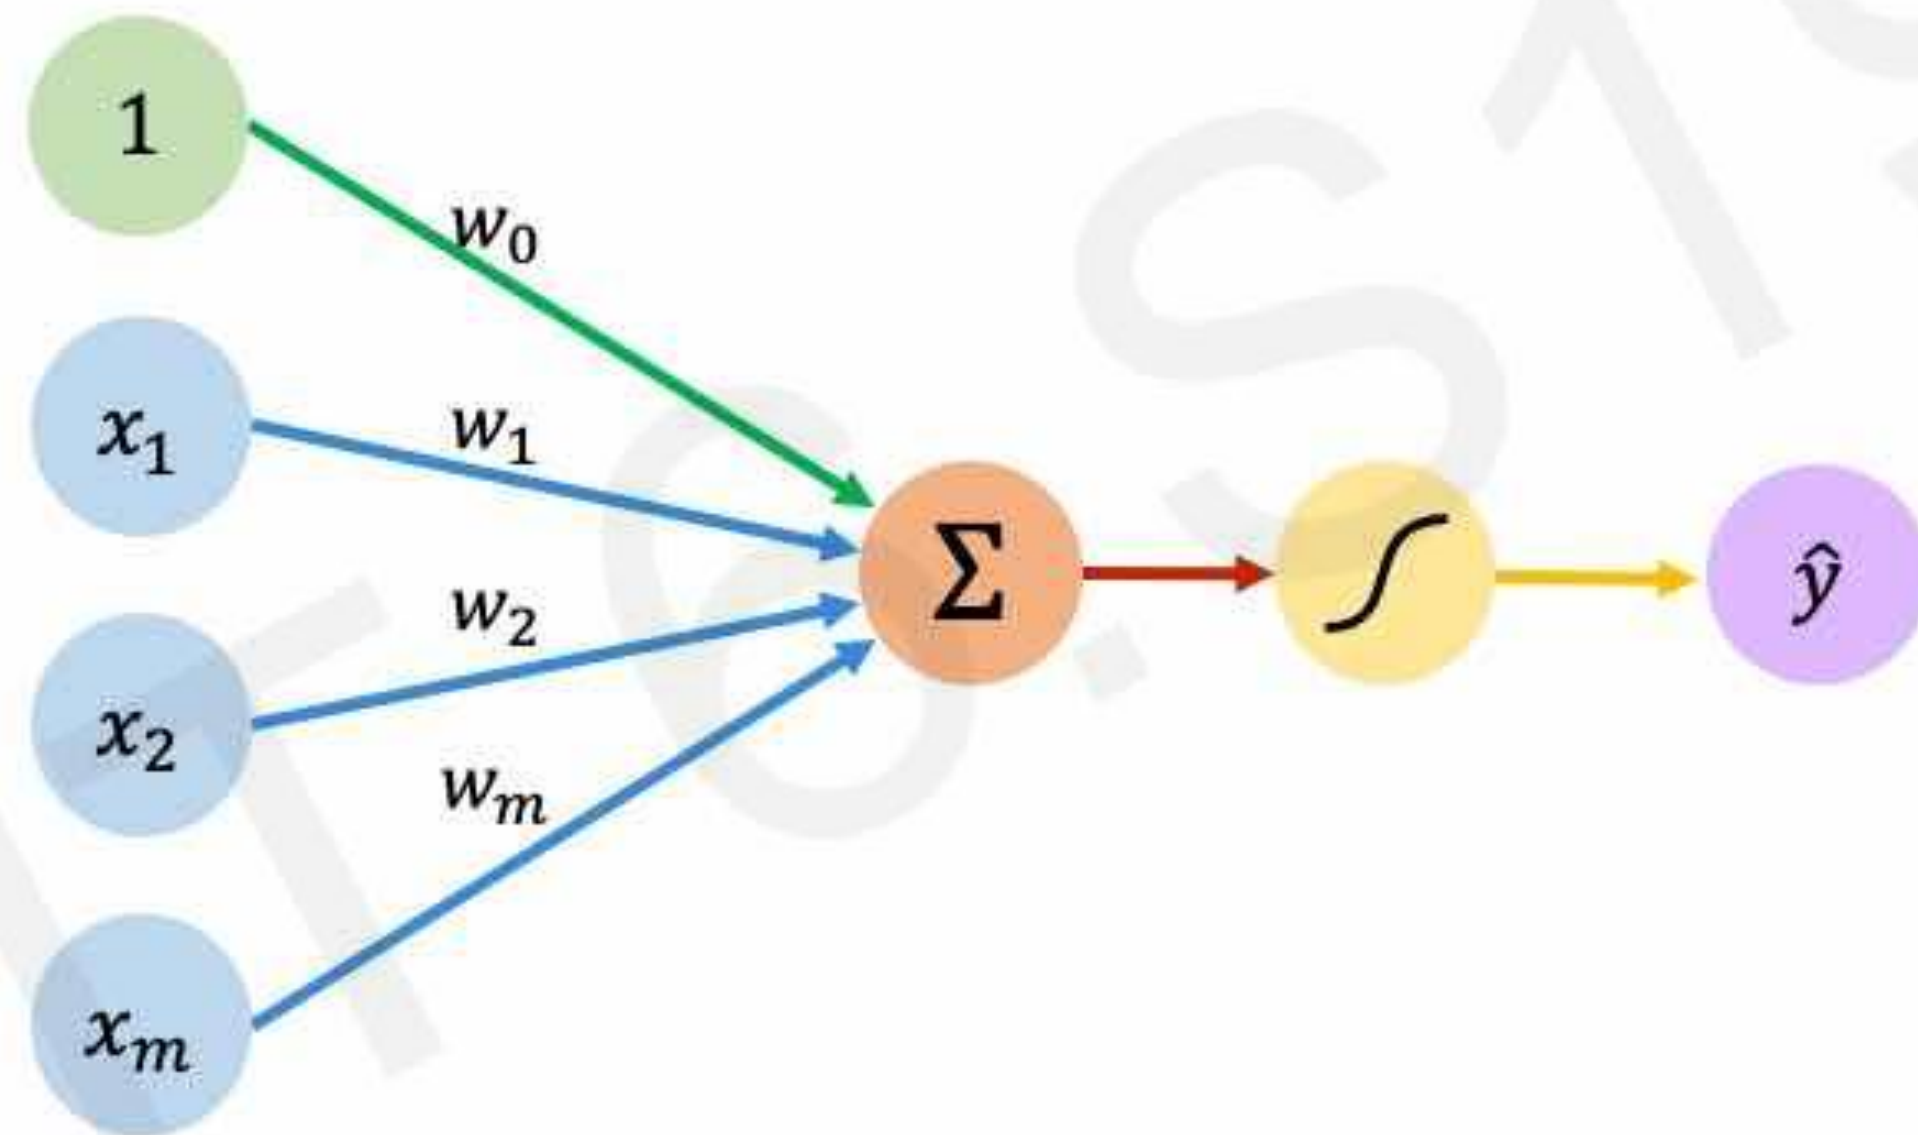


Backward-Pass  
a.k.a. Backpropagation



# The Perceptron: Simplified

$$\hat{y} = g(w_0 + \mathbf{X}^T \mathbf{W})$$



Inputs

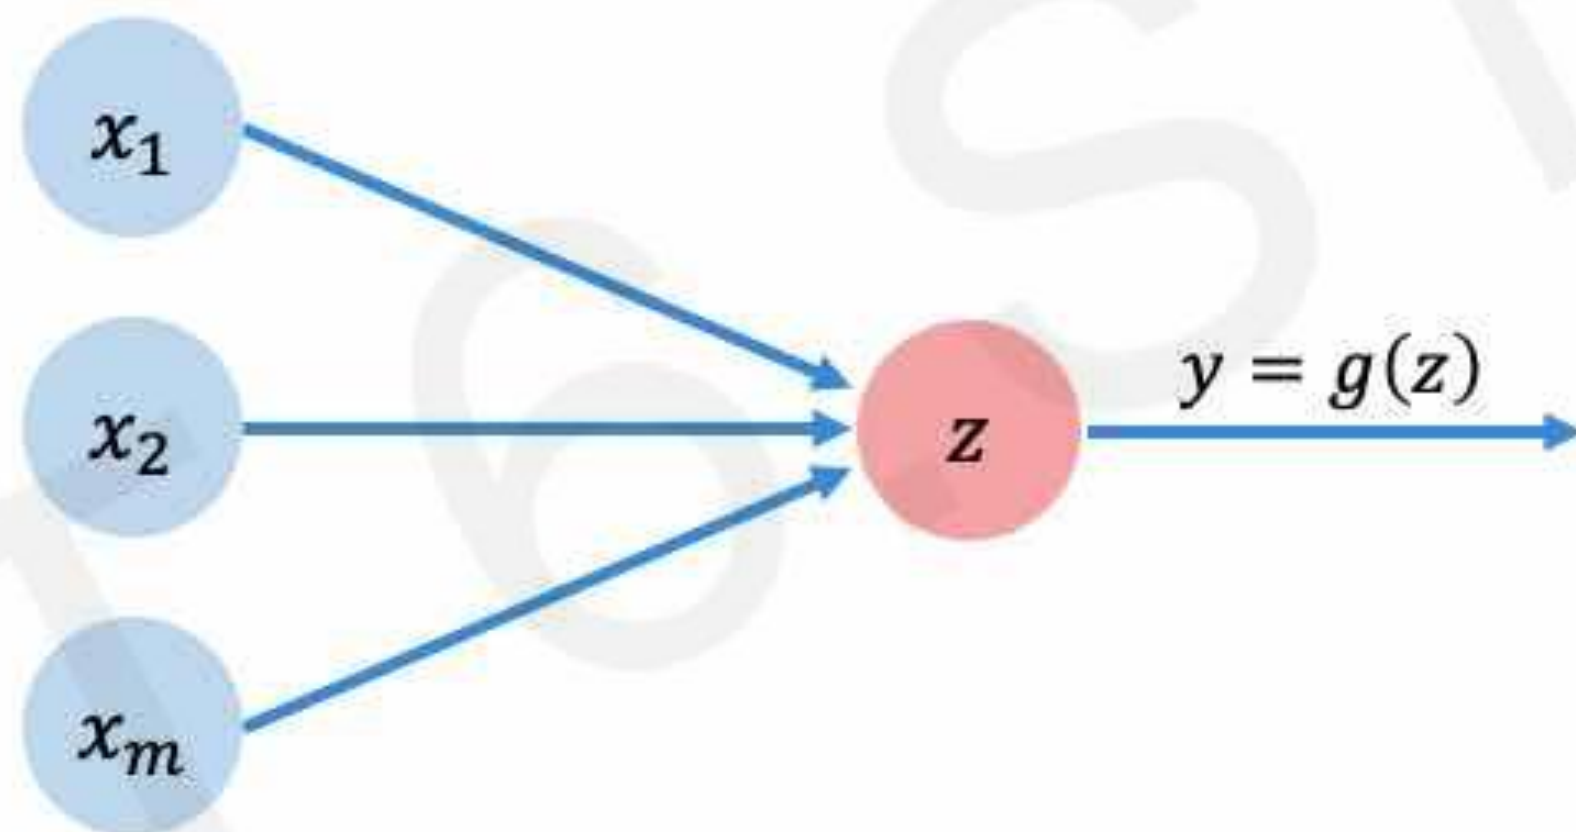
Weights

Sum

Non-Linearity

Output

# The Perceptron: Simplified

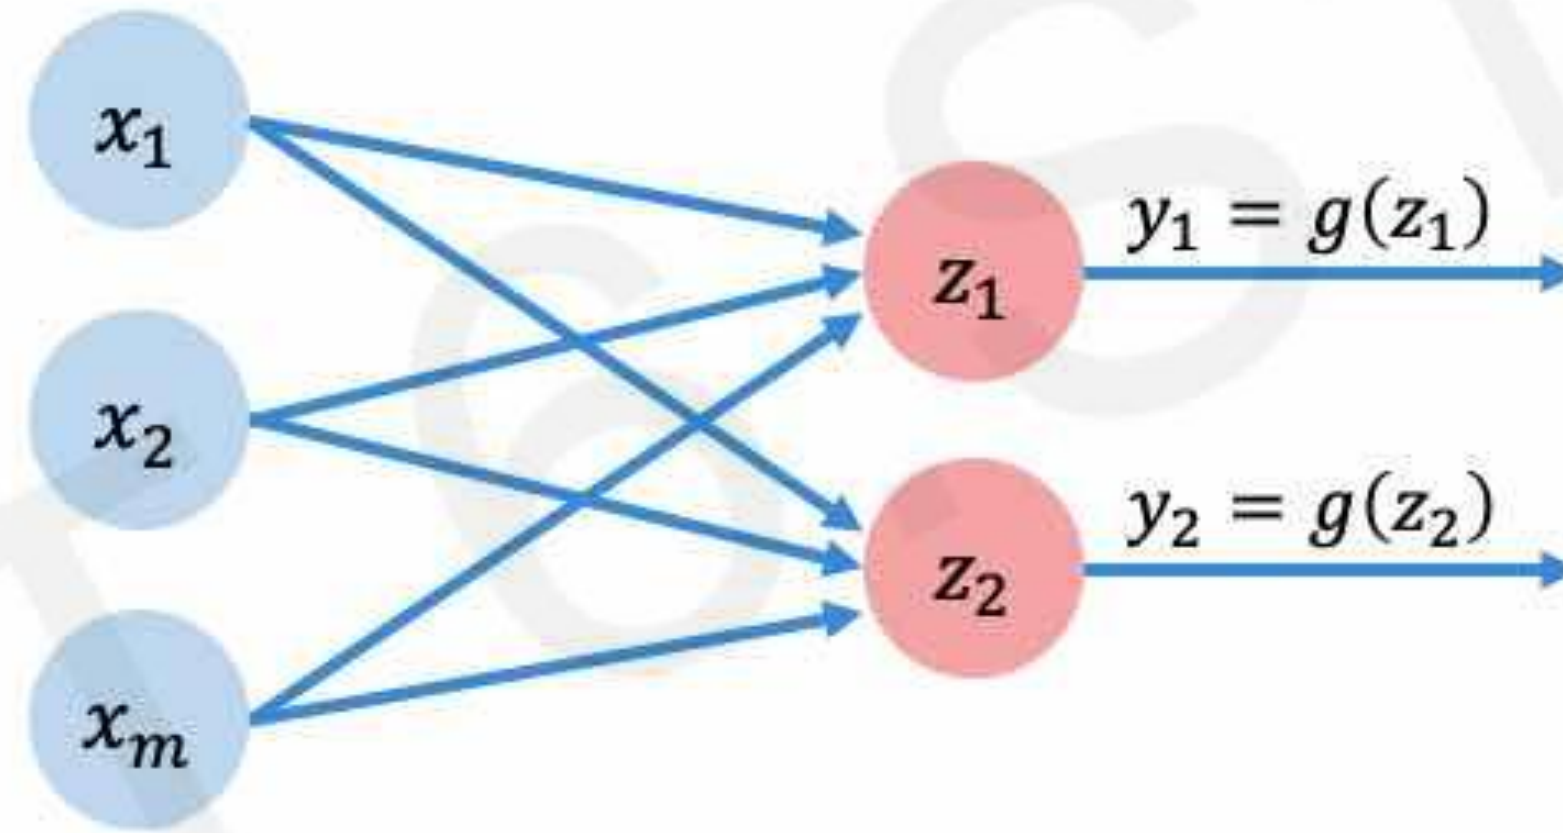


$$z = w_0 + \sum_{j=1}^m x_j w_j$$



# Multi Output Perceptron

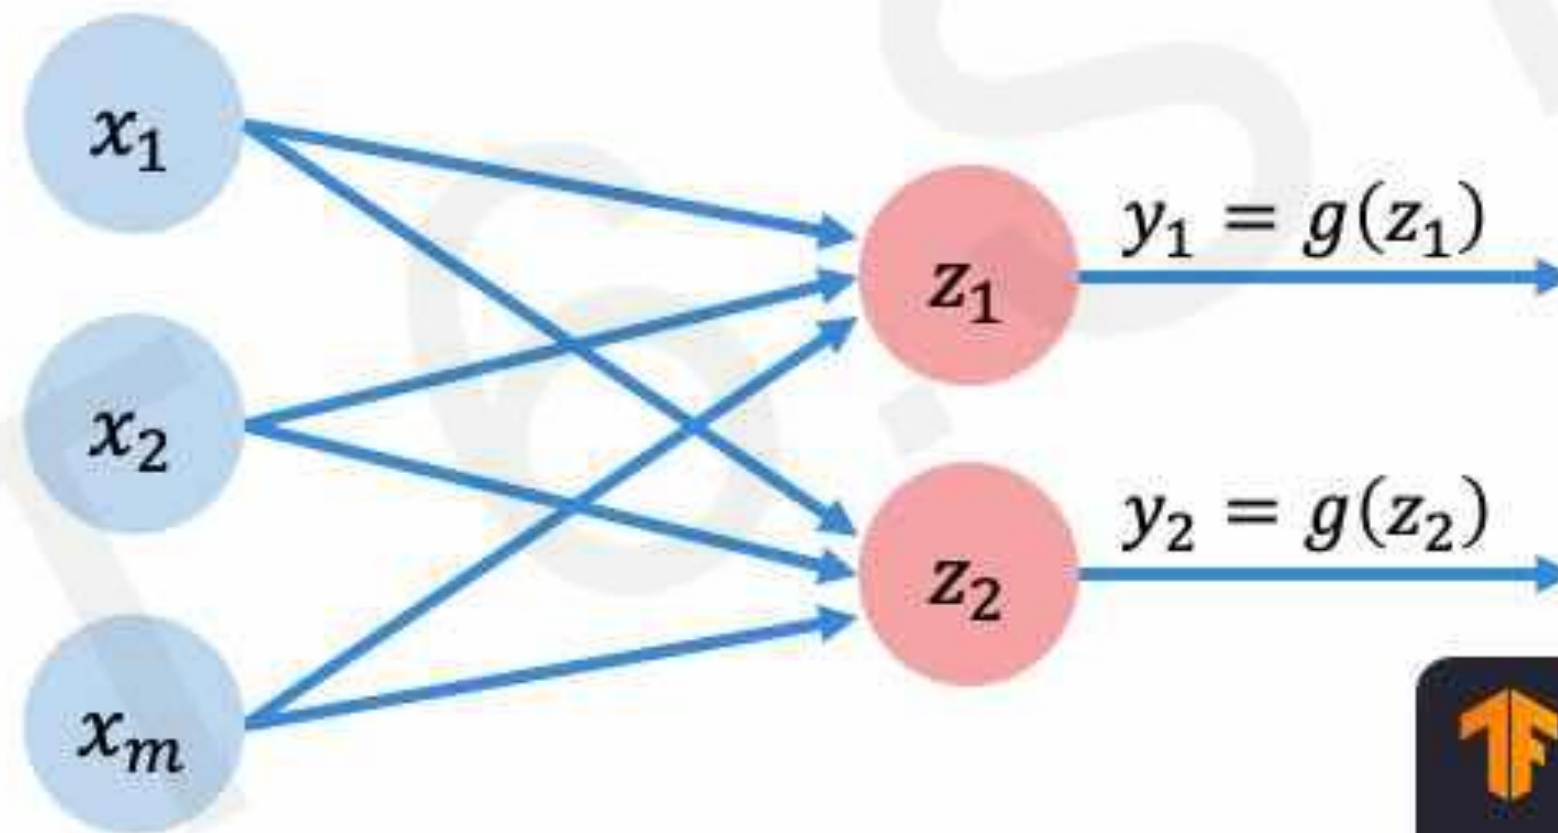
Because all inputs are densely connected to all outputs, these layers are called **Dense** layers




$$z_i = w_{0,i} + \sum_{j=1}^m x_j w_{j,i}$$

# Multi Output Perceptron

Because all inputs are densely connected to all outputs, these layers are called **Dense** layers

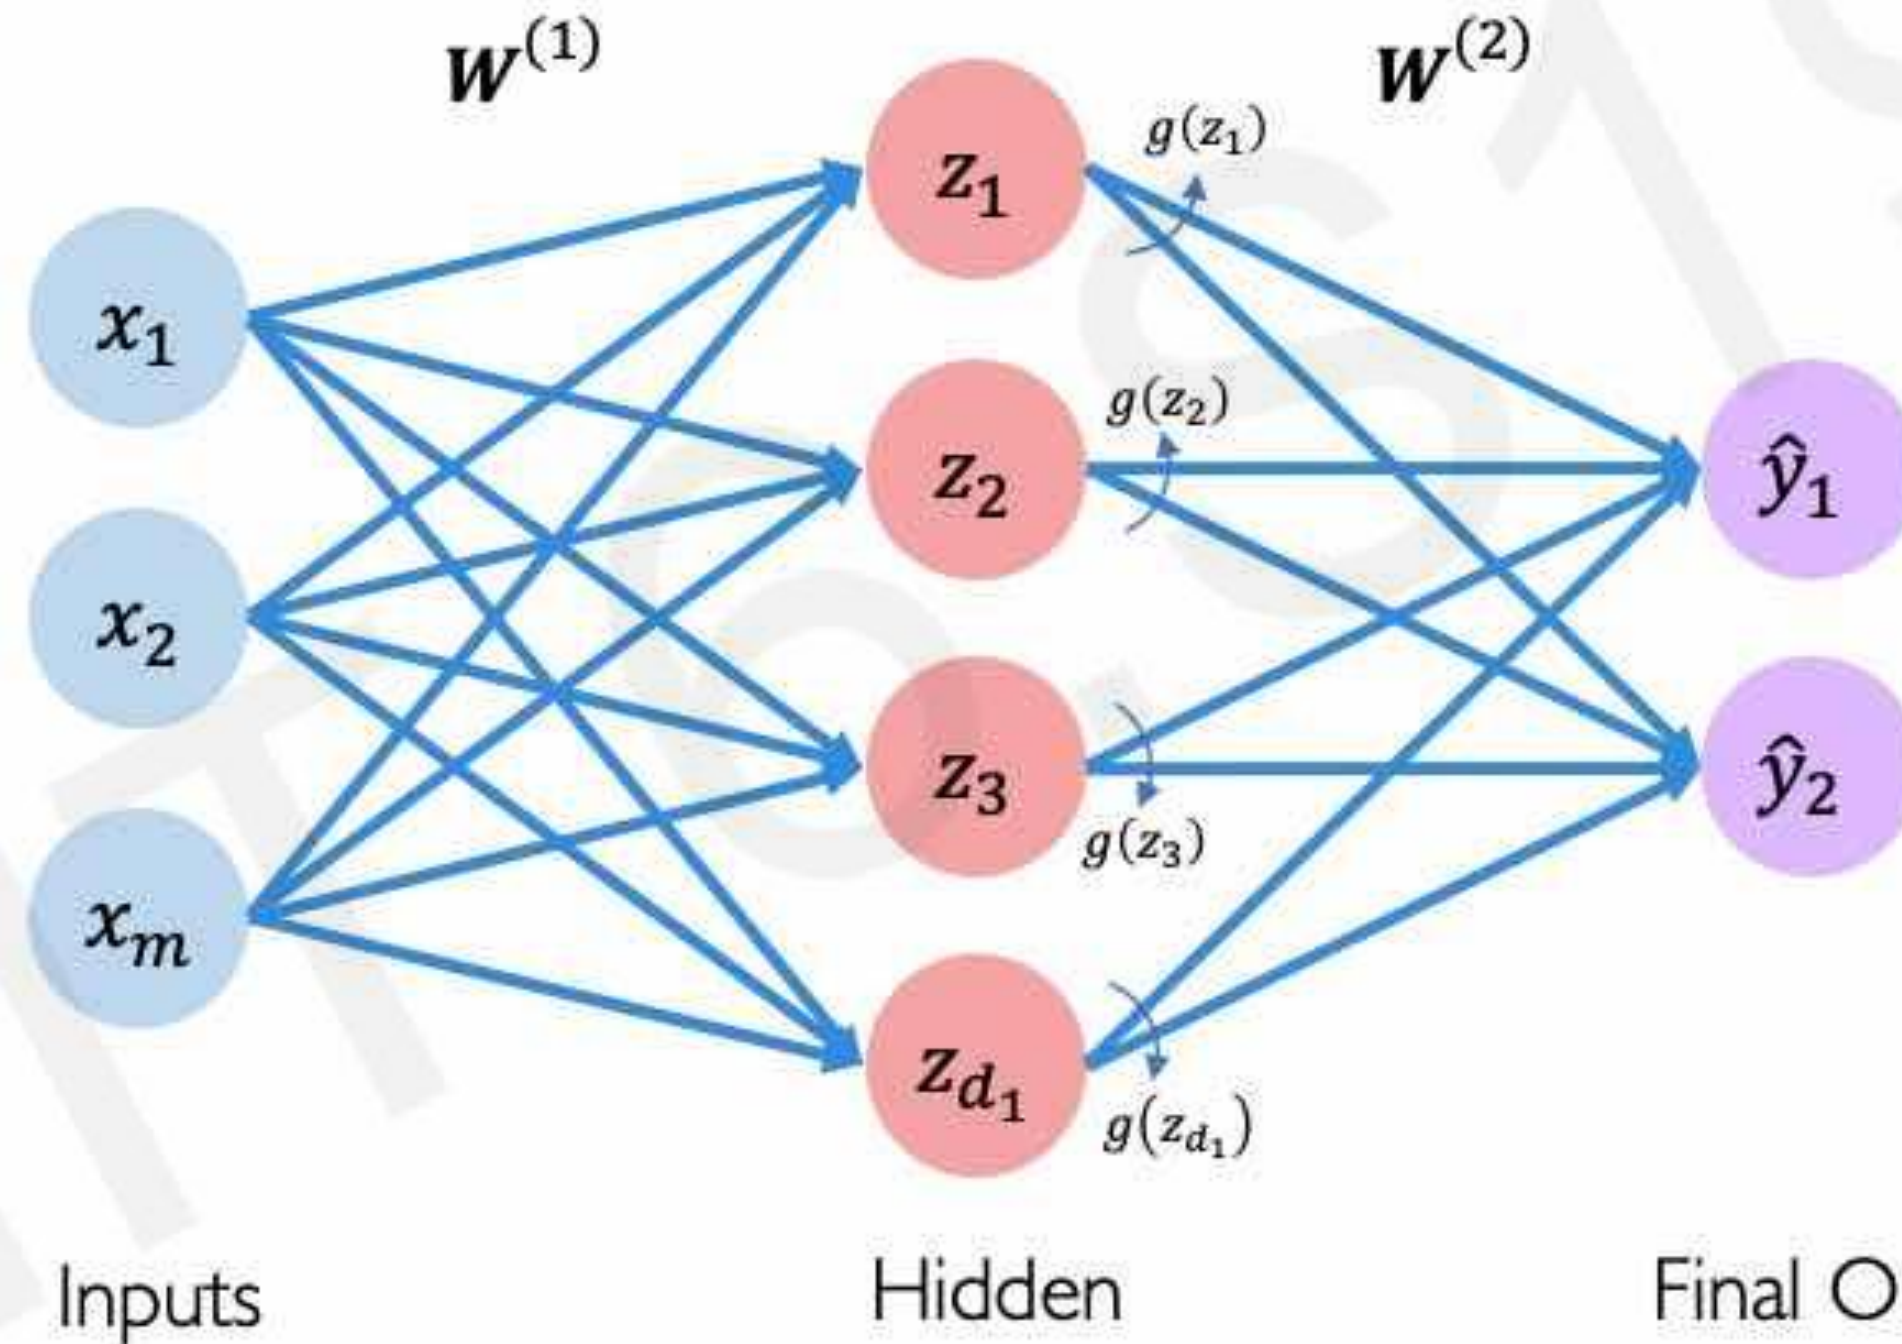


```
 import tensorflow as tf  
  
layer = tf.keras.layers.Dense(  
    units=2)
```

$$z_i = w_{0,i} + \sum_{j=1}^m x_j w_{j,i}$$



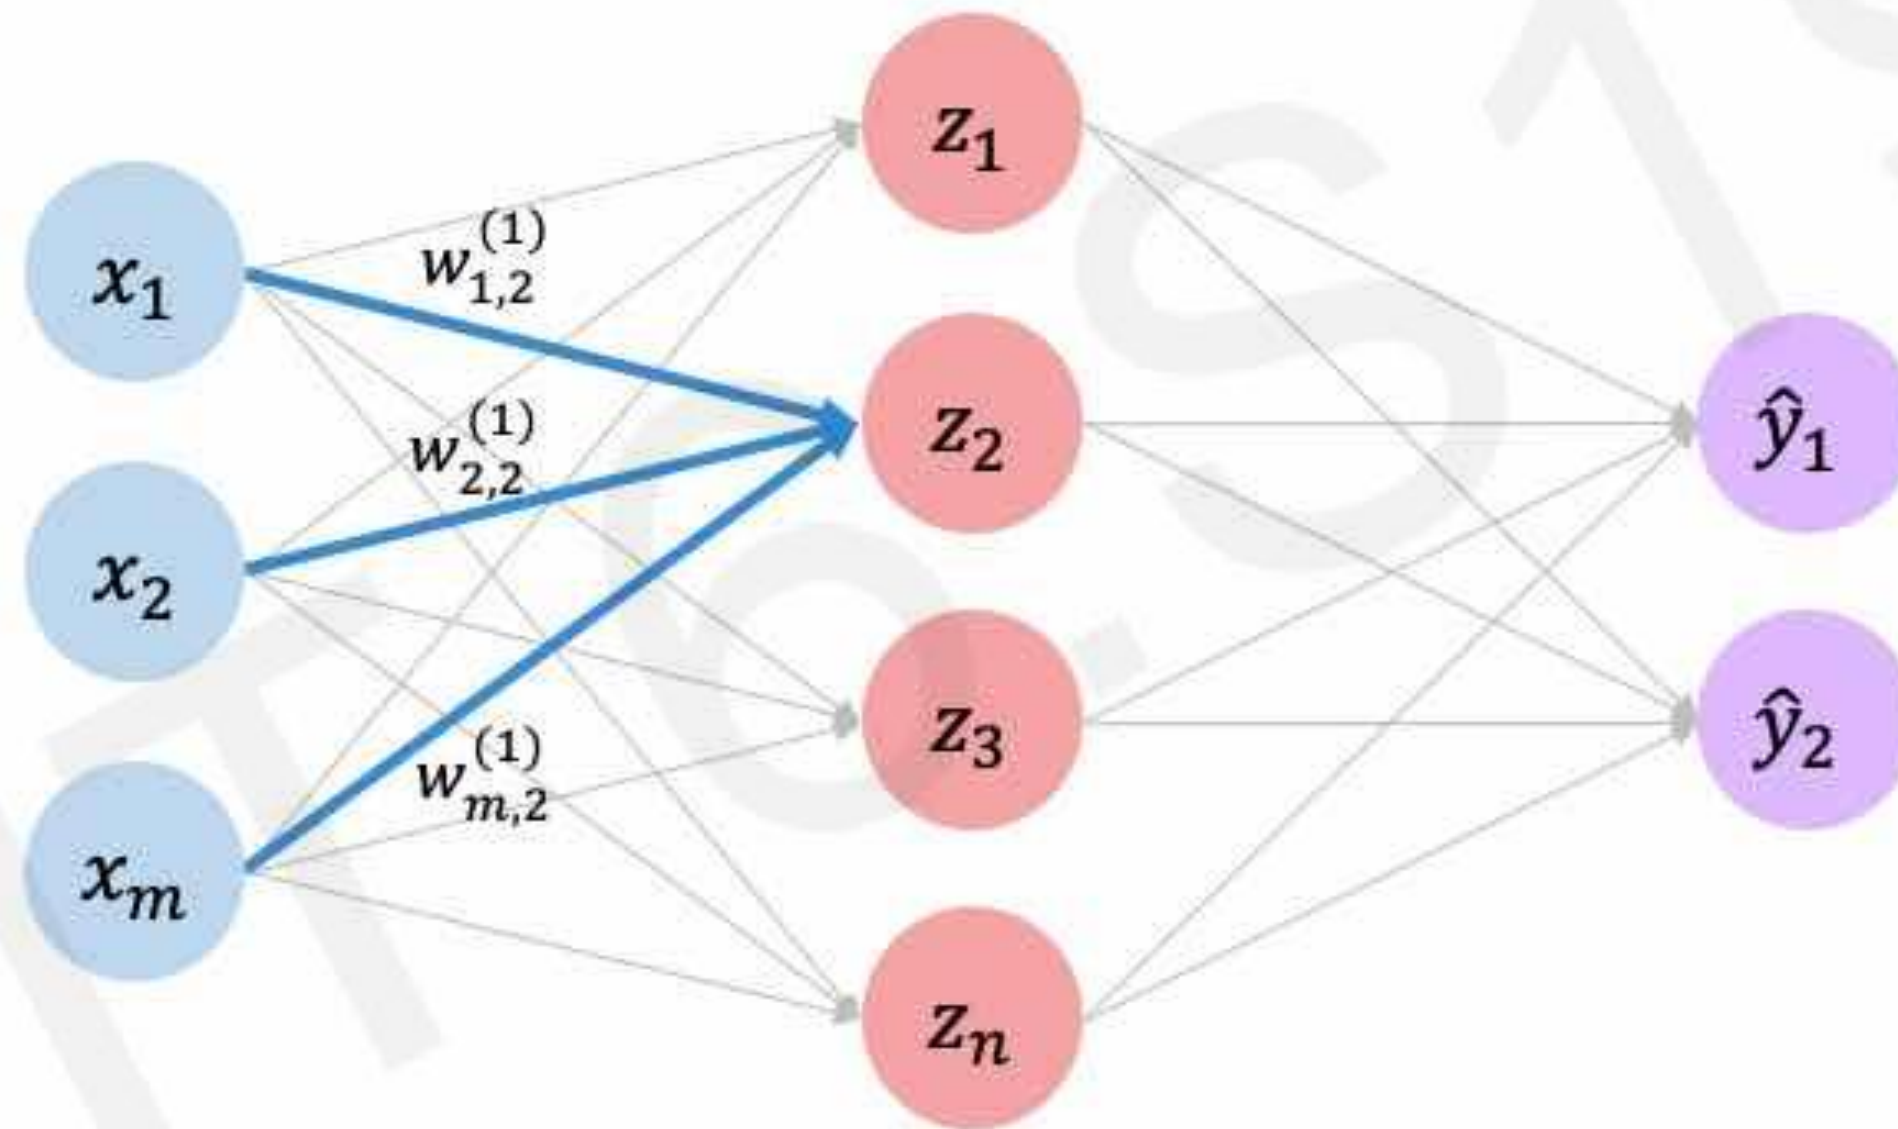
# Single Layer Neural Network



$$z_i = w_{0,i}^{(1)} + \sum_{j=1}^m x_j w_{j,i}^{(1)}$$

$$\hat{y}_i = g \left( w_{0,i}^{(2)} + \sum_{j=1}^{d_1} g(z_j) w_{j,i}^{(2)} \right)$$

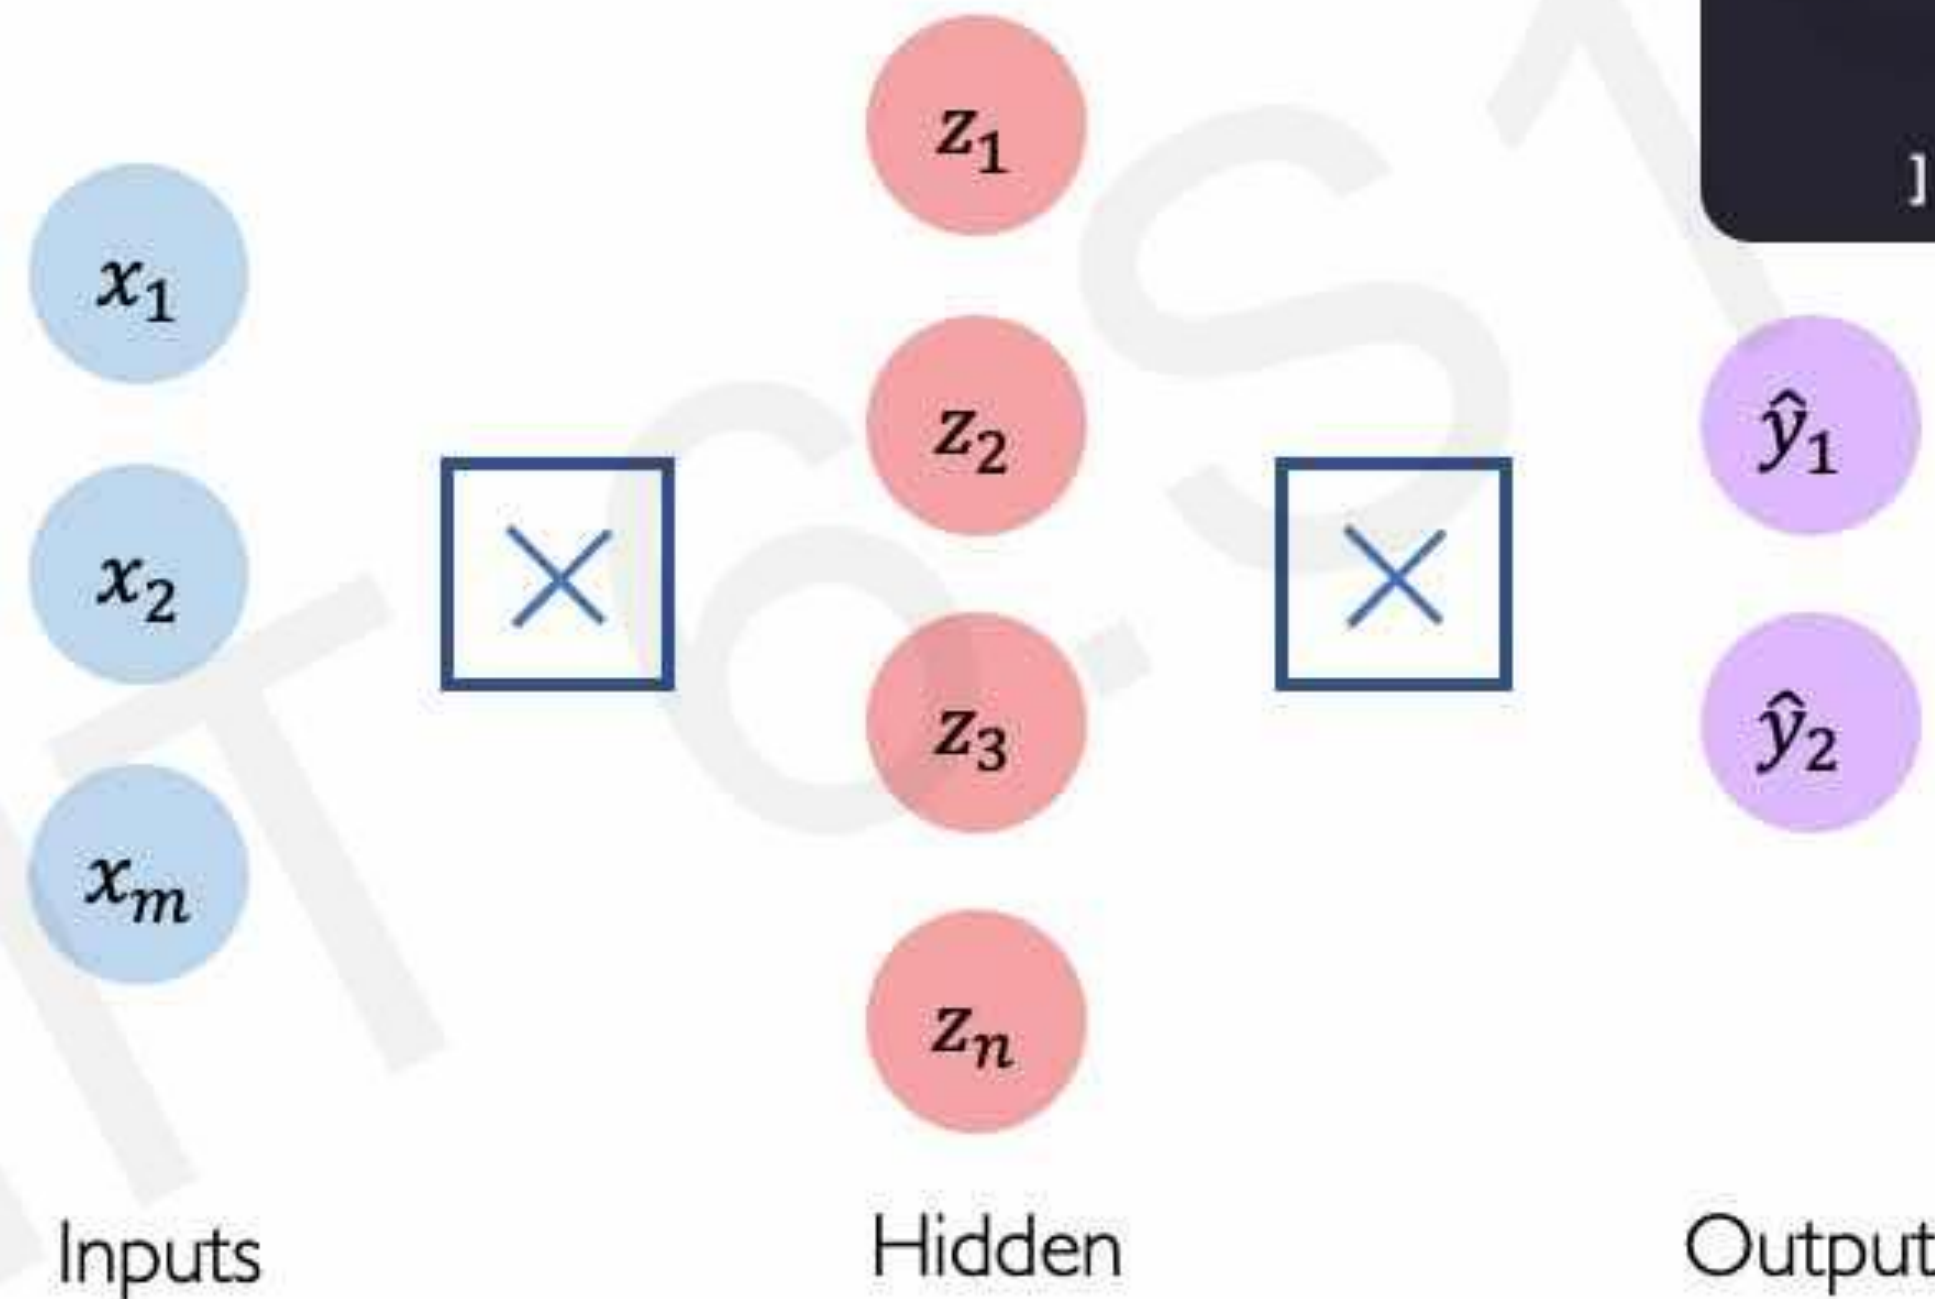
# Single Layer Neural Network



$$\begin{aligned} z_2 &= w_{0,2}^{(1)} + \sum_{j=1}^m x_j w_{j,2}^{(1)} \\ &= w_{0,2}^{(1)} + x_1 w_{1,2}^{(1)} + x_2 w_{2,2}^{(1)} + x_m w_{m,2}^{(1)} \end{aligned}$$



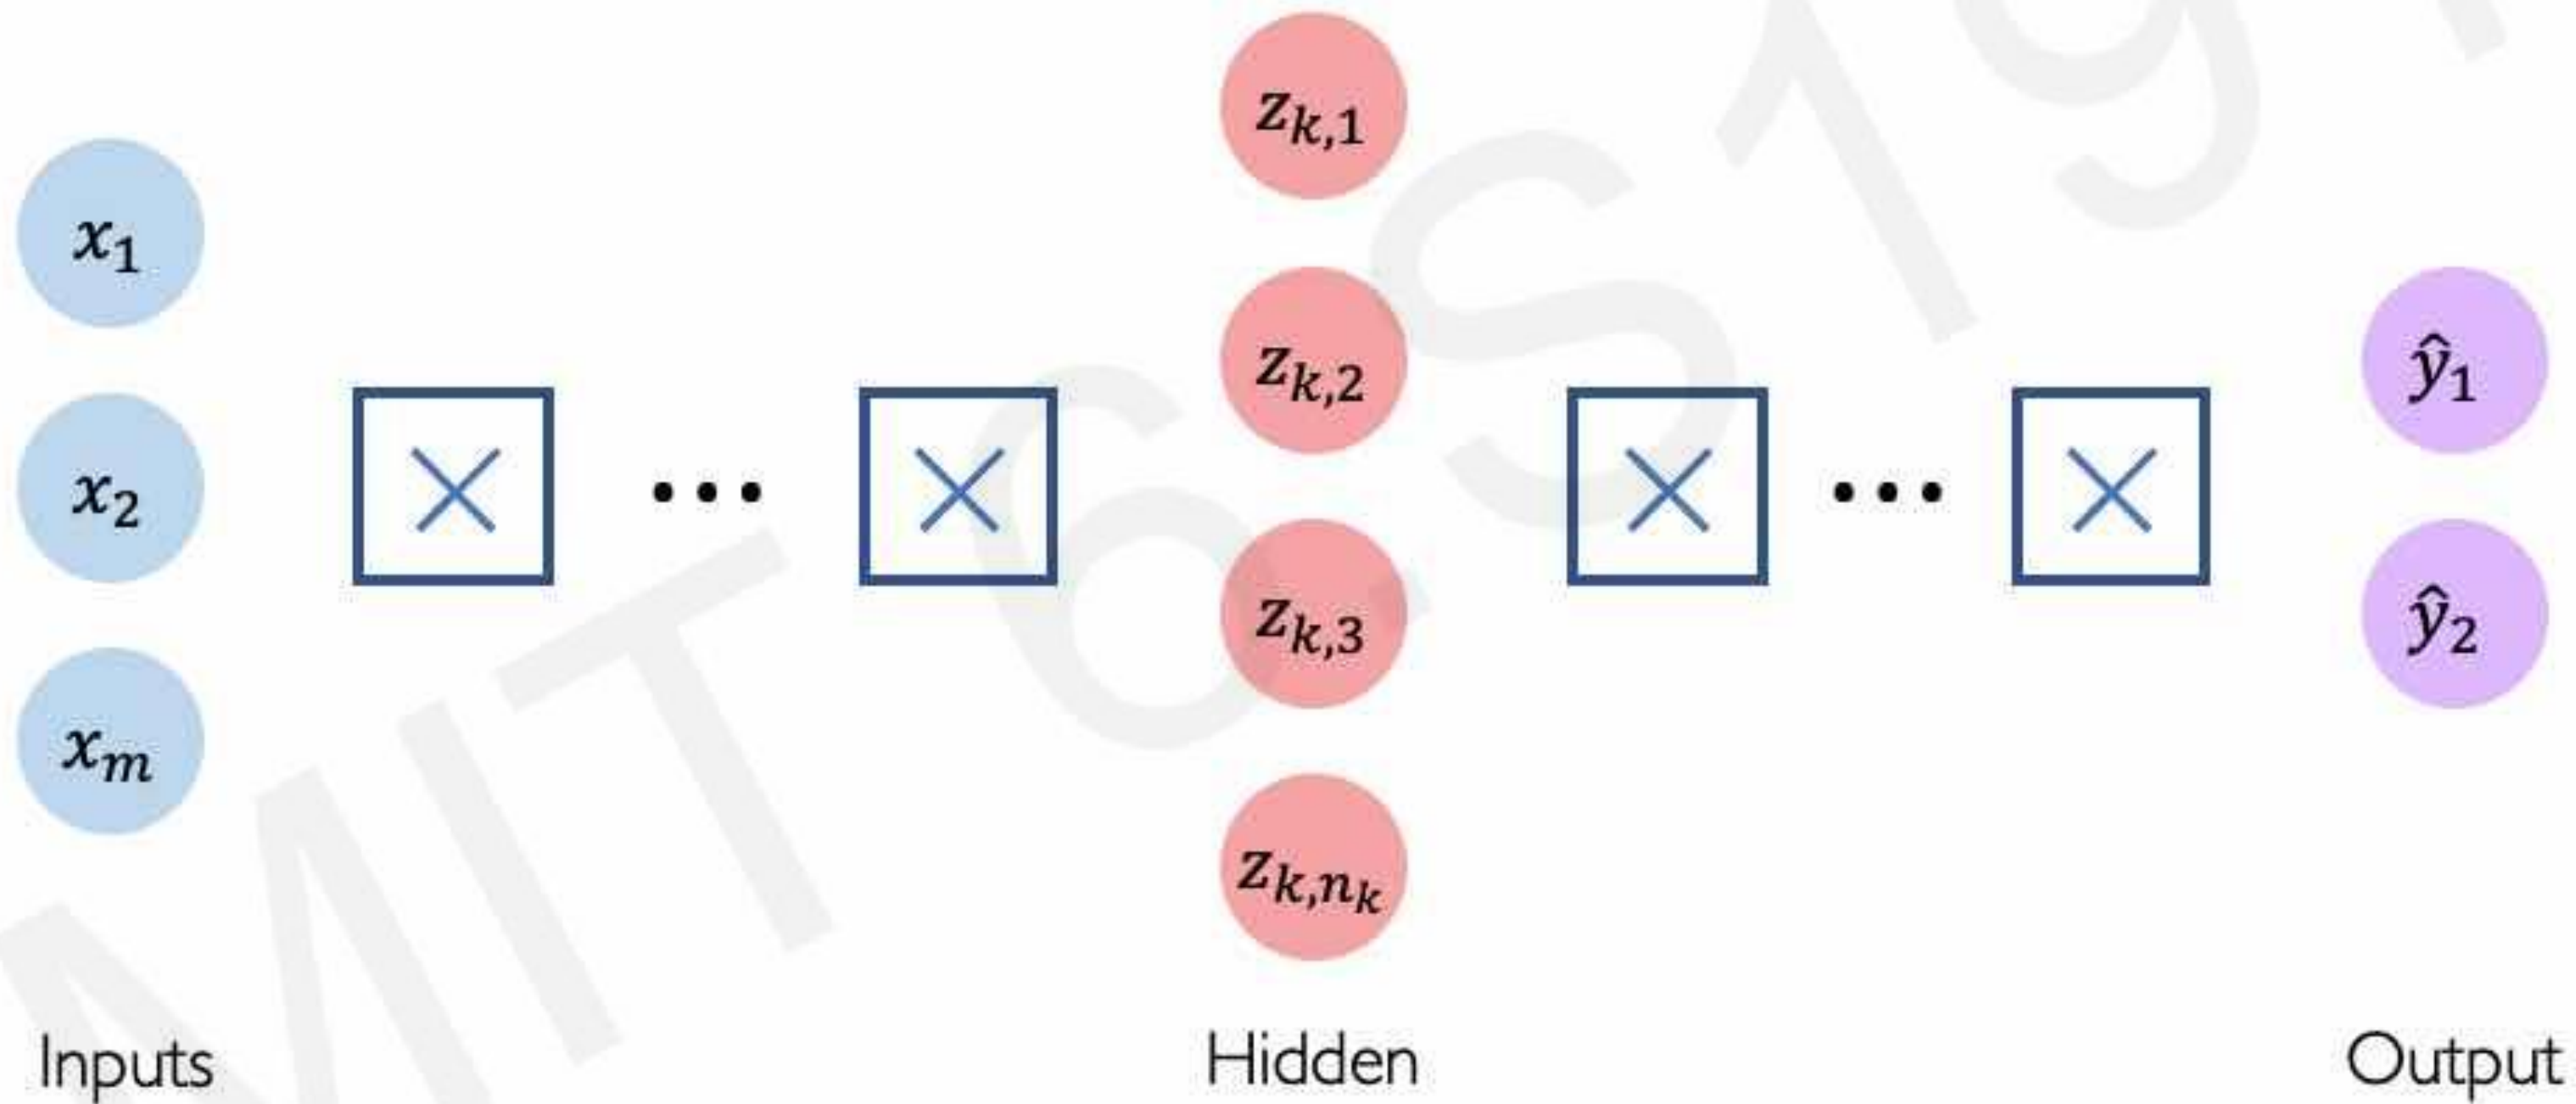
# Multi Output Perceptron



```
import tensorflow as tf

model = tf.keras.Sequential([
    tf.keras.layers.Dense(n),
    tf.keras.layers.Dense(2)
])
```

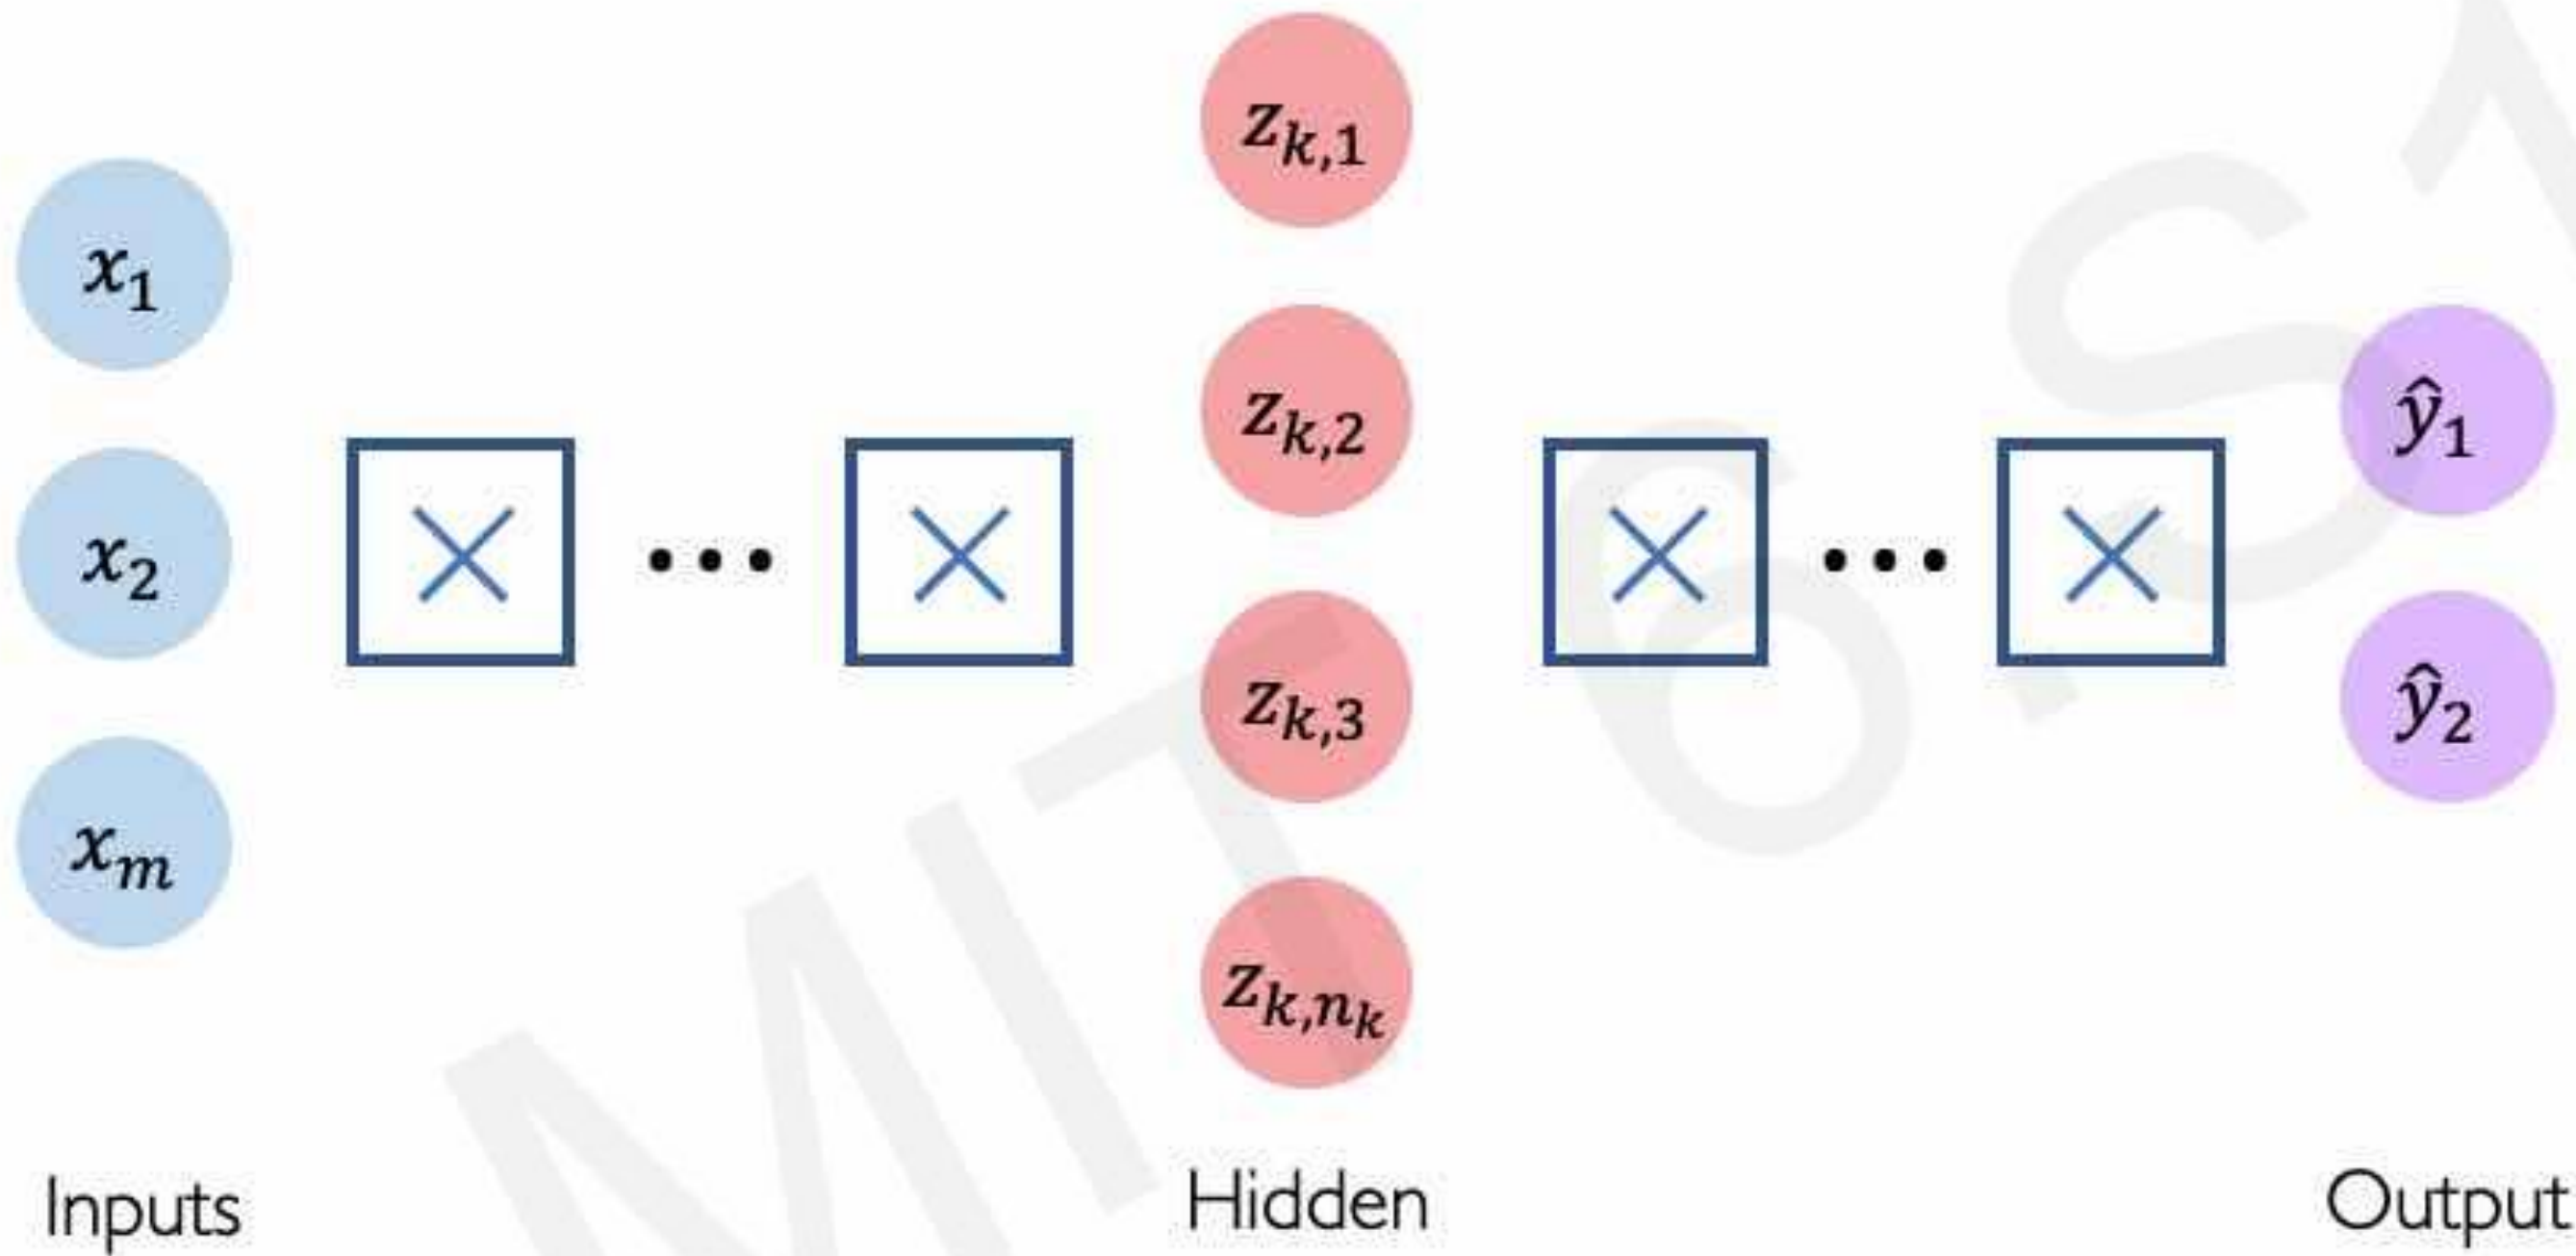
# Deep Neural Network



$$z_{k,i} = w_{0,i}^{(k)} + \sum_{j=1}^{n_{k-1}} g(z_{k-1,j}) w_{j,i}^{(k)}$$



# Deep Neural Network



```
import tensorflow as tf

model = tf.keras.Sequential([
    tf.keras.layers.Dense(n_1),
    tf.keras.layers.Dense(n_2),
    ...,
    tf.keras.layers.Dense(2)
])
```

$$z_{k,i} = w_{0,i}^{(k)} + \sum_{j=1}^{n_{k-1}} g(z_{k-1,j}) w_{j,i}^{(k)}$$

# Applying Neural Networks



# Example Problem

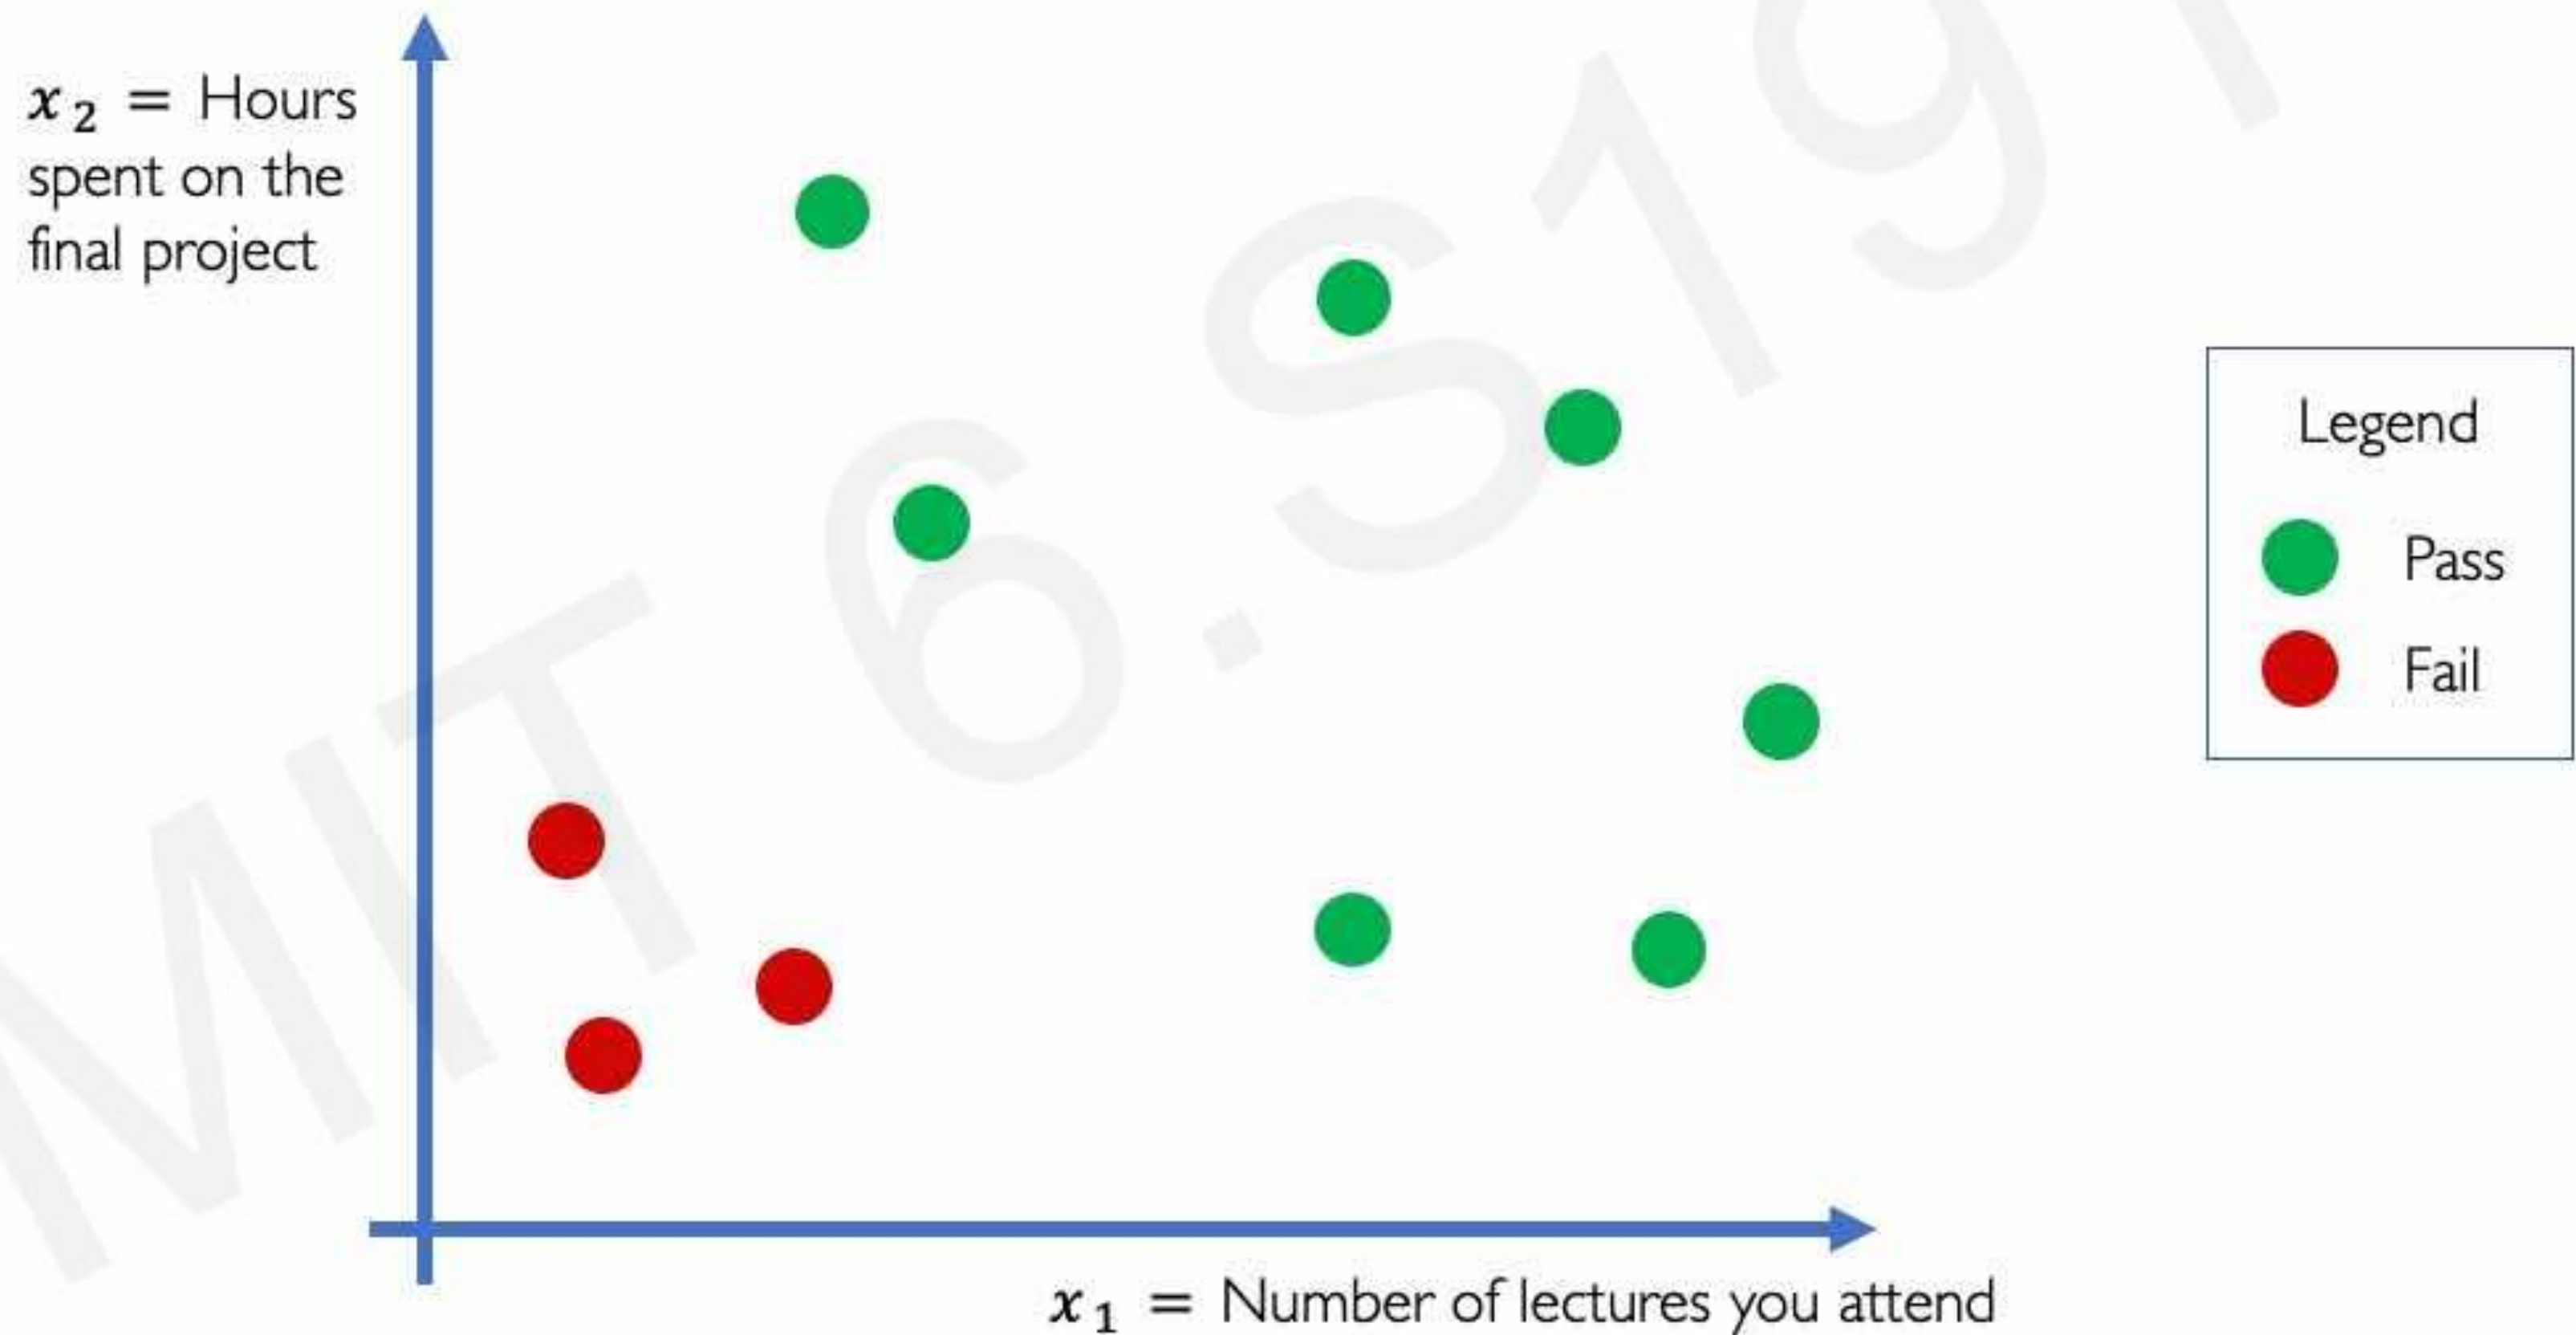
Will I pass this class?

Let's start with a simple two feature model

$x_1$  = Number of lectures you attend

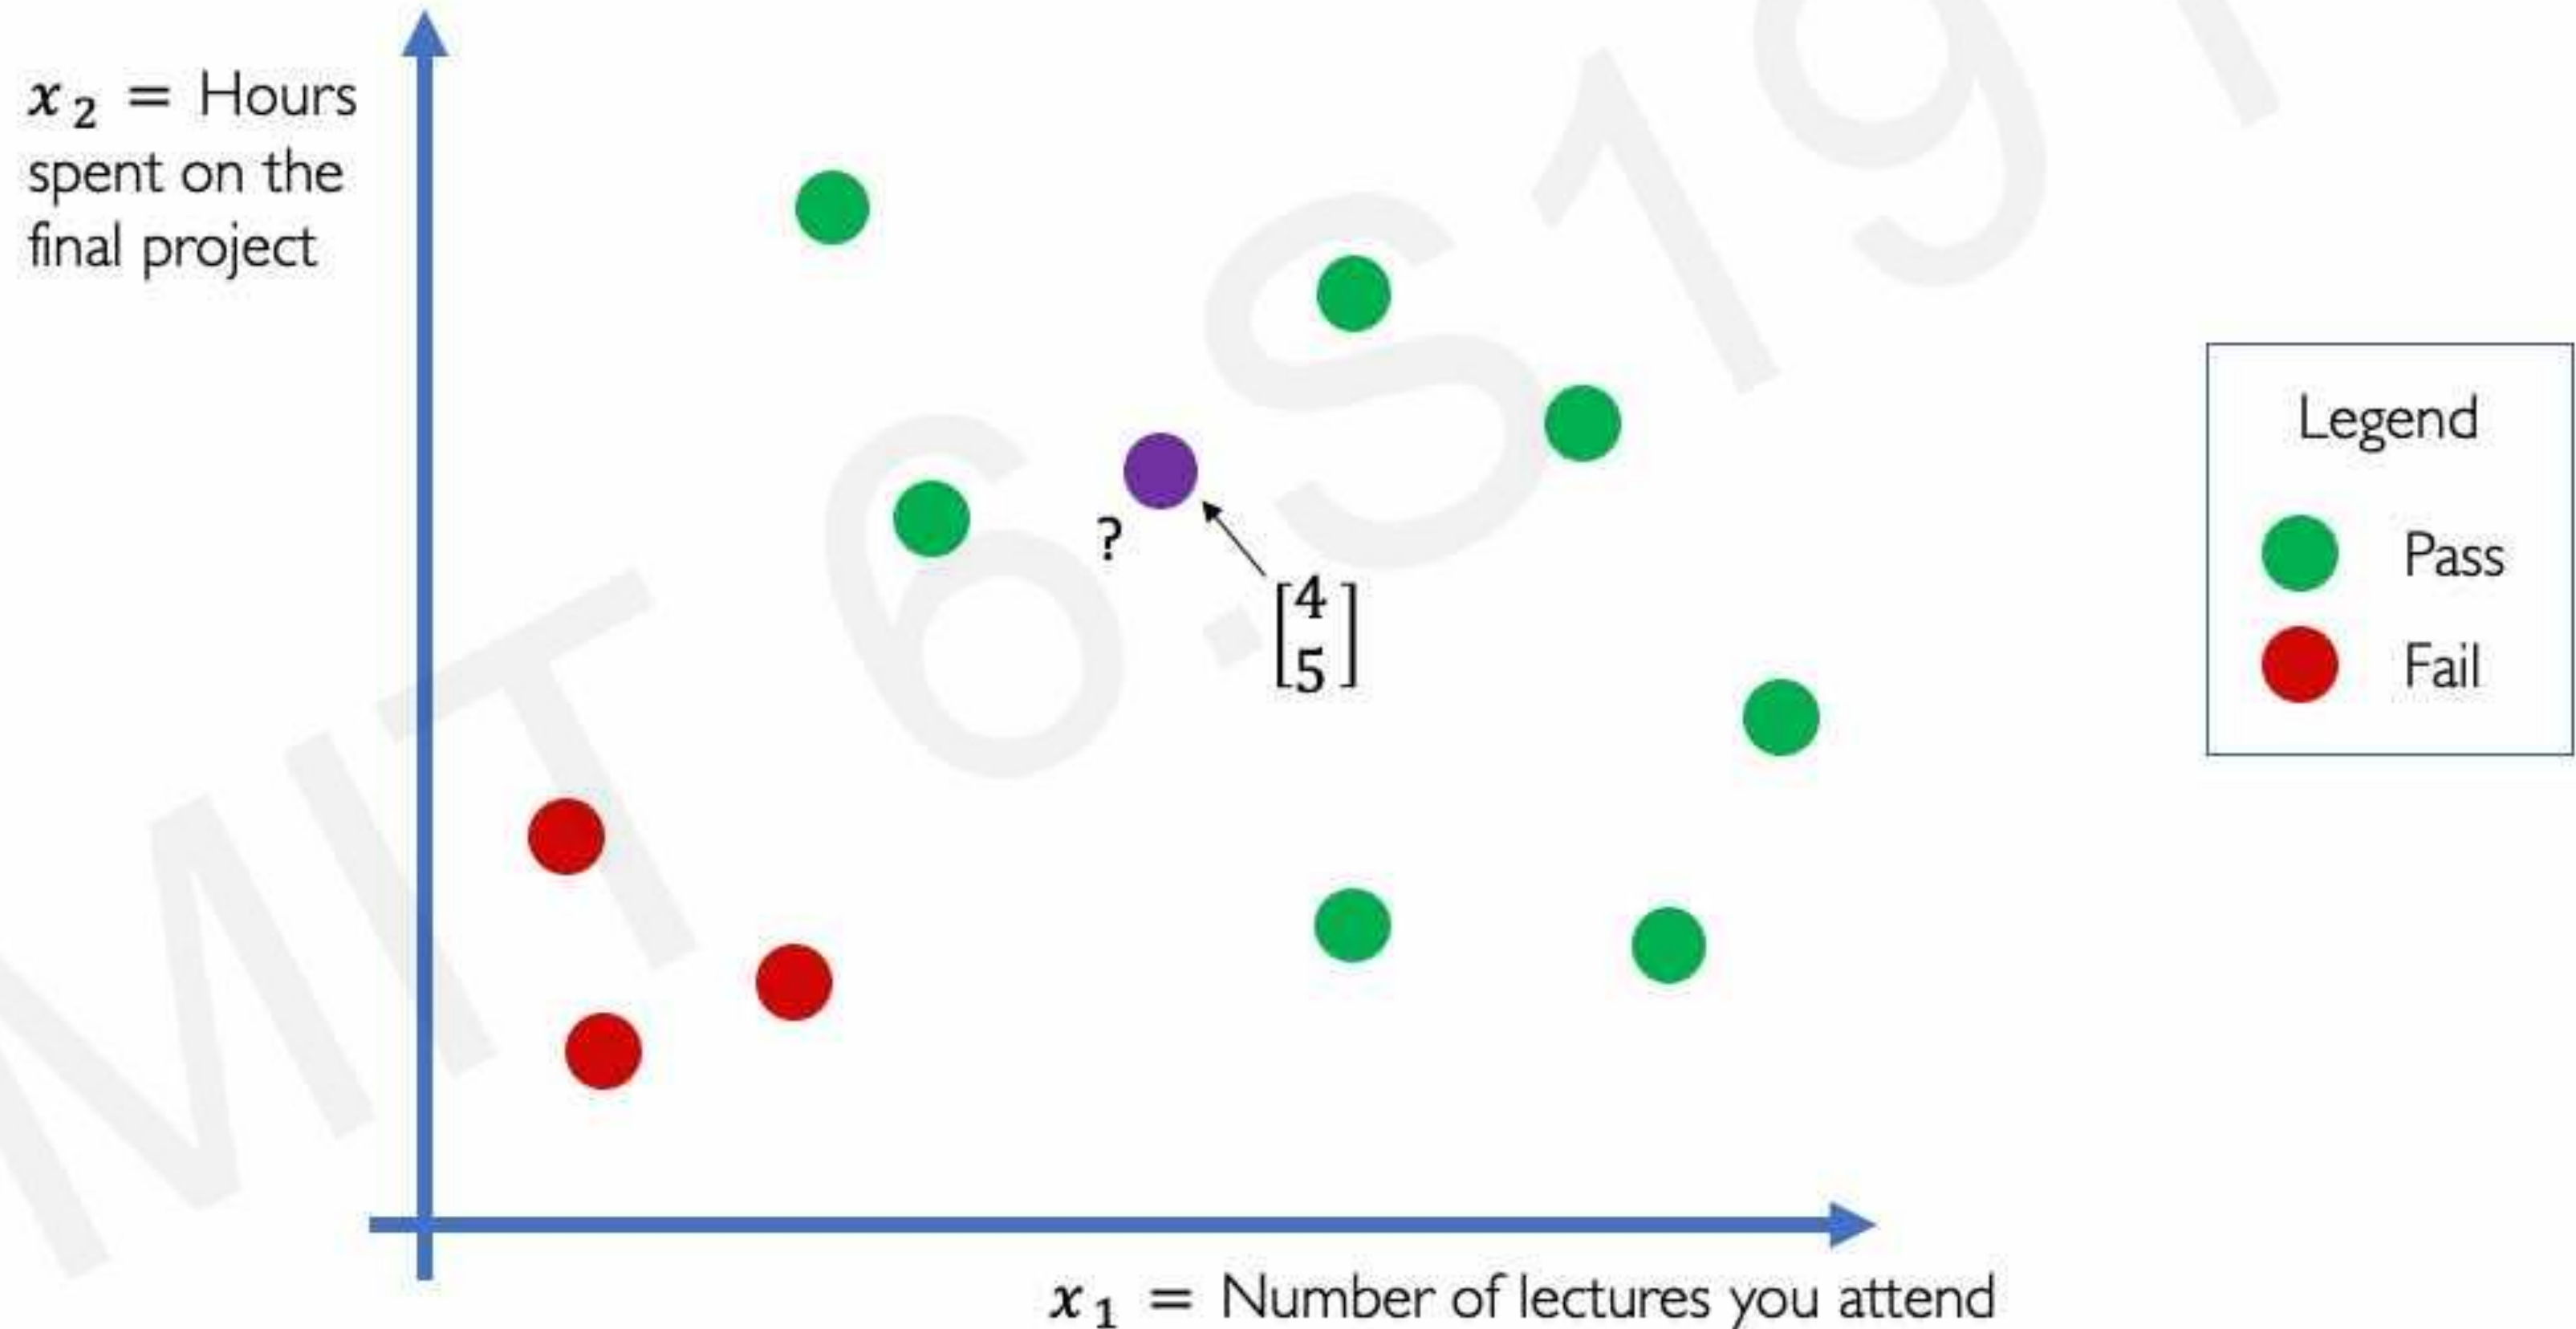
$x_2$  = Hours spent on the final project

# Example Problem: Will I pass this class?

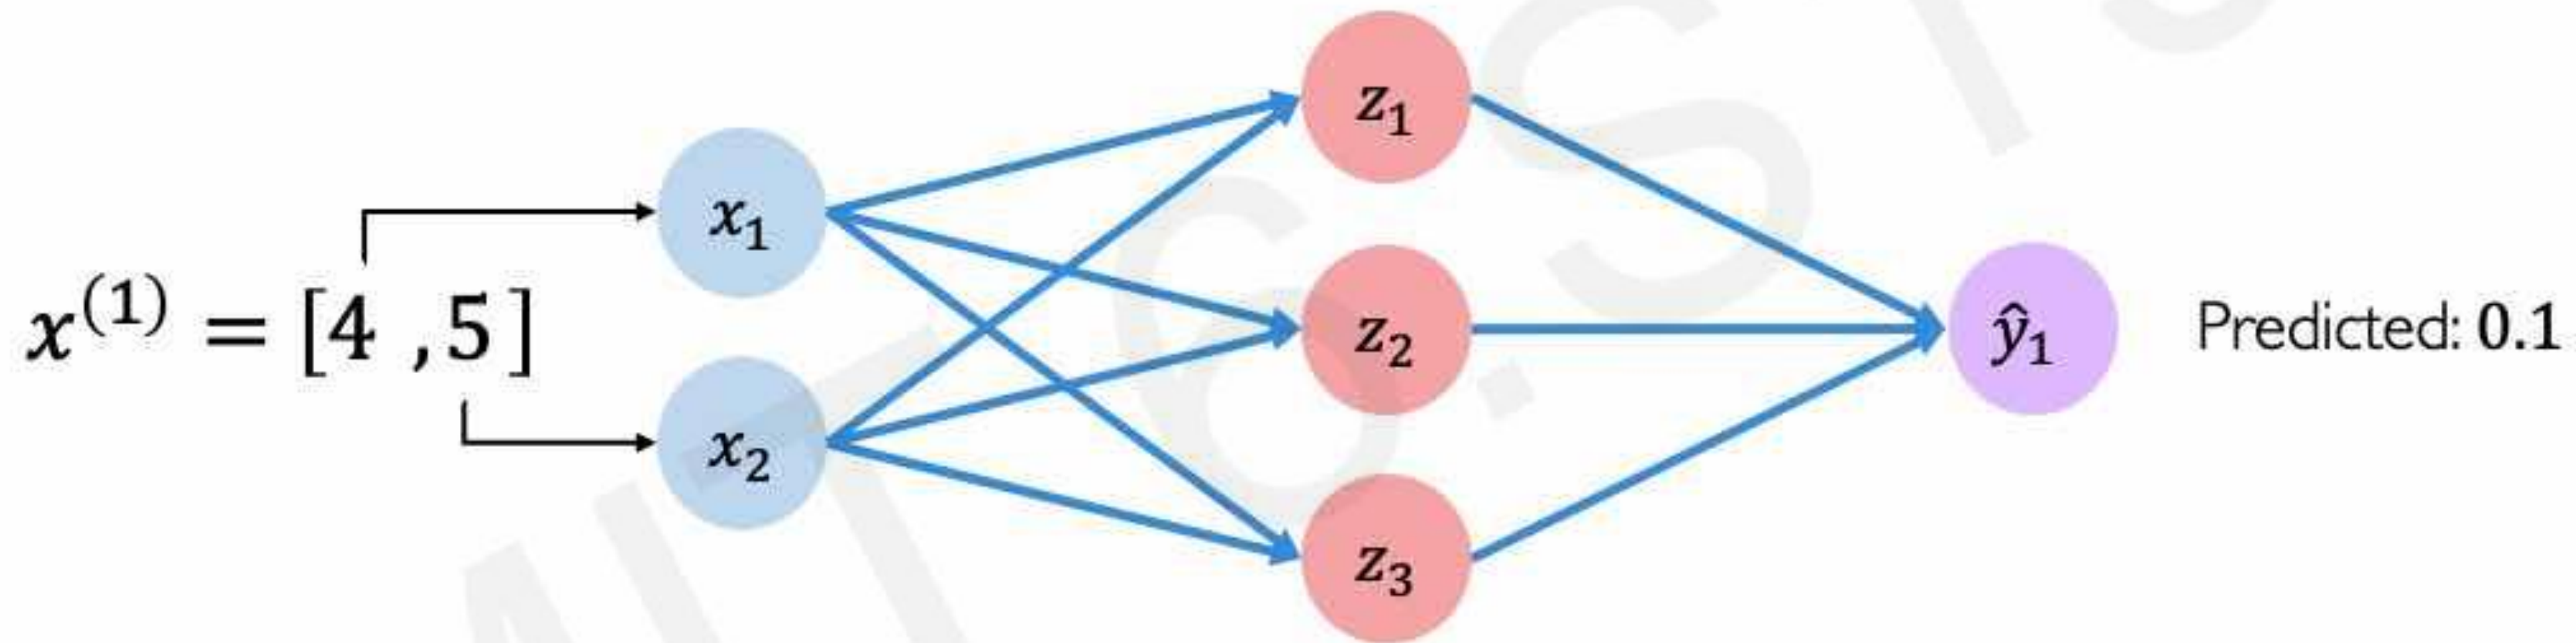




# Example Problem: Will I pass this class?

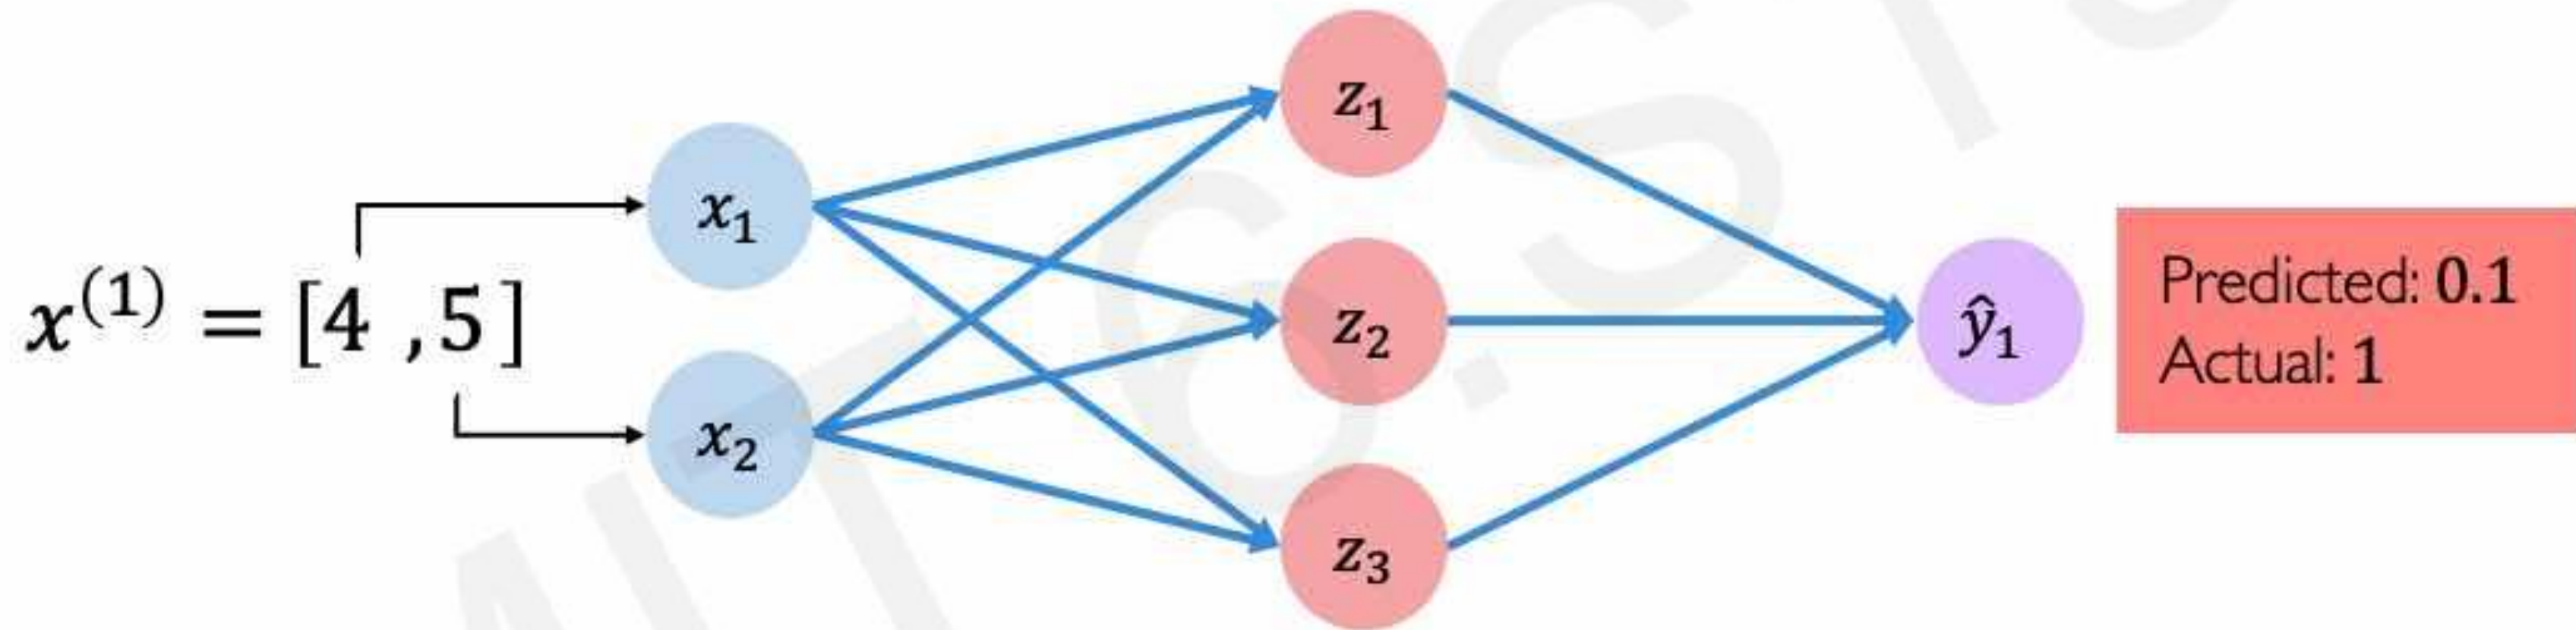


# Example Problem: Will I pass this class?



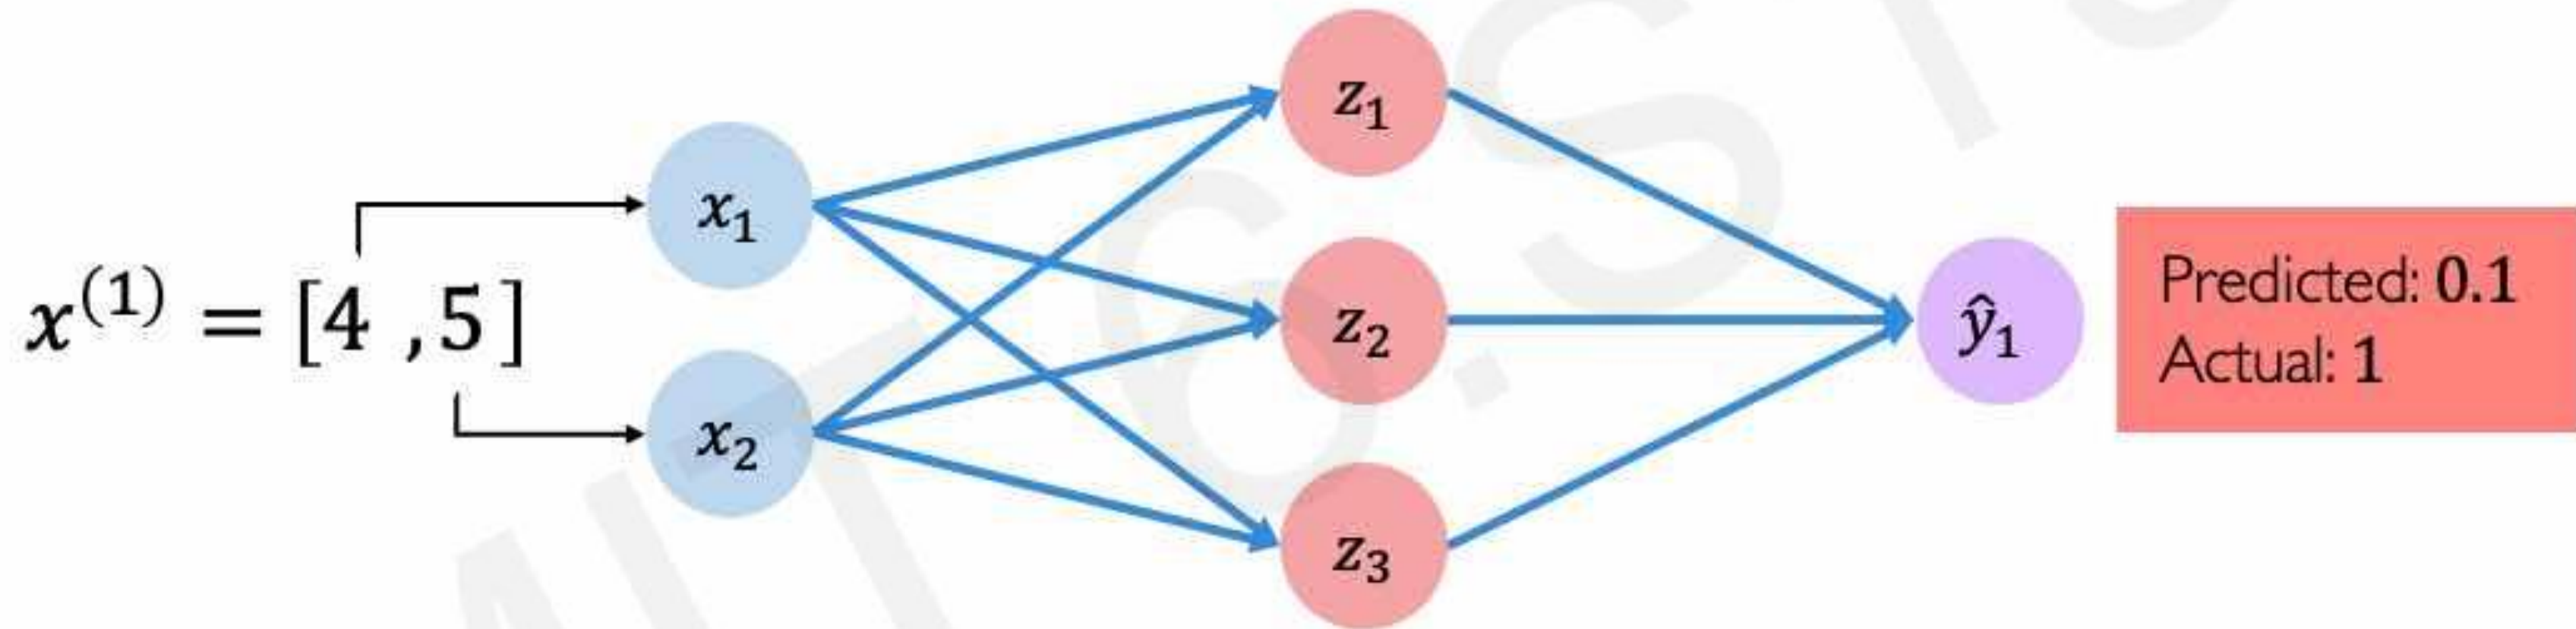


# Example Problem: Will I pass this class?



# Quantifying Loss

The **loss** of our network measures the cost incurred from incorrect predictions



$$\mathcal{L}(\underbrace{f(x^{(i)}; \mathbf{W})}_{\text{Predicted}}, \underbrace{y^{(i)}}_{\text{Actual}})$$